

Probabilistic Programming Languages

Build Your Owl PPL

Guillaume Baudart
Christine Tasson

MPRI 2021-2022

Reminders

BYO-PPL

Probabilistic Programming

Probabilistic Programming

Programming and reasoning with uncertainty

- Sample from probability distributions
- Condition on observed data

Probabilistic Programming

Programming and reasoning with uncertainty

- Sample from probability distributions
- Condition on observed data

Bayesian Inference: learn parameters from data

- Latent parameter θ
- Observed data x_1, \dots, x_n



Thomas Bayes (1701-1761)

Probabilistic Programming

Programming and reasoning with uncertainty

- Sample from probability distributions
- Condition on observed data

Bayesian Inference: learn parameters from data

- Latent parameter θ
- Observed data x_1, \dots, x_n

$$p(\theta \mid x_1, \dots, x_n) = \frac{p(\theta) p(x_1, \dots, x_n \mid \theta)}{p(x_1, \dots, x_n)} \quad (\text{Bayes' theorem})$$

$$\propto p(\theta) p(x_1, \dots, x_n \mid \theta) \quad (\text{Data are constants})$$



Thomas Bayes (1701-1761)

Probabilistic Programming

Programming and reasoning with uncertainty

- Sample from probability distributions
- Condition on observed data

Bayesian Inference: learn parameters from data

- Latent parameter θ
- Observed data x_1, \dots, x_n

$$p(\theta \mid x_1, \dots, x_n) = \frac{p(\theta) p(x_1, \dots, x_n \mid \theta)}{p(x_1, \dots, x_n)} \quad (\text{Bayes' theorem})$$

posterior

$$\propto p(\theta) p(x_1, \dots, x_n \mid \theta) \quad (\text{Data are constants})$$



Thomas Bayes (1701-1761)

Probabilistic Programming

Programming and reasoning with uncertainty

- Sample from probability distributions
- Condition on observed data

Bayesian Inference: learn parameters from data

- Latent parameter θ
- Observed data x_1, \dots, x_n

$$p(\theta \mid x_1, \dots, x_n) = \frac{p(\theta) p(x_1, \dots, x_n \mid \theta)}{p(x_1, \dots, x_n)} \quad (\text{Bayes' theorem})$$

posterior

$$\propto p(\theta) p(x_1, \dots, x_n \mid \theta) \quad (\text{Data are constants})$$

prior



Thomas Bayes (1701-1761)

Probabilistic Programming

Programming and reasoning with uncertainty

- Sample from probability distributions
- Condition on observed data

Bayesian Inference: learn parameters from data

- Latent parameter θ
- Observed data x_1, \dots, x_n

$$p(\theta \mid x_1, \dots, x_n) = \frac{p(\theta) p(x_1, \dots, x_n \mid \theta)}{p(x_1, \dots, x_n)} \quad (\text{Bayes' theorem})$$

posterior

$$\propto p(\theta) p(x_1, \dots, x_n \mid \theta) \quad (\text{Data are constants})$$

prior

likelihood



Thomas Bayes (1701-1761)

Example: Coin



Consider a series of coin tosses

- Observations: head or tail
- Each toss is independant
- What is the probability of getting head at the next toss?

Probabilistic model

- Prior: $z \sim \text{Uniform}(0,1)$
- Observations: $\forall i \in [1, n] . x_i \sim \text{Bernoulli}(z)$
- Posterior: $p(z | x_1, x_2, \dots, x_n)$?

Example: Coin



Consider a series of coin tosses

- Observations: head or tail
- Each toss is independant
- What is the probability of getting head at the next toss?

Probabilistic model

- Prior: $z \sim \text{Uniform}(0,1)$
- Observations: $\forall i \in [1, n]. x_i \sim \text{Bernoulli}(z)$
- Posterior: $p(z | x_1, x_2, \dots, x_n)$?

$$\begin{aligned} p(z | x_1, \dots, x_n) &= \frac{p(x_1, \dots, x_n | z)p(z)}{p(x_1, \dots, x_n)} \\ &= \frac{p(x_1, \dots, x_n | z)p(z)}{\int_z p(x_1, \dots, x_n | z)} \end{aligned}$$

$$\begin{aligned} p(x_1, \dots, x_n | z) &= \prod_{i=1}^n p(x_i | z) \\ &= \prod_{i=1}^n z^{x_i} (1 - z)^{1-x_i} \\ &= z^{\sum_{i=1}^n x_i} (1 - z)^{\sum_{i=1}^n (1-x_i)} \\ &= z^{\text{\#heads}} (1 - z)^{\text{\#tails}} \\ p(z | x_1, \dots, x_n) &= \frac{z^{\text{\#heads}} (1 - z)^{\text{\#tails}}}{\int_z z^{\text{\#heads}} (1 - z)^{\text{\#tails}}} \\ &= \frac{z^{\text{\#heads}} (1 - z)^{\text{\#tails}}}{B(\text{\#heads} + 1, \text{\#tails} + 1)} \\ &= \text{Beta}_{\text{pdf}}(\text{\#heads} + 1, \text{\#tails} + 1) \end{aligned}$$

Example: Coin



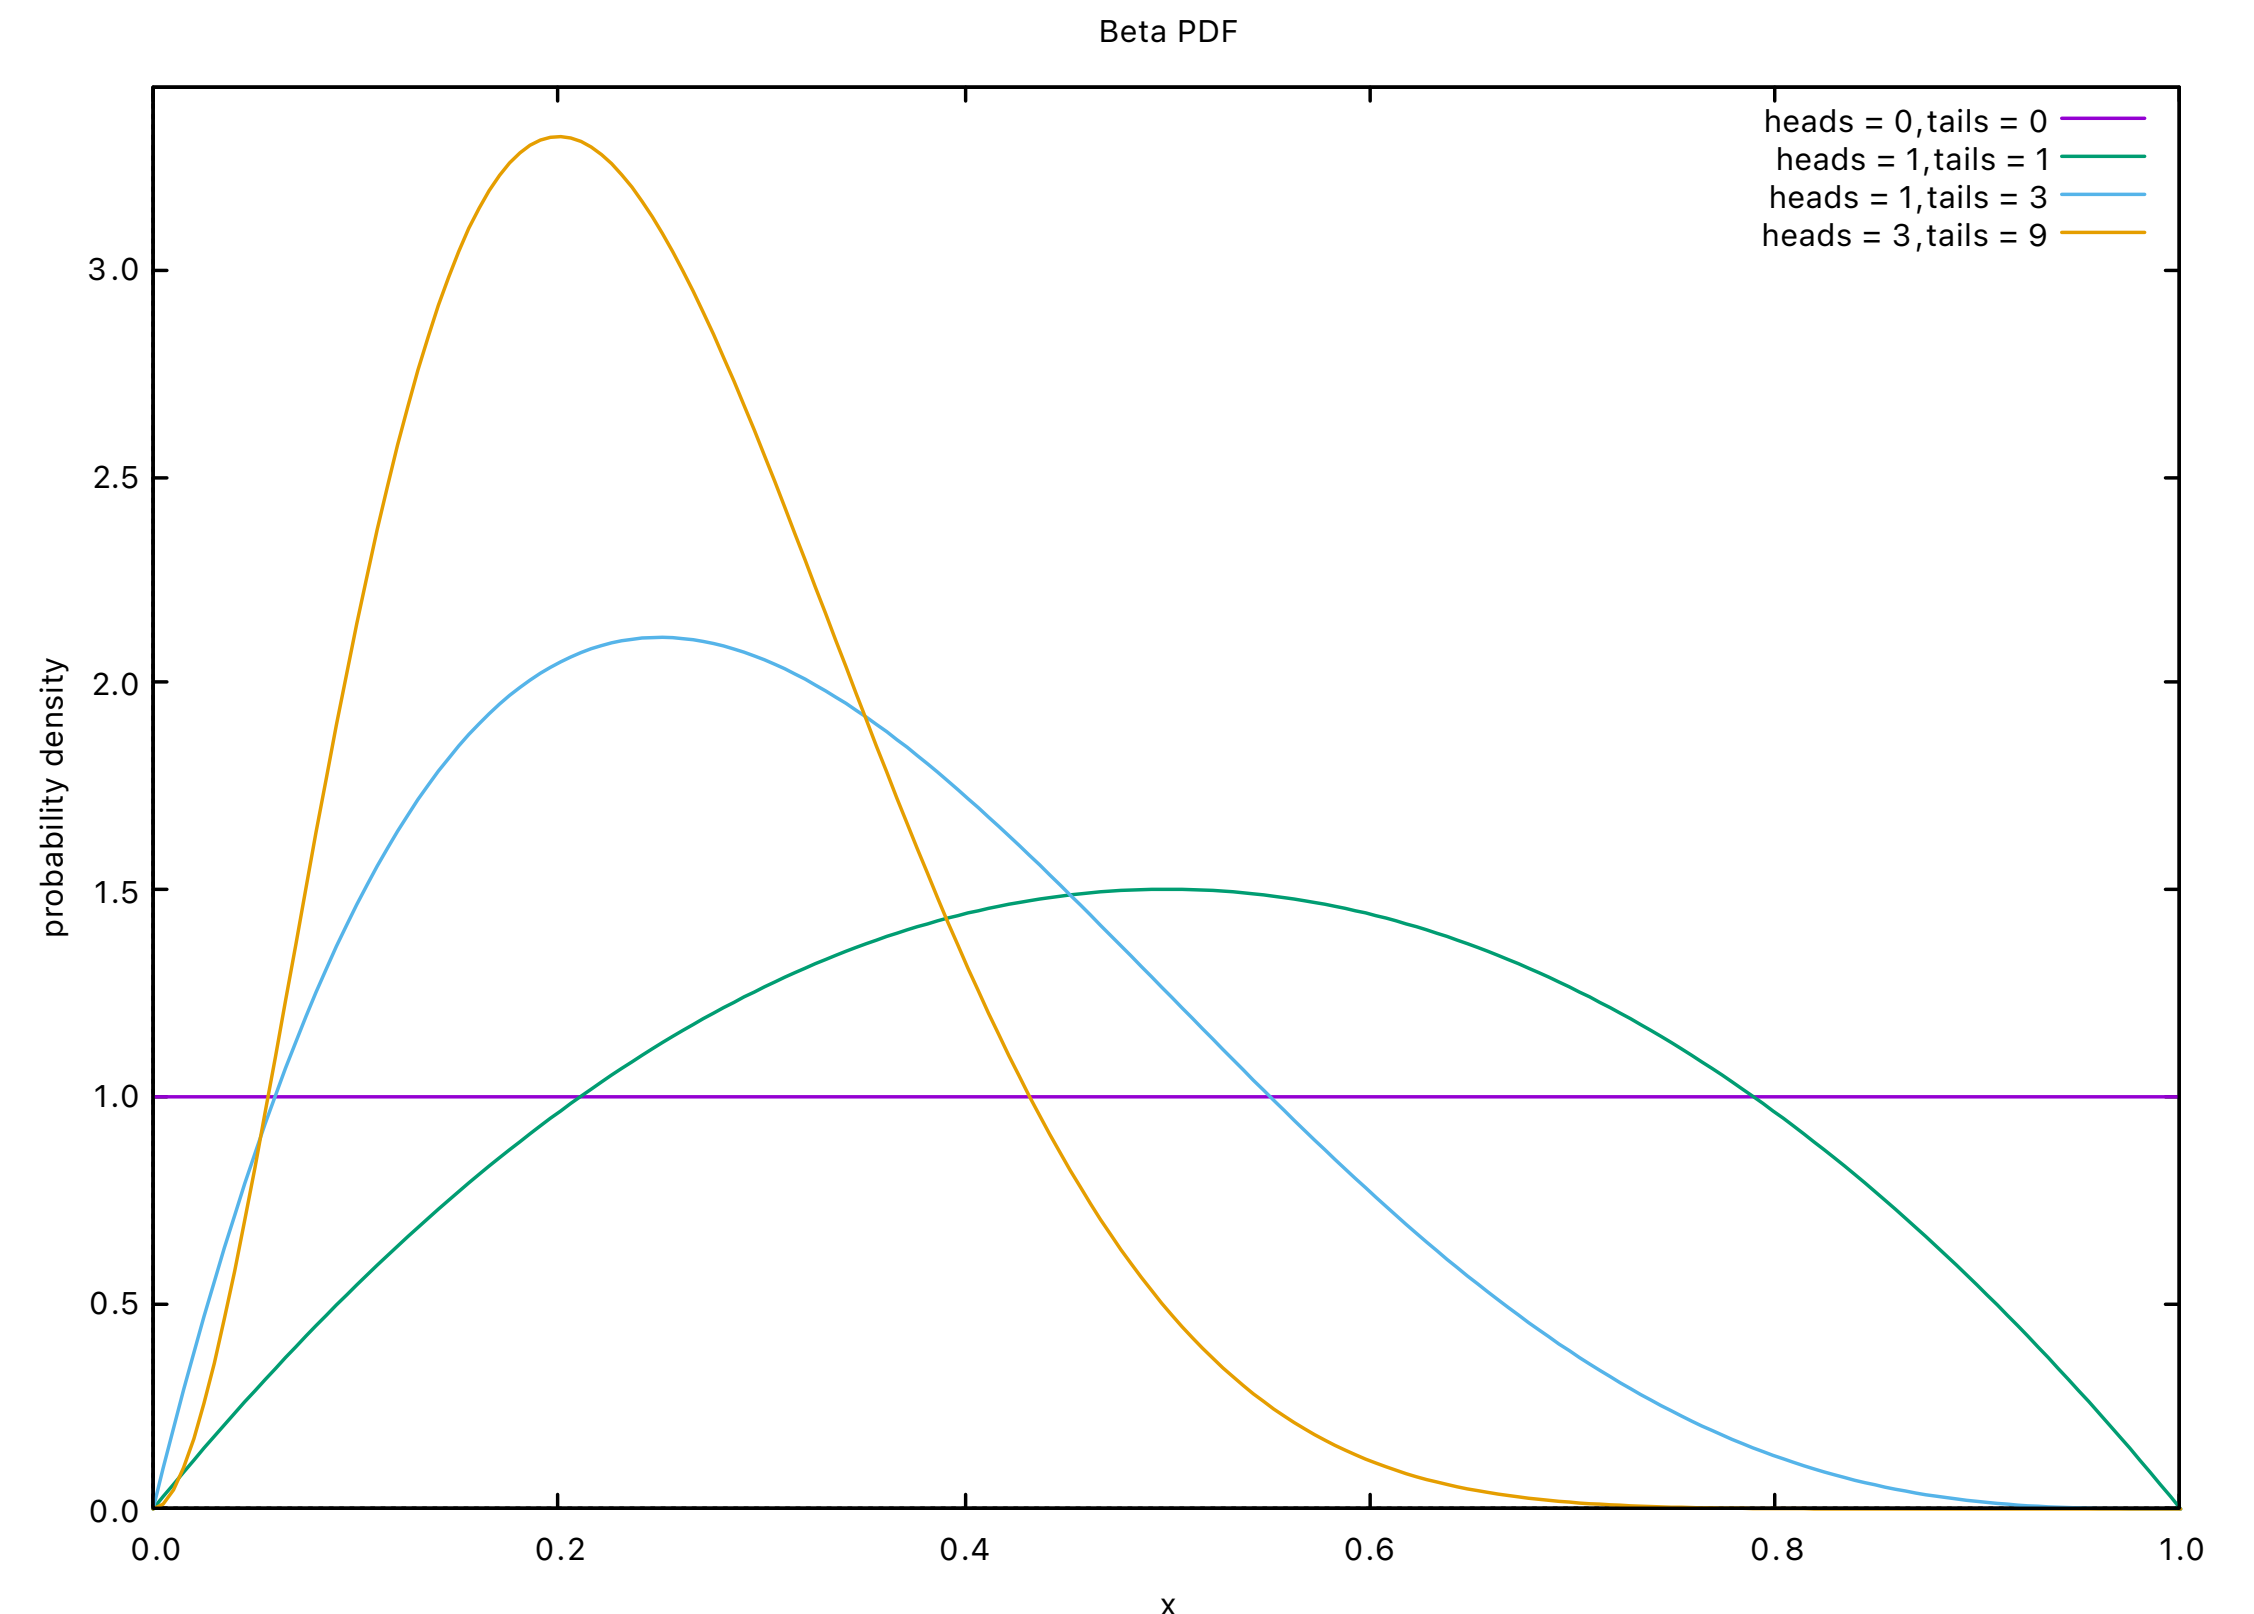
Consider a series of coin tosses

- Observations: head or tail
- Each toss is independant
- What is the probability of getting head at the next toss?

Probabilistic model

- Prior: $z \sim \text{Uniform}(0,1)$
- Observations: $\forall i \in [1, n]. x_i \sim \text{Bernoulli}(z)$
- Posterior: $p(z | x_1, x_2, \dots, x_n)$?

$$z \sim \text{Beta}(\text{\#heads} + 1, \text{\#tails} + 1)$$



Example: Coin



Consider a series of coin tosses

- Observations: head or tail
- Each toss is independant
- What is the probability of getting head at the next toss?

Probabilistic model

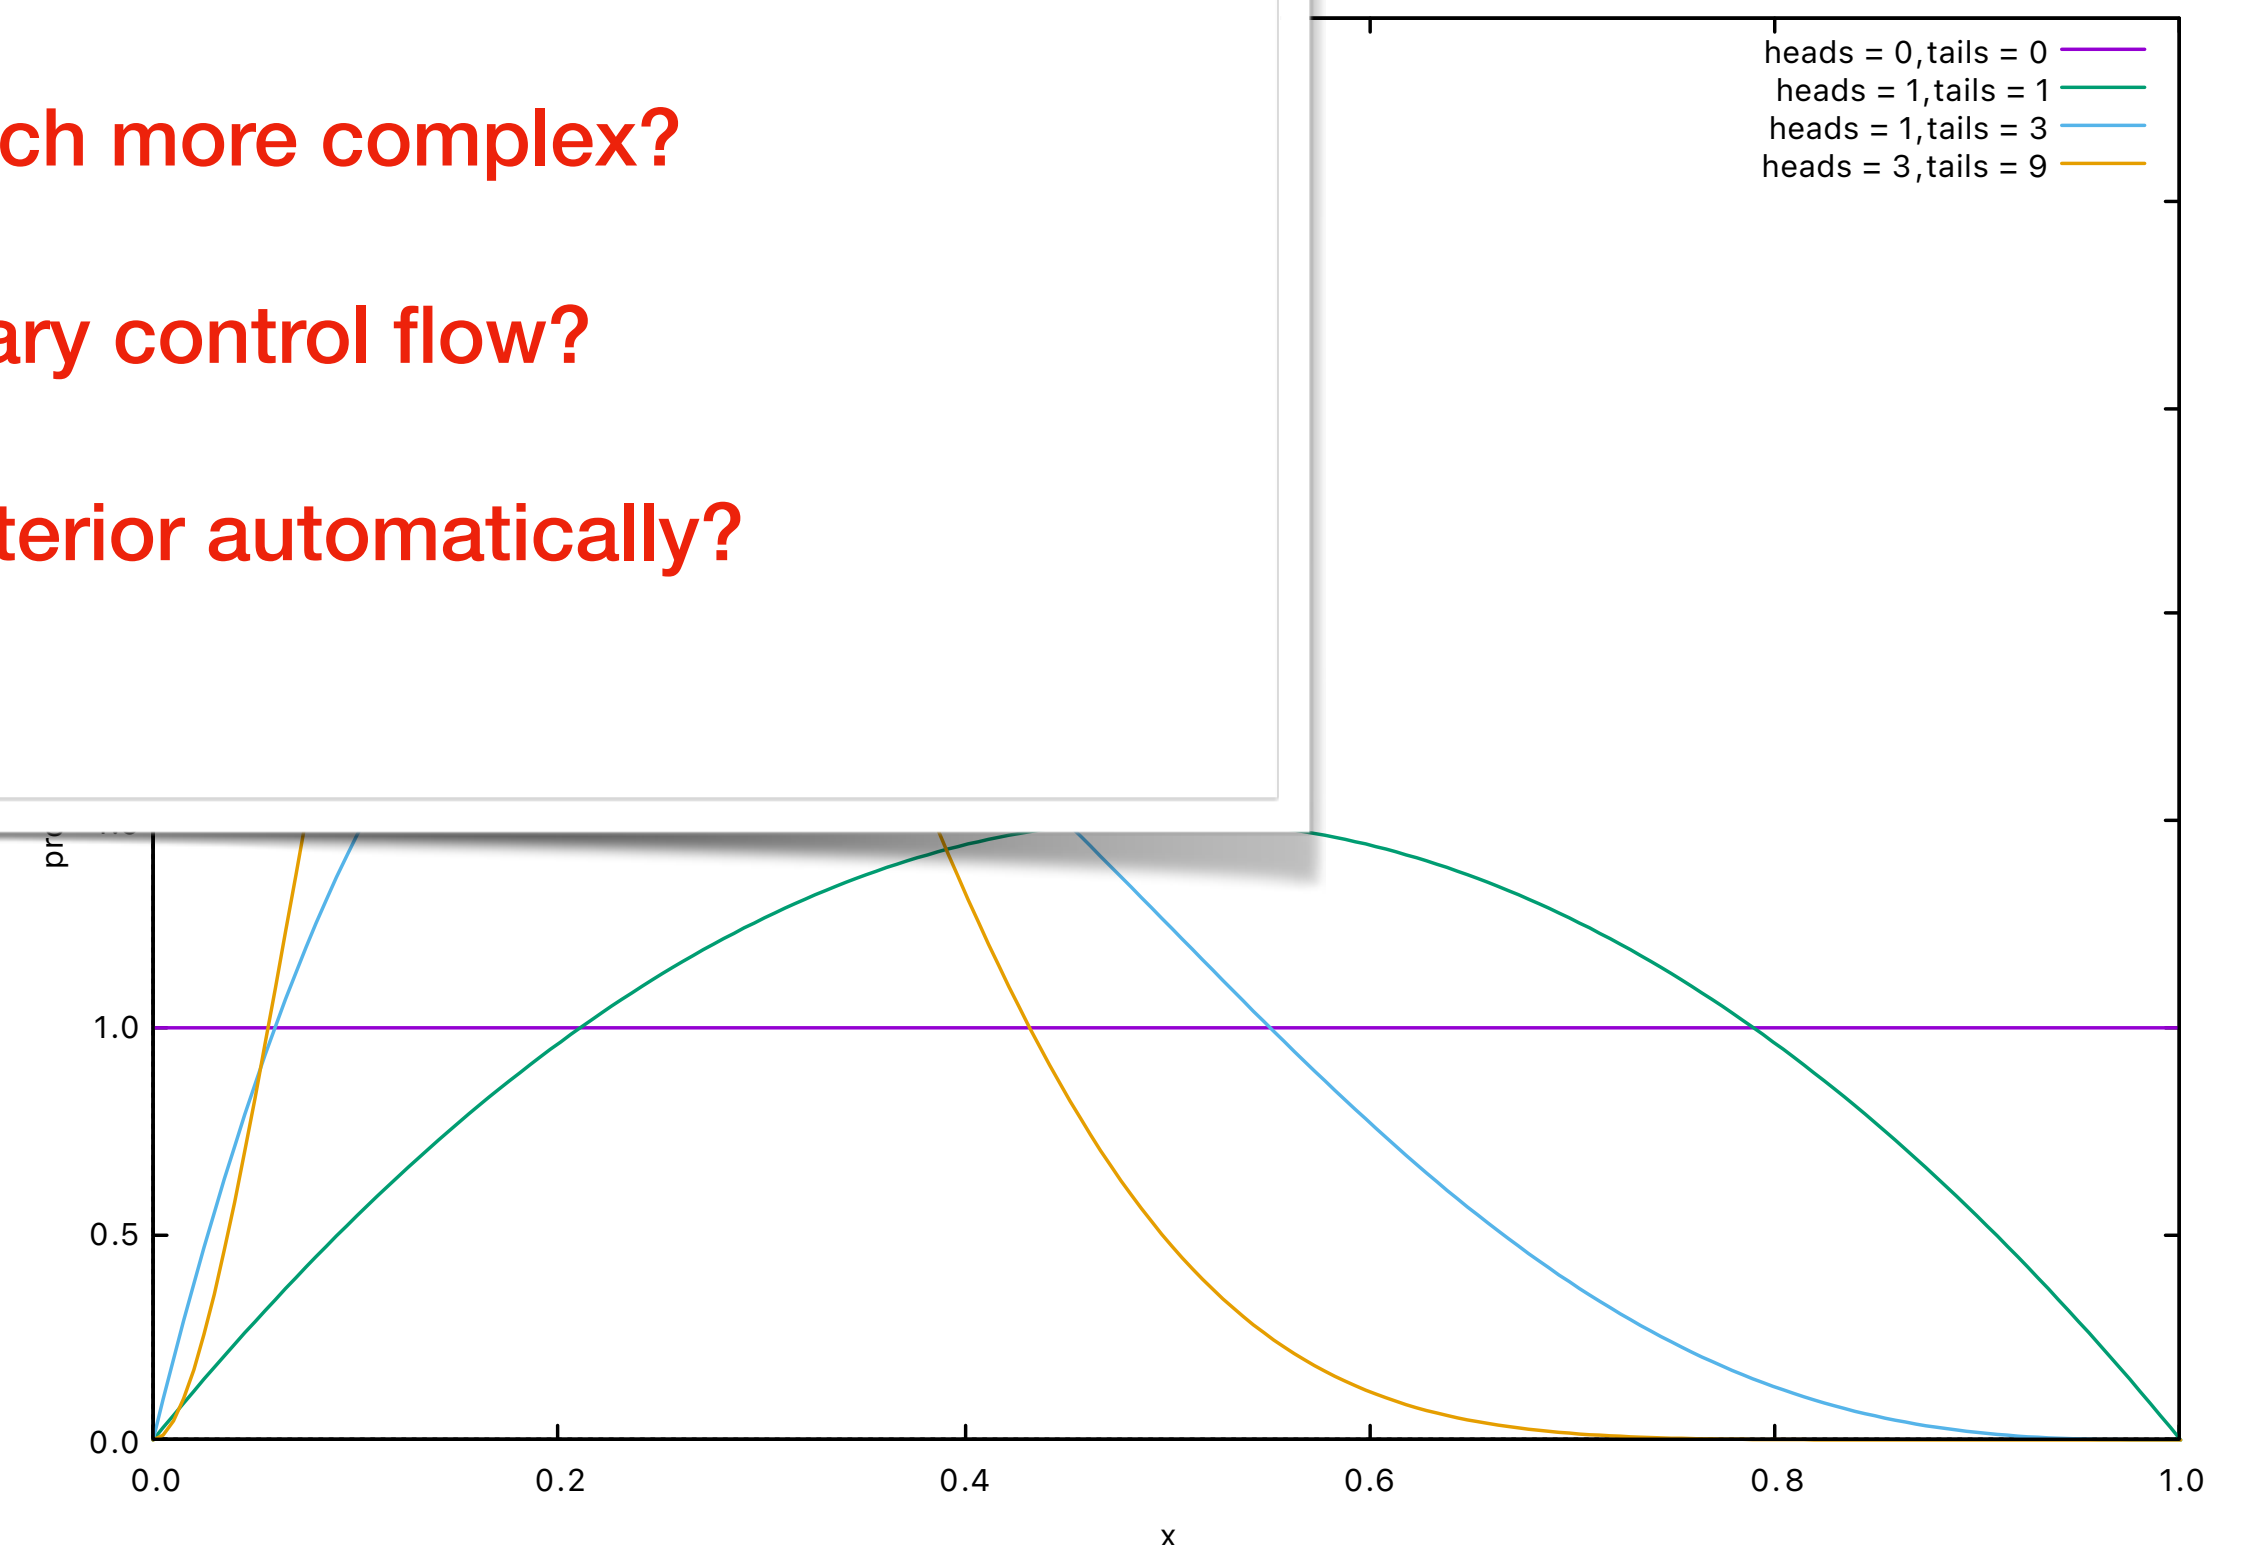
- Prior: $z \sim \text{Uniform}$
- Observations: $\forall i \in$
- Posterior: $p(z | x_1, x$

What if the model is much more complex?

What if we use arbitrary control flow?

Can we compute the posterior automatically?

$z \sim \text{Beta}(\text{\#heads} + 1, \text{\#tails} + 1)$



Probabilistic Programming Languages

Probabilistic Programming Languages

General purpose programming languages extended with probabilistic constructs

- `sample`: draw a sample from a distribution
- `assume`, `factor`, `observe`: condition the model on inputs (e.g., observed data)
- `infer`: compute the posterior distribution of a model given the inputs

Probabilistic Programming Languages

General purpose programming languages extended with probabilistic constructs

- **sample**: draw a sample from a distribution
- **assume**, **factor**, **observe**: condition the model on inputs (e.g., observed data)
- **infer**: compute the posterior distribution of a model given the inputs

Multiple examples:

- Church, Anglican (lisp, clojure), 2008
- WebPPL (javascript), 2014
- Pyro/NumPyro (python), 2017/2019
- Gen (julia), 2018
- ProbZelus (Zelus), 2019
- ...

Probabilistic Programming Languages

General purpose programming languages extended with probabilistic constructs

- **sample**: draw a sample from a distribution
- **assume**, **factor**, **observe**: condition the model on inputs (e.g., observed data)
- **infer**: compute the posterior distribution of a model given the inputs

Multiple examples:

- Church, Anglican (lisp, clojure), 2008
- WebPPL (javascript), 2014
- Pyro/NumPyro (python), 2017/2019
- Gen (julia), 2018
- ProbZelus (Zelus), 2019
- ...

More and more, incorporating new ideas:

- New inference techniques, e.g., stochastic variational inference (SVI)
- Interaction with neural nets (deep probabilistic programming)

Build Your Owl PPL

BYO-PPL

Simplified Syntax

$x ::= \text{variables}$

$c ::= \text{constants}$

$d ::= \text{let } x = e \mid d \ d$

$p ::= x \mid (p, p)$

$e ::= c \mid x \mid (e, e) \mid \text{op}(e) \mid e \ e$

$\mid \text{if } e \text{ then } e \text{ else } e \mid \text{let } p = e \text{ in } e \mid \text{fun } p \rightarrow e$

$\mid \text{sample } e \mid \text{assume } e \mid \text{factor } e \mid \text{observe } e \ e \mid \text{infer } e$

A first-order functional programming language extended with probabilistic constructs

Outline

For a given inference algorithm, how to implement `sample`, `assume`, `factor`, `observe`, and `infer`?

I - Basic inference

- Rejection sampling
- Importance sampling

II - Continuation Passing Style models

III - Inference on CPS models

- Sample generation
- Importance sampling
- Particle filter

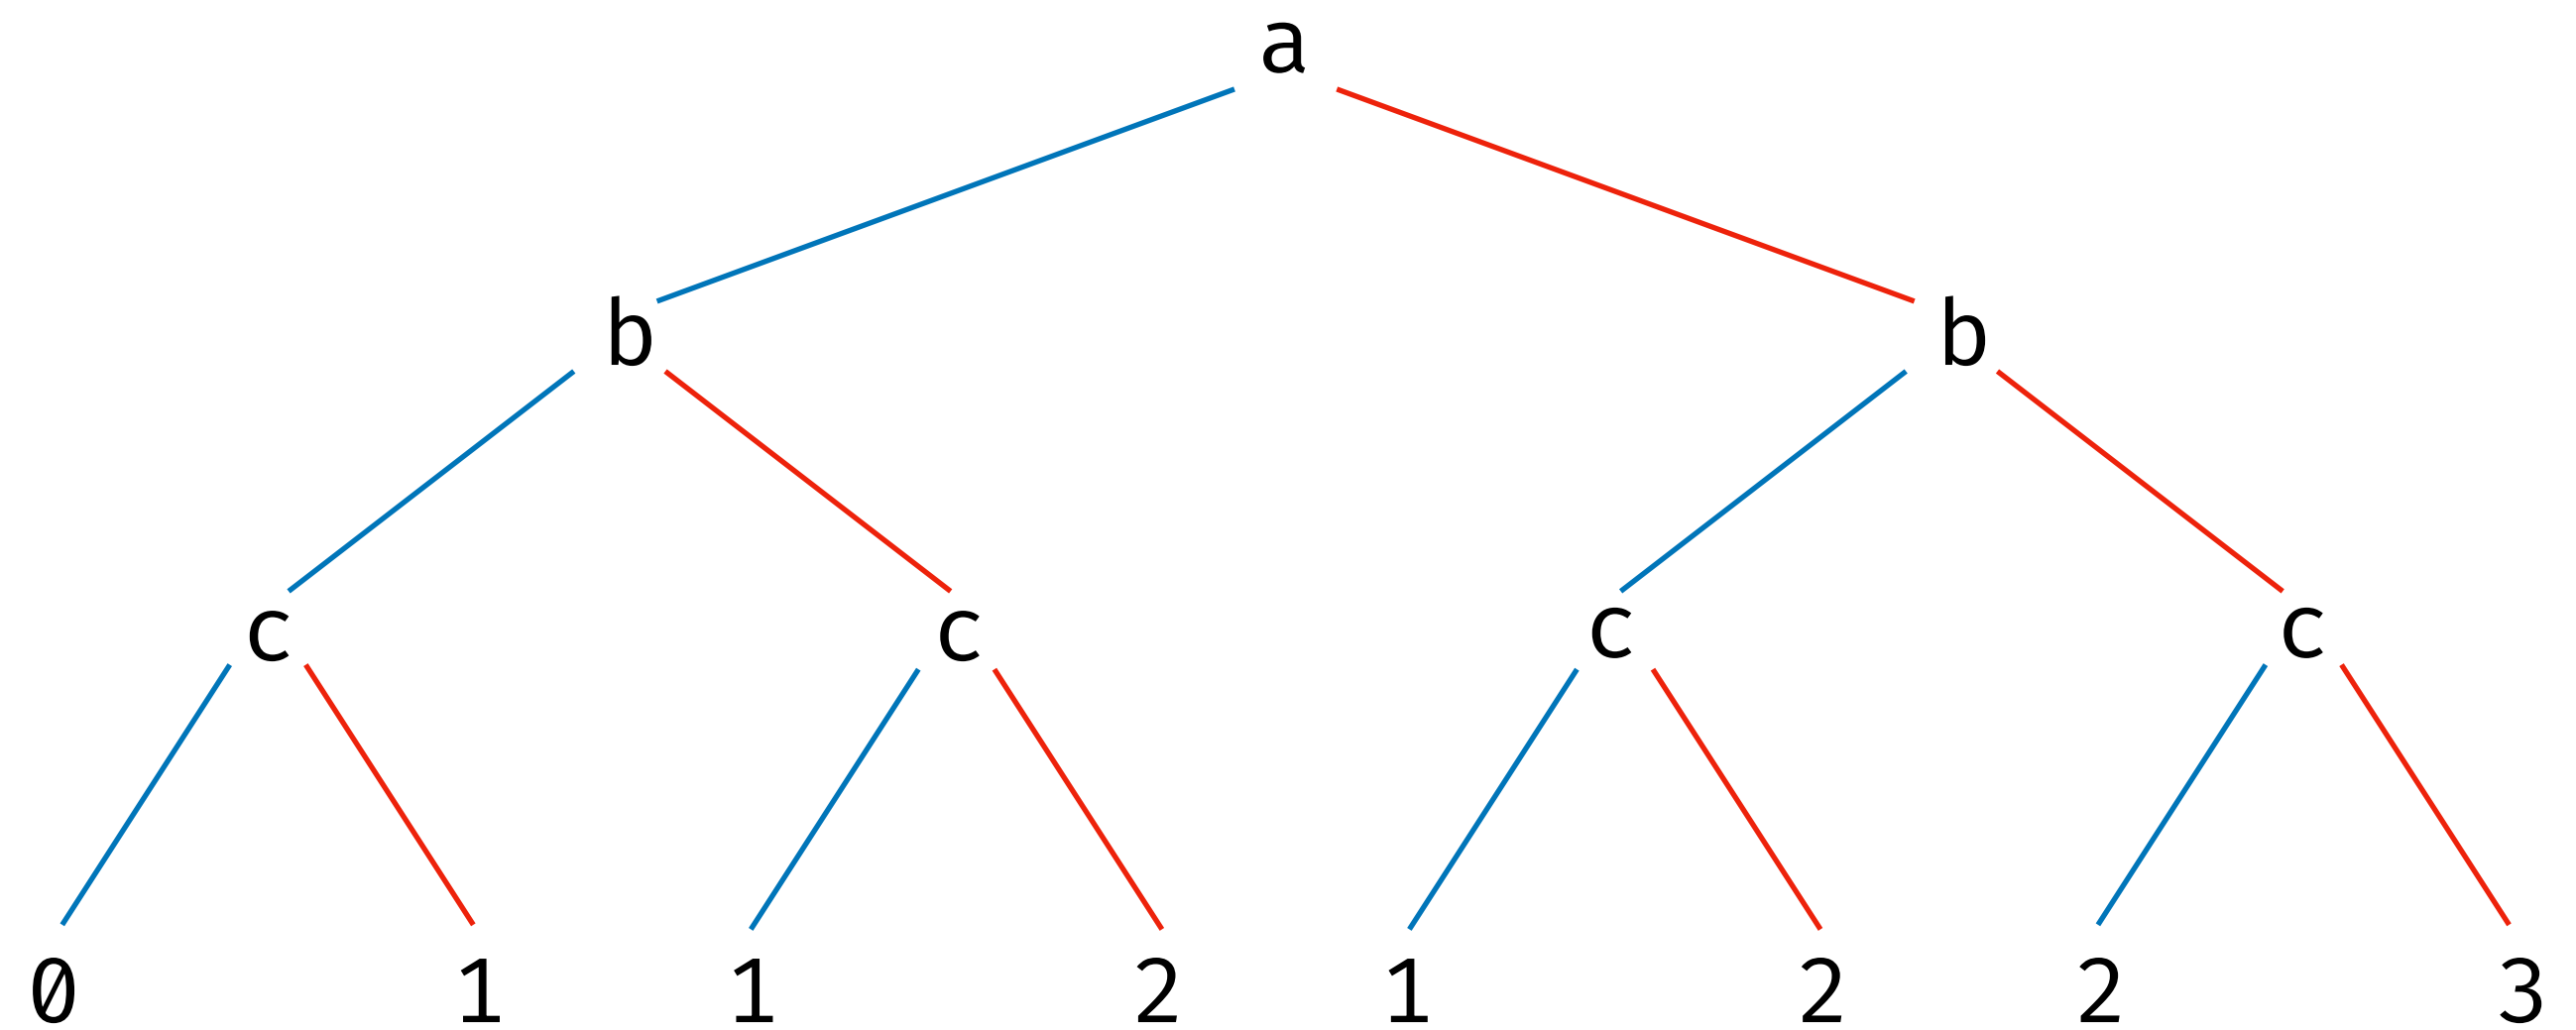
Warm-up : Rejection Sampling

BYO-PPL

Funny Bernoulli

funny_bernoulli.ml

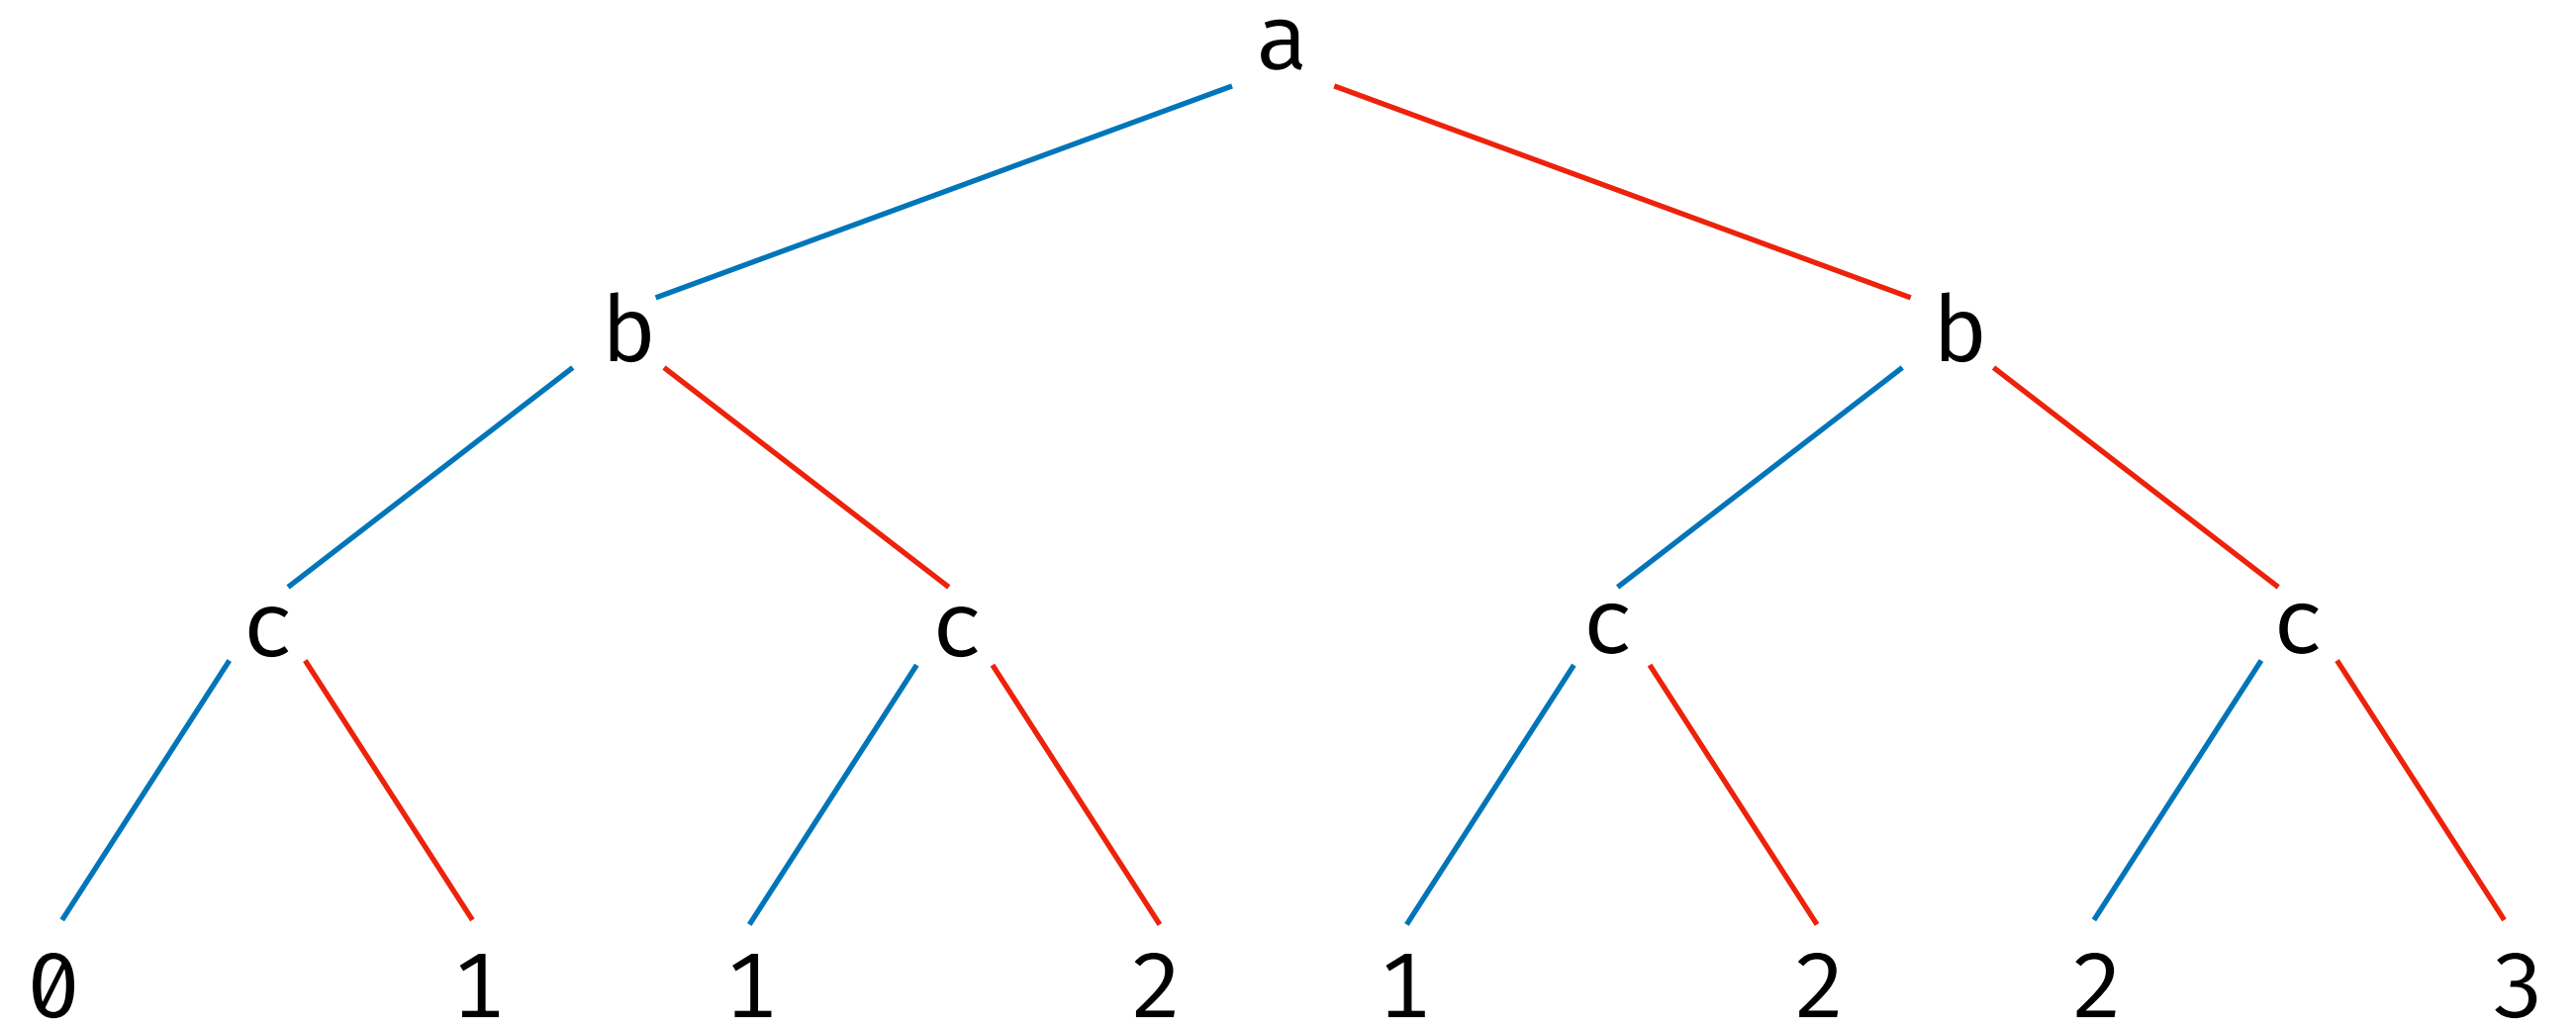
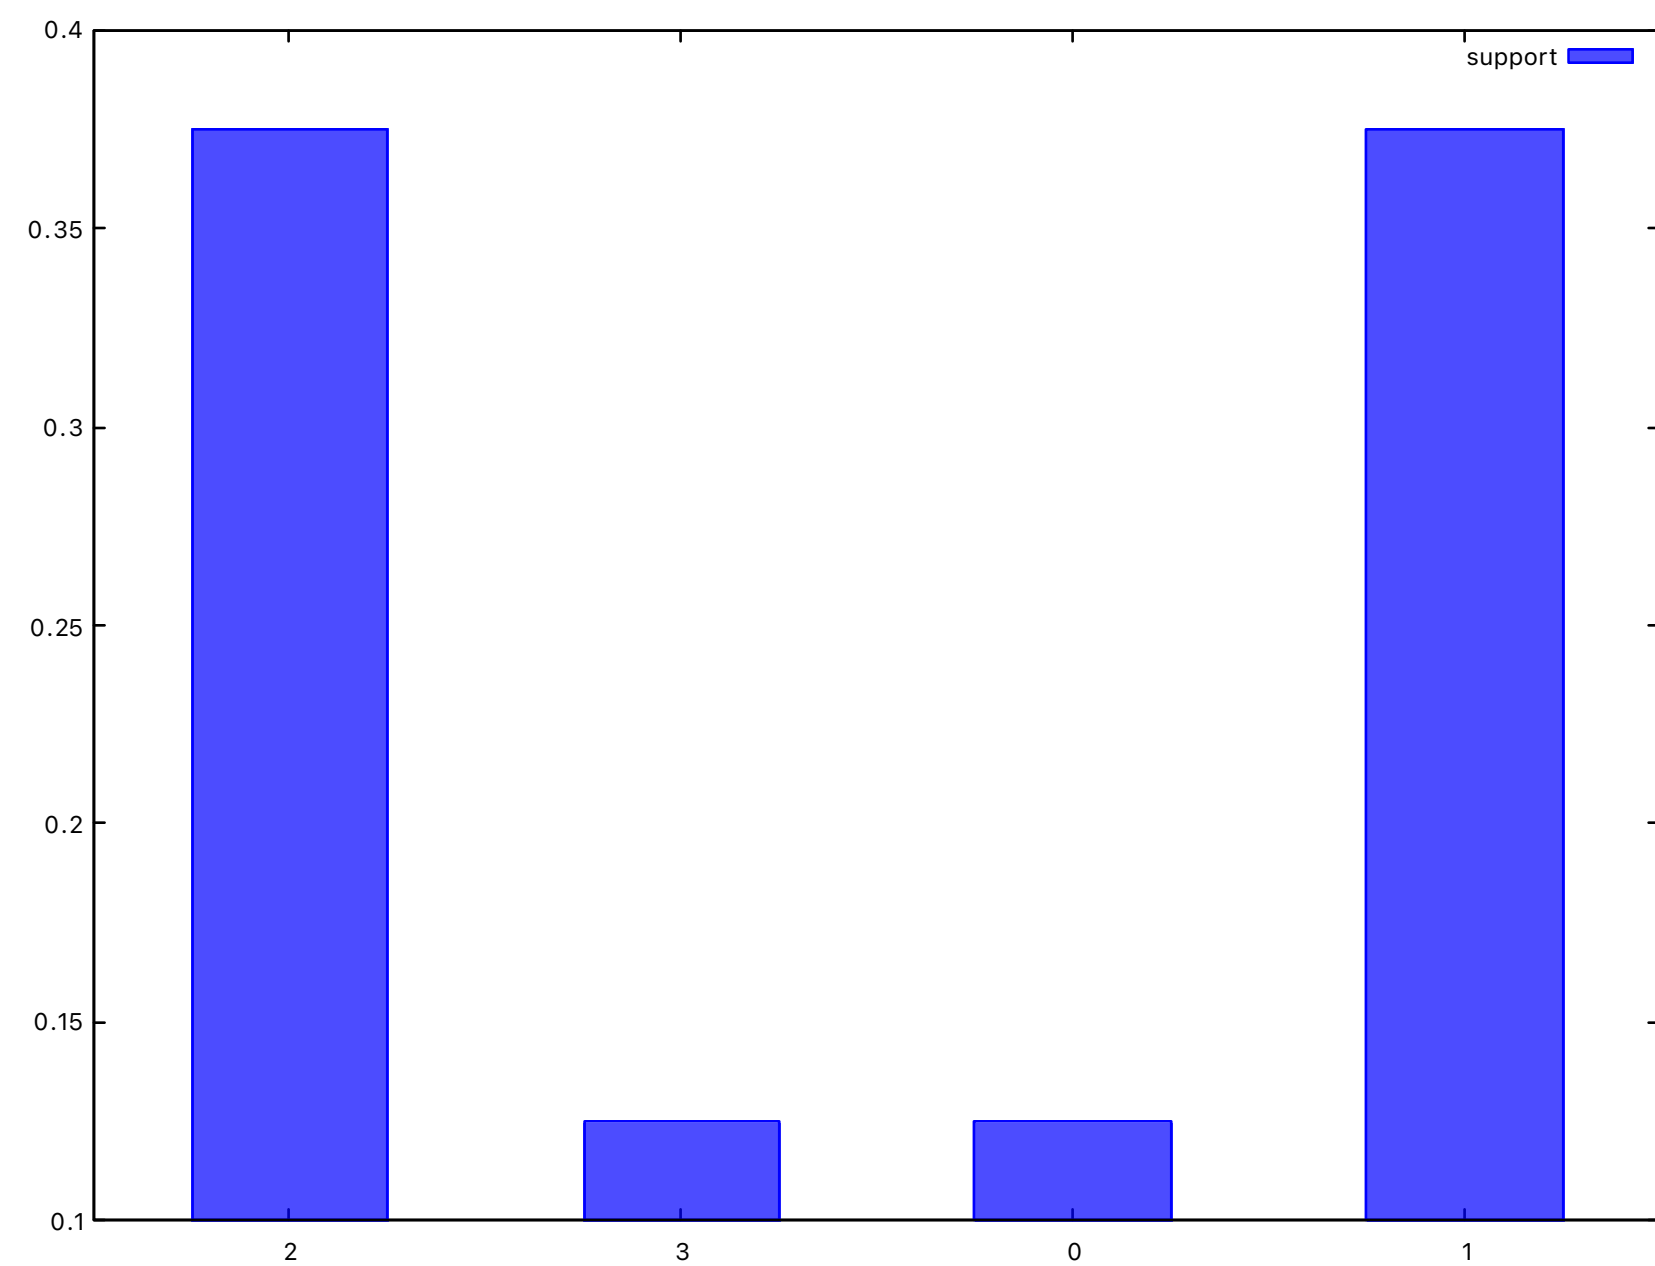
```
let funny_bernoulli () =  
  let a = sample (bernoulli ~p:0.5) in  
  let b = sample (bernoulli ~p:0.5) in  
  let c = sample (bernoulli ~p:0.5) in  
  a + b + c
```



Funny Bernoulli

funny_bernoulli.ml

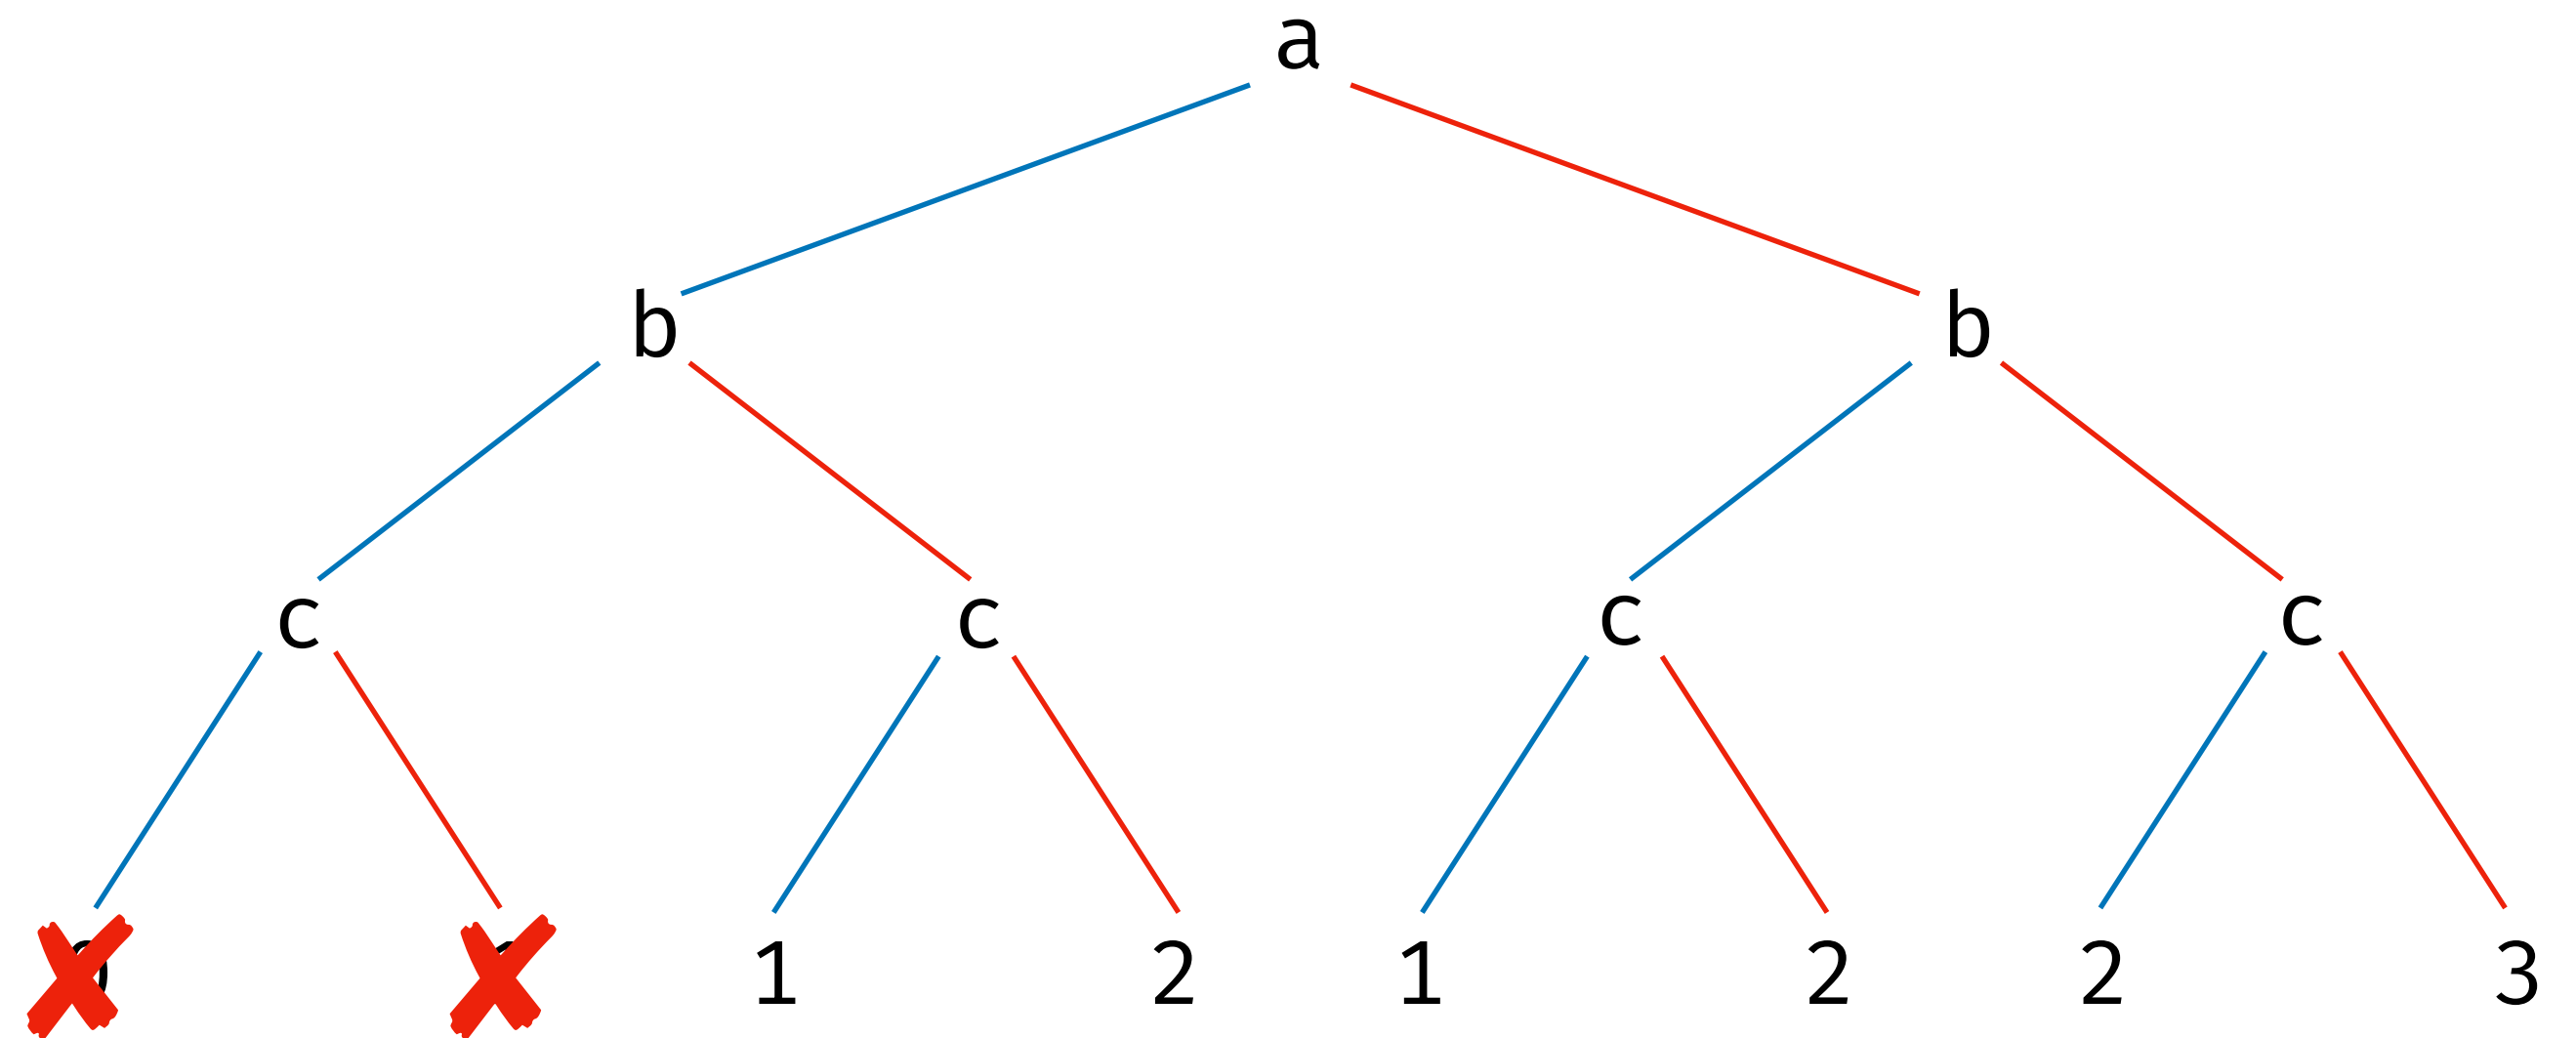
```
let funny_bernoulli () =  
  let a = sample (bernoulli ~p:0.5) in  
  let b = sample (bernoulli ~p:0.5) in  
  let c = sample (bernoulli ~p:0.5) in  
  a + b + c
```



Funny Bernoulli

funny_bernoulli.ml

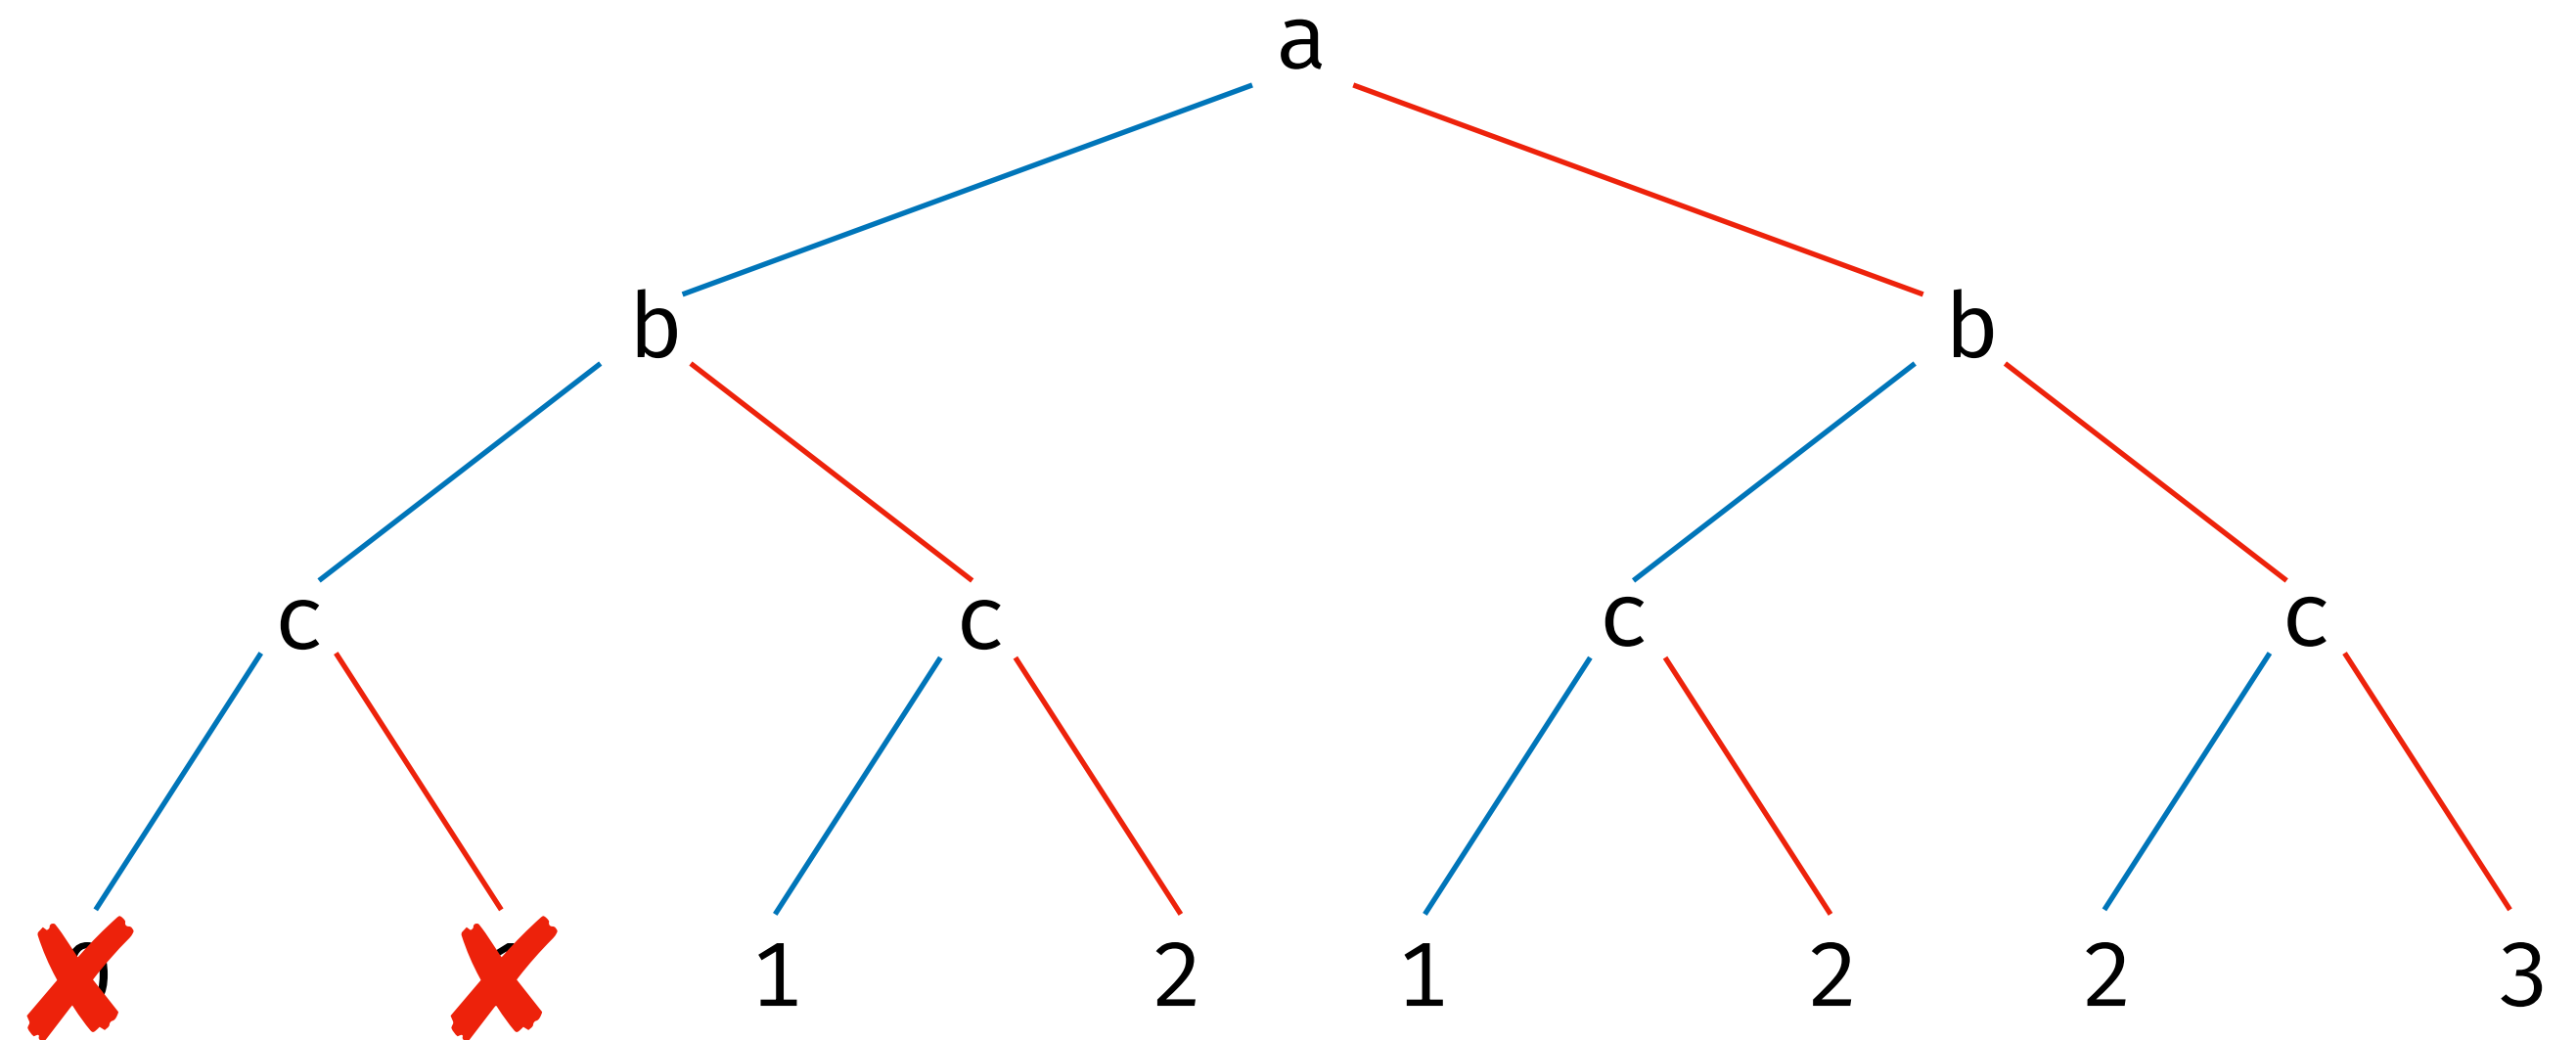
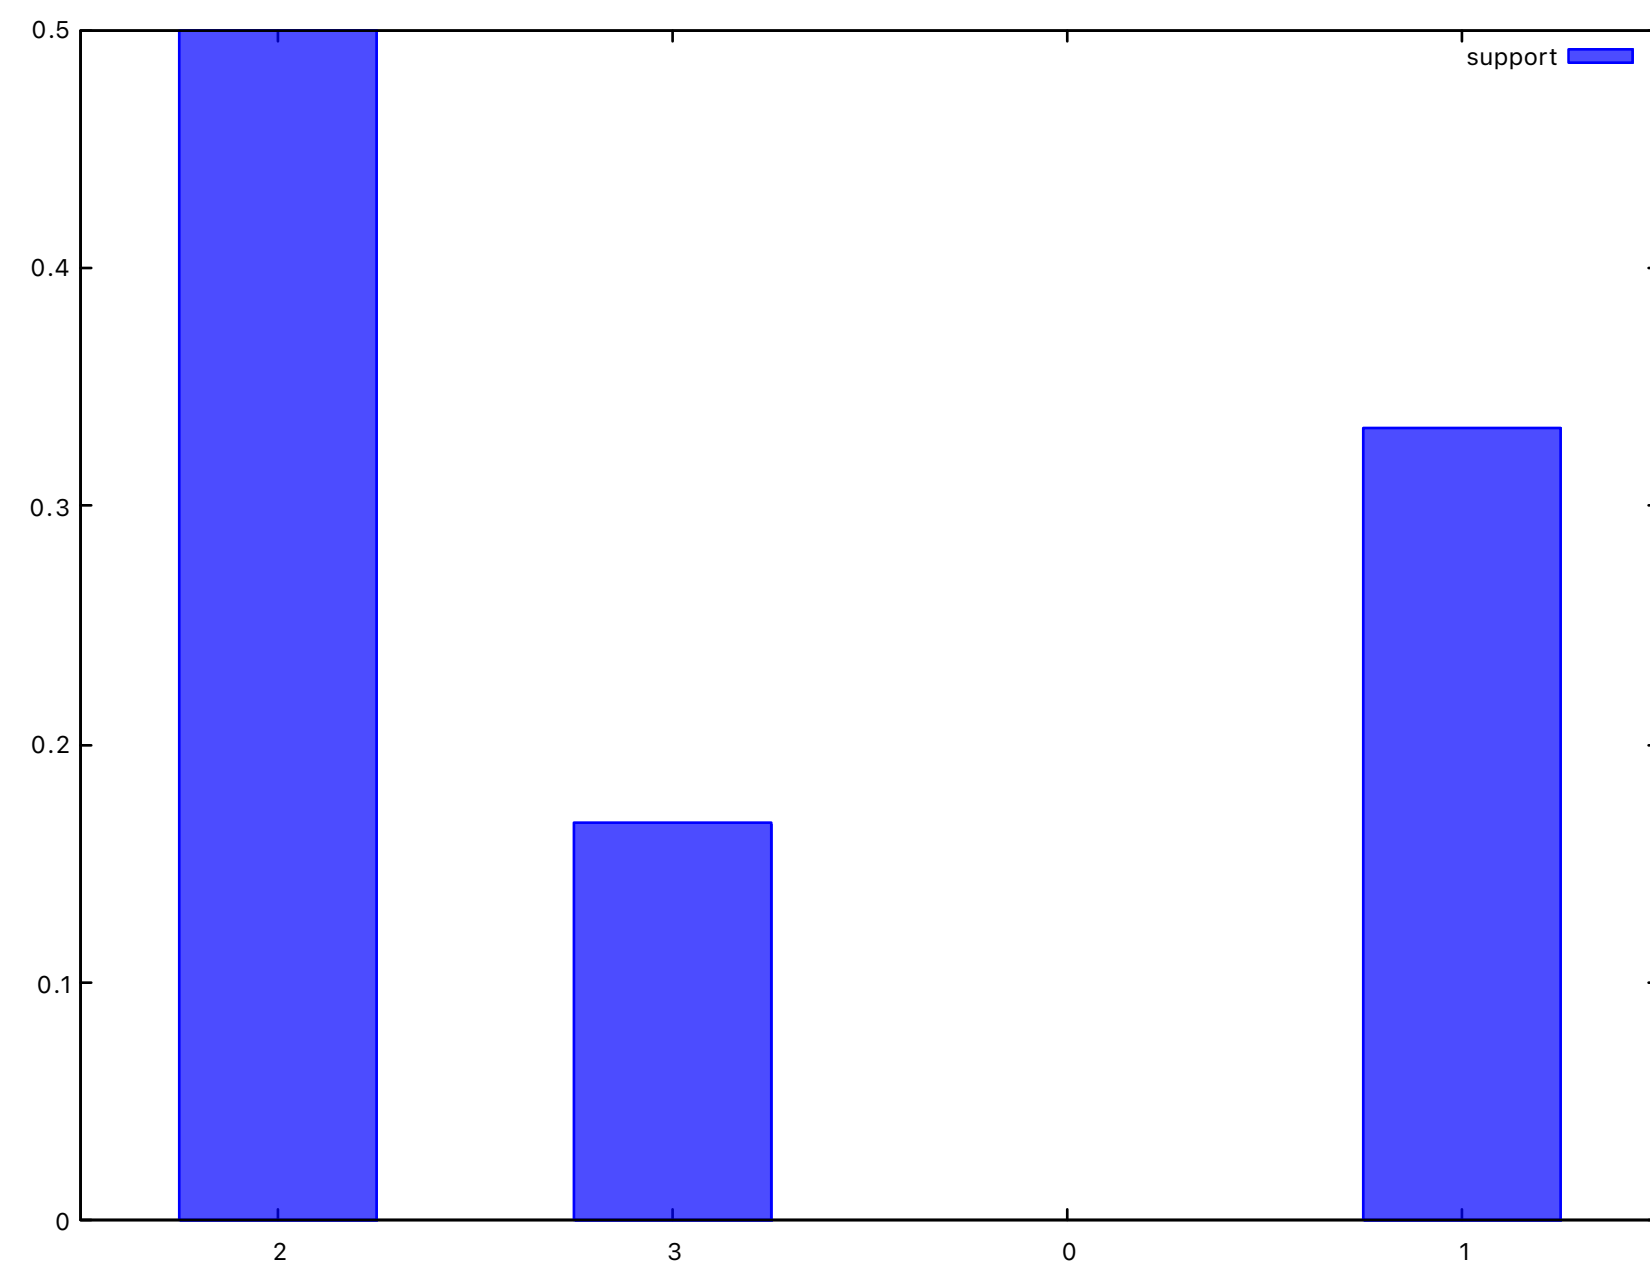
```
let funny_bernoulli () =  
  let a = sample (bernoulli ~p:0.5) in  
  let b = sample (bernoulli ~p:0.5) in  
  let c = sample (bernoulli ~p:0.5) in  
  let () = assume (a = 1 || b = 1) in  
  a + b + c
```



Funny Bernoulli

funny_bernoulli.ml

```
let funny_bernoulli () =  
  let a = sample (bernoulli ~p:0.5) in  
  let b = sample (bernoulli ~p:0.5) in  
  let c = sample (bernoulli ~p:0.5) in  
  let () = assume (a = 1 || b = 1) in  
  a + b + c
```



Rejection Sampling

basic.ml

```
module Rejection_sampling : sig
  val sample : 'a Distribution.t → 'a
  val assume : bool → unit
  val infer : ?n:int → ('a → 'b) → 'a → 'b Distribution.t
end = struct ... end
```

Inference algorithm

- Run the model to get a sample
- **sample** : draw a value from a distribution
- **assume** : accept / reject a sample
- If a sample is rejected, re-run the model to get another sample

Rejection Sampling

basic.ml

```
module Rejection_sampling = struct

  let sample d = assert false
  let assume p = assert false

  let infer ?(n = 1000) model obs = assert false
end
```

Rejection Sampling

basic.ml

```
module Rejection_sampling = struct
  exception Reject

  let sample d = Distribution.draw d
  let assume p = if not p then raise Reject

  let infer ?(n = 1000) model obs =
    let rec exec i = try model obs with Reject → exec i in
    let values = Array.init n exec in
    Distribution.uniform_support ~values
end
```


The type `prob` trick

```
module Rejection_sampling : sig
  type prob
  val sample : prob → 'a Distribution.t → 'a
  val assume : prob → bool → unit
  val infer : ?n:int → (prob → 'a → 'b) → 'a → 'b Distribution.t
end = struct ... end
```

Forbid the use of probabilistic construct outside a model

- Define a simple abstract type `prob`
- Probabilistic constructs and models all require an argument of type `prob`
- Such a value can only be build by `infer`

Rejection Sampling

basic.ml

```
module Rejection_sampling = struct
  type prob = Prob

  exception Reject

  let sample _prob d = Distribution.draw d
  let assume _prob p = if not p then raise Reject

  let infer ?(n = 1000) model obs =
    let rec exec i = try model Prob obs with Reject → exec i in
    let values = Array.init n exec in
    Distribution.uniform_support ~values
end
```

Funny Bernoulli

funny_bernoulli.ml

```
open Byoppl
open Distribution
open Basic.Rejection_sampling

let funny_bernoulli prob () =
  let a = sample prob (bernoulli ~p:0.5) in
  let b = sample prob (bernoulli ~p:0.5) in
  let c = sample prob (bernoulli ~p:0.5) in
  let () = assume prob (a = 1 || b = 1) in
  a + b + c

let _ =
  let dist = infer funny_bernoulli () in
  let { values; probs; _ } = get_support ~shrink:true dist in
  Array.iteri (fun i x → Format.printf "%d %f@." x probs.(i)) values
```

› dune exec ./examples/funny_bernoulli.exe

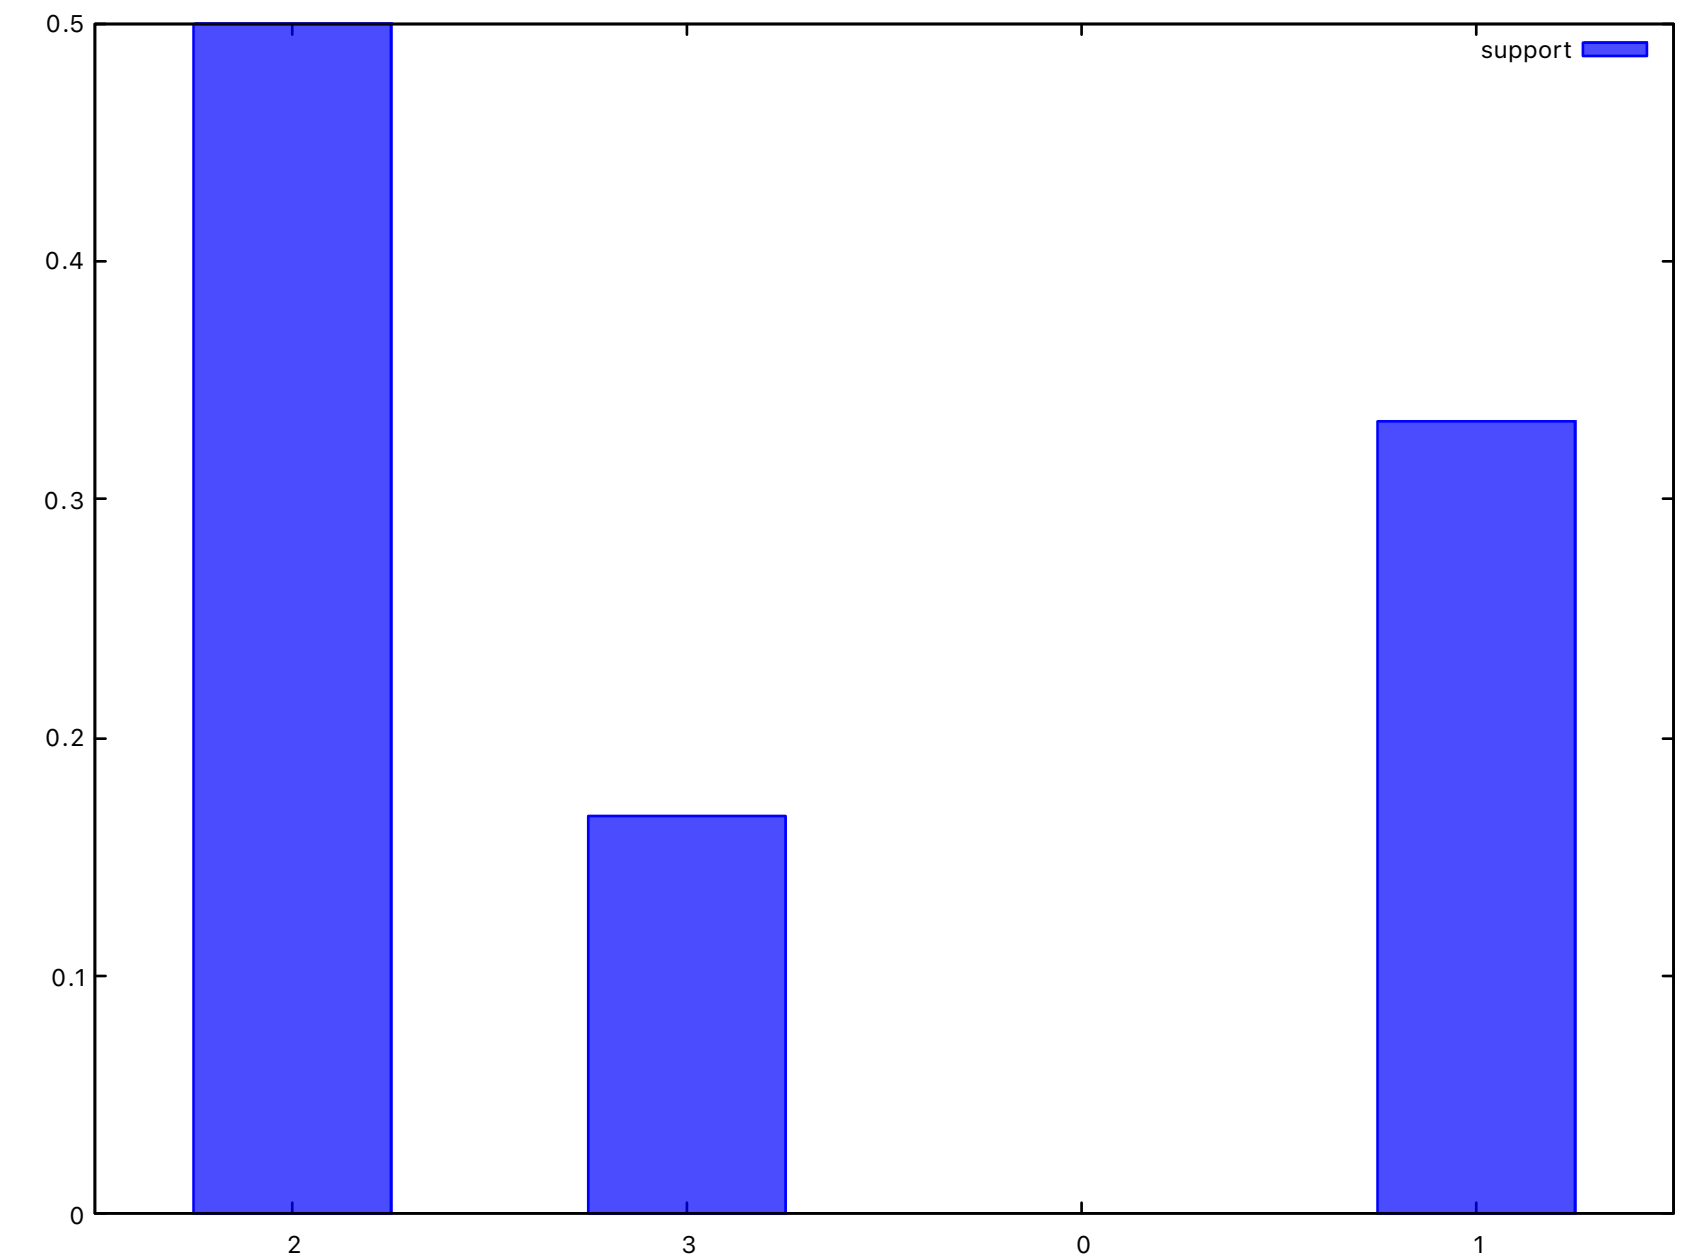
Funny Bernoulli

funny_bernoulli.ml

```
open Byoppl
open Distribution
open Basic.Rejection_sampling
```

```
let funny_bernoulli prob () =
  let a = sample prob (bernoulli ~p:0.5) in
  let b = sample prob (bernoulli ~p:0.5) in
  let c = sample prob (bernoulli ~p:0.5) in
  let () = assume prob (a = 1 || b = 1) in
  a + b + c
```

```
let _ =
  let dist = infer funny_bernoulli () in
  let { values; probs; _ } = get_support ~shrink:true dist in
  Array.iteri (fun i x → Format.printf "%d %f@" x probs.(i)) values
```



```
> dune exec ./examples/funny_bernoulli.exe
```

Importance Sampling


BYO-PPL

Laplace and Gender Bias

laplace.ml

```
open Basic.Rejection_sampling
```

```
let laplace prob () =  
  let p = sample prob (uniform ~a:0. ~b:1.) in  
  let g = sample prob (binomial ~p ~n:493_472) in  
  let () = assume prob (g = 241_945) in  
  p
```



```
let () = observe prob  
  (binomial ~p ~n:493_472) 241_945
```

```
let _ =  
  let dist = infer ~n:1000 laplace () in  
  let m, s = Distribution.stats dist in  
  Format.printf "Gender bias, mean:%f std:%f@." m s
```


```
> dune exec ./examples/laplace.exe
```

Laplace and Gender Bias

laplace.ml

```
open Basic.Rejection_sampling
```

```
let laplace prob () =  
  let p = sample prob (uniform ~a:0. ~b:1.) in  
  let g = sample prob (binomial ~p ~n:493_472) in  
  let () = assume prob (g = 241_945) in  
  p
```



```
let () = observe prob  
  (binomial ~p ~n:493_472) 241_945
```

```
let _ =  
  let dist = infer ~n:1000 laplace () in  
  let m, s = Distribution.stats dist in  
  Format.printf "Gender bias, mean:%f std:%f@." m s
```

```
> dune exec ./examples/laplace.exe
```

Never terminate!

Coin

coin.ml

```
open Basic.Rejection_sampling
```

```
let coin prob x =  
  let z = sample prob (uniform ~a:0. ~b:1.) in  
  let () = List.iter (observe prob (bernoulli ~p:z)) x in  
  z
```

```
let _ =  
  let dist = infer coin [1; 1; 0; 0; 0; 0; 0; 0; 0; 0] in  
  let m, s = Distribution.stats dist in  
  Format.printf "Coin bias, mean:%f std:%f@." m s
```

```
> dune exec ./examples/coin.exe
```

```
Coin bias, mean:0.246161, std:0.119687
```


Coin

coin.ml

```
open Basic.Rejection_sampling
```

```
let coin prob x =  
  let z = sample prob (uniform ~a:0. ~b:1.) in  
  let () = List.iter (observe prob (bernoulli ~p:z)) x in  
  z
```

```
let _ =  
  let dist = infer coin [1; 1; 0; 0; 0; 0; 0; 0; 0; 0] in  
  let m, s = Distribution.stats dist in  
  Format.printf "Coin bias, mean:%f std:%f@." m s
```

```
> dune exec ./examples/coin.exe
```

```
Coin bias, mean:0.246161, std:0.119687
```

Slow!

Importance Sampling

basic.ml

```
module Importance_sampling : sig
  type prob
  val sample : prob → 'a Distribution.t → 'a
  val factor : prob → float → unit
  val infer : ?n:int → (prob → 'a → 'b) → 'a → 'b Distribution.t
end = struct ... end
```

Inference algorithm

- Run a set of n independent executions
- **sample**: draw a sample from a distribution
- **factor**: associate a score to the current execution
- Gather output values and score to approximate the posterior distribution

Likelihood weighting

- **observe** $d\ x := \text{factor} (\text{logpdf } d\ x)$
- **assume** $p := \text{factor} (\text{if } p \text{ then } 0. \text{ else } -.infinity)$

Importance Sampling

basic.ml

```
module Importance_sampling = struct
  type prob = ...

  let sample prob d = assert false
  let factor prob s = assert false
  let observe prob d x = factor prob (Distribution.logpdf d x)
  let assume prob p = factor prob (if p then 0. else -. infinity)

  let infer ?(n = 1000) model obs = assert false
end
```

Importance Sampling

basic.ml

```
module Importance_sampling = struct
  type prob = { id : int; scores : float array }

  let sample _prob d = Distribution.draw d
  let factor prob s = prob.scores.(prob.id) ← prob.scores.(prob.id) +. s
  let observe prob d x = factor prob (Distribution.logpdf d x)
  let assume prob p = factor prob (if p then 0. else -. infinity)

  let infer ?(n = 1000) model obs =
    let scores = Array.make n 0. in
    let values = Array.mapi (fun i _ → model { id = i; scores } obs) scores in
    Distribution.support ~values ~logits:scores
end
```

Coin

coin.ml

```
open Basic.Importance_sampling
```

```
let coin prob x =  
  let z = sample prob (uniform ~a:0. ~b:1.) in  
  let () = List.iter (observe prob (bernoulli ~p:z)) x in  
  z
```

```
let _ =  
  let dist = infer coin [1; 1; 0; 0; 0; 0; 0; 0; 0; 0] in  
  let m, s = Distribution.stats dist in  
  Format.printf "Coin bias, mean:%f, std:%f@." m s
```

```
> dune exec ./examples/coin.exe
```

```
Coin bias, mean:0.247876, std:0.118921  
Beta(2+1, 8+1), mean:0.250000, std:0.120096
```

Particle Filter

BYO-PPL

HMM: Hidden Markov Model

Track the position of an agent from noisy observations

- The current position should not be too far from the previous position
- The observations should not be too far from the current position

Probabilistic model: $\forall t \in \mathbb{N}$.

- $x_t \sim \mathcal{N}(x_{t-1}, \text{speed})$
- $y_t \sim \mathcal{N}(x_t, \text{noise})$

HMM: Hidden Markov Model

hmm.ml

```
open Basic.Importance_sampling

let hmm prob data =
  let rec gen states data =
    match (states, data) with
    | [], y :: data → gen [ y ] data
    | states, [] → states
    | pre_x :: _, y :: data →
      let x = sample prob (gaussian ~mu:pre_x ~sigma:1.0) in
      let () = observe prob (gaussian ~mu:x ~sigma:1.0) y in
      gen (x :: states) data
  in
  gen [] data

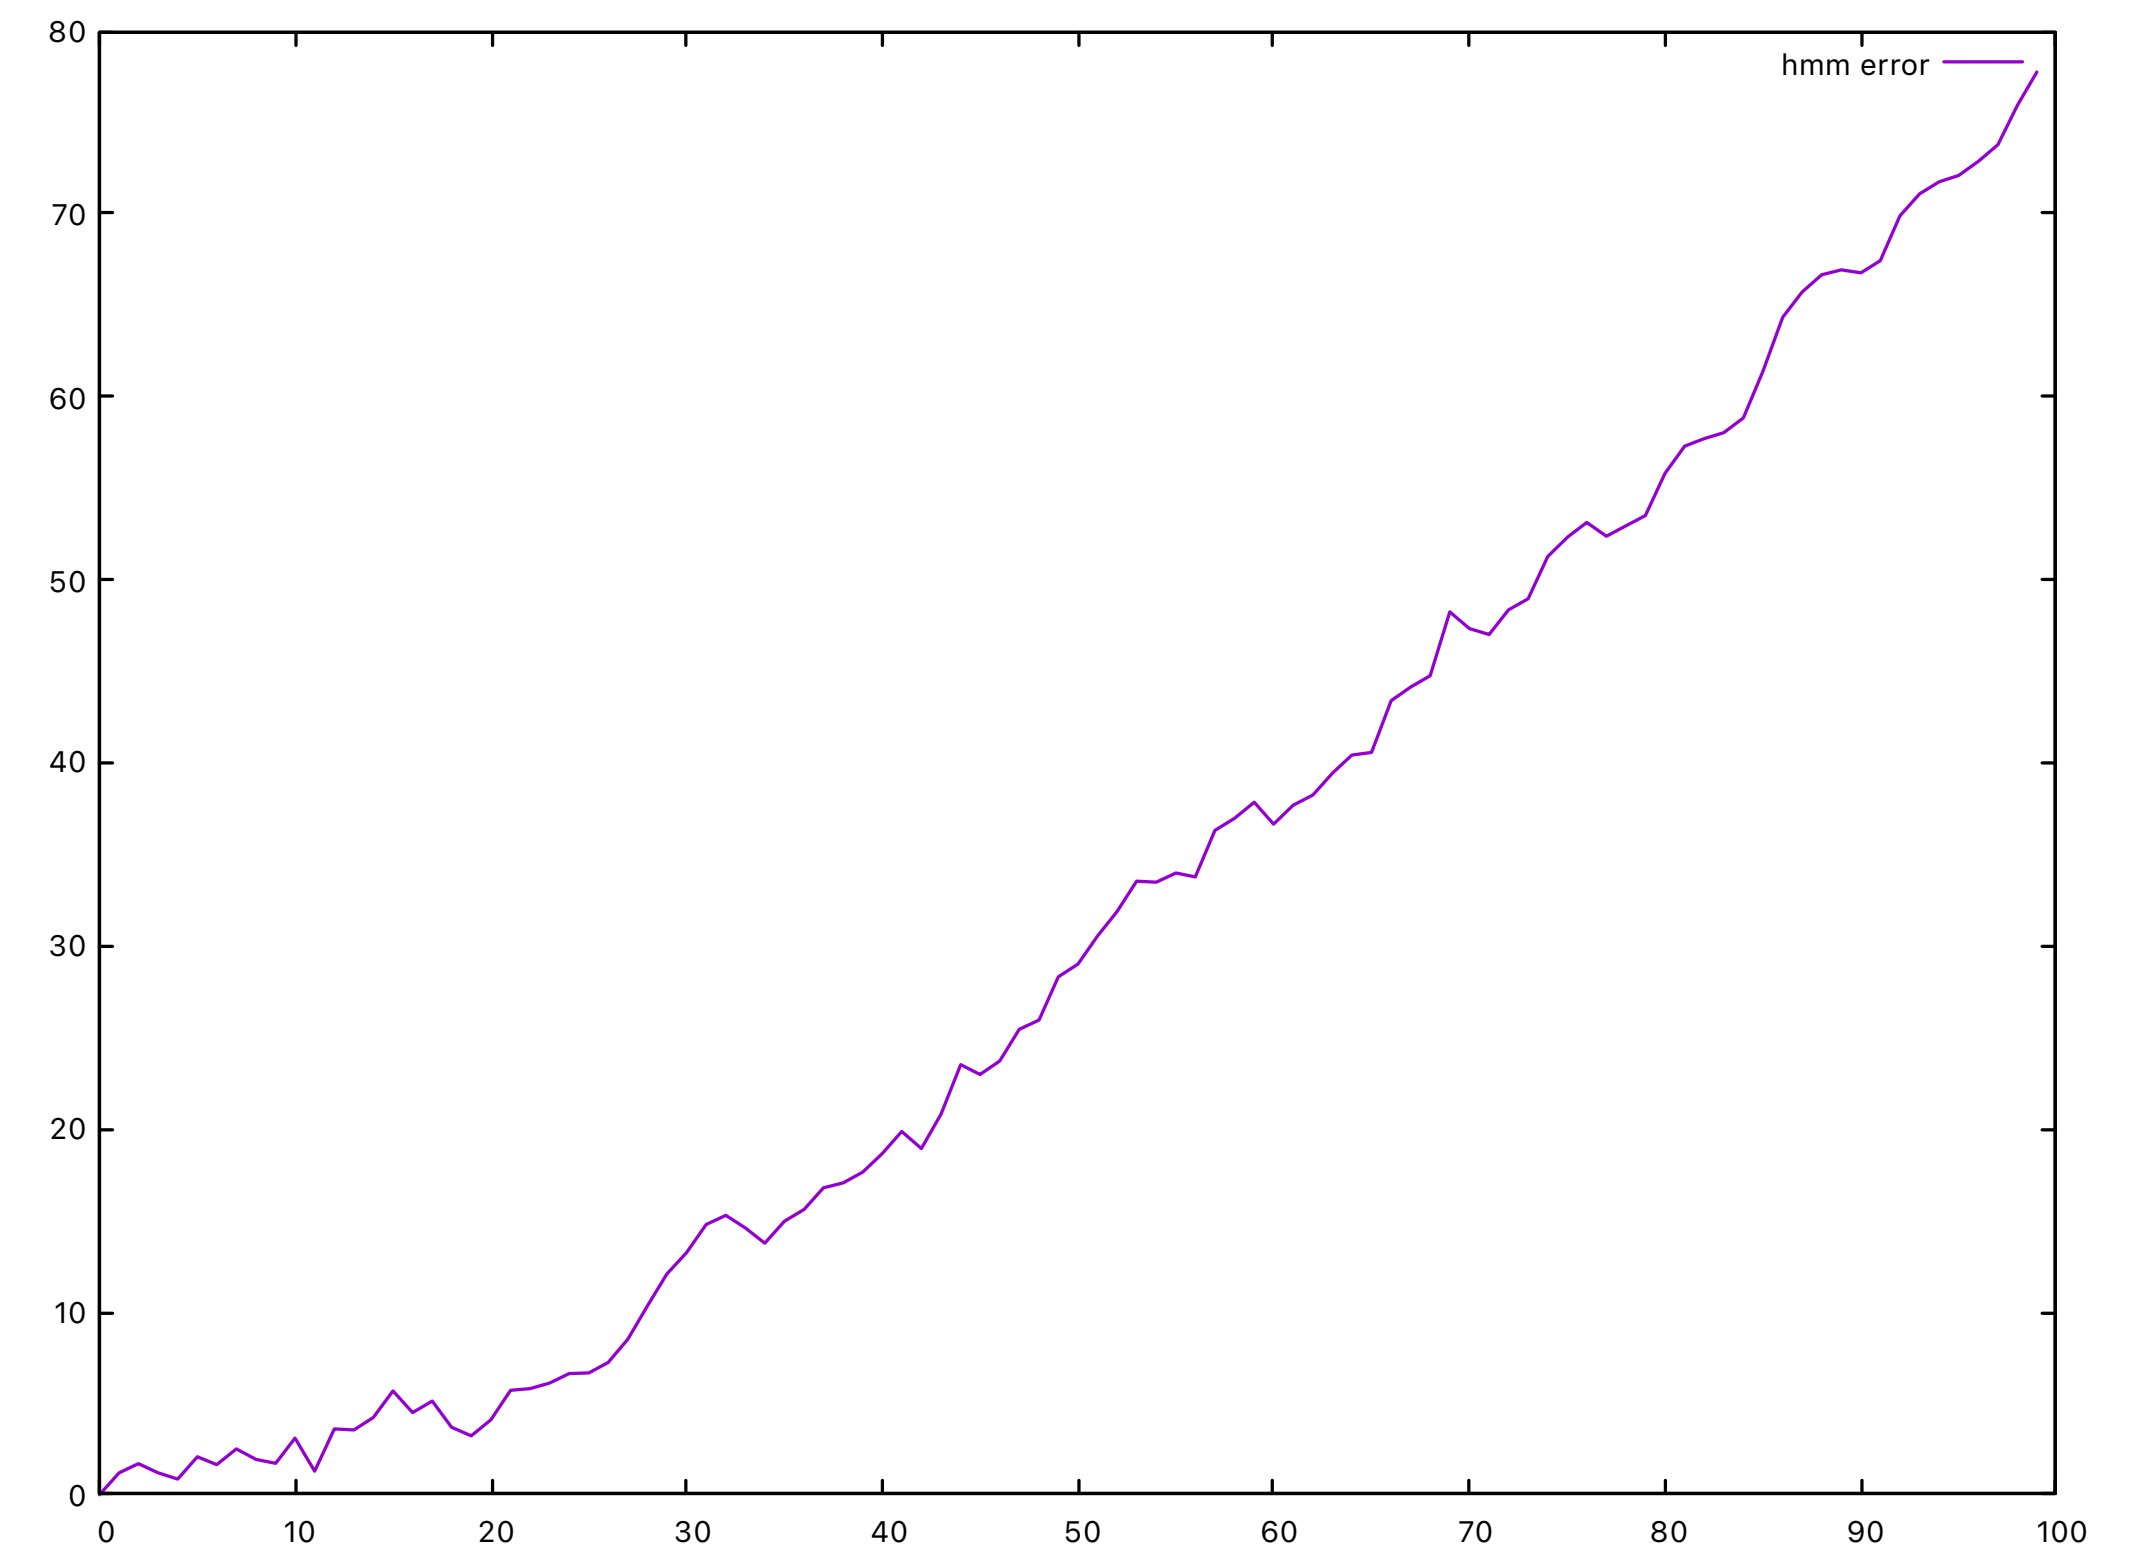
let _ =
  let data = Owl.Arr.linspace 0. 20. 20 ▷ Owl.Arr.to_array ▷ Array.to_list in
  let dist = Distribution.split_list (infer hmm data) in
  let m_x = List.rev (List.map Distribution.mean dist) in
  List.iter2 (Format.printf "%f >> %f@.") data m_x
```


HMM: Hidden Markov Model

```
› dune exec ./examples/hmm.exe
```

```
0.000000 >> 0.000000
1.052632 >> 0.278989
2.105263 >> 2.923428
3.157895 >> 2.812035
4.210526 >> 2.328341
5.263158 >> 1.742109
6.315789 >> 2.518105
7.368421 >> 3.958375
8.421053 >> 5.946233
9.473684 >> 7.329554
10.526316 >> 9.293653
11.578947 >> 10.181831
12.631579 >> 8.549409
13.684211 >> 9.323073
14.736842 >> 9.280692
15.789474 >> 9.352218
```

...



HMM: Importance Sampling

Problem:

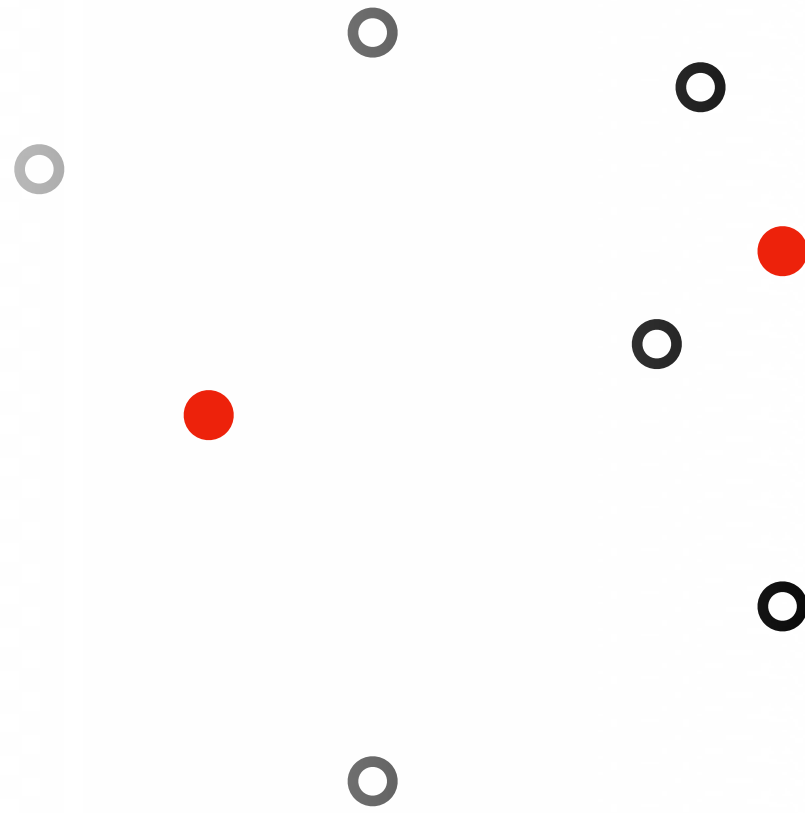
- Each particle does a random walk and compute a score
- No re-calibration based on observations
- Score decreases at each factor statement



HMM: Importance Sampling

Problem:

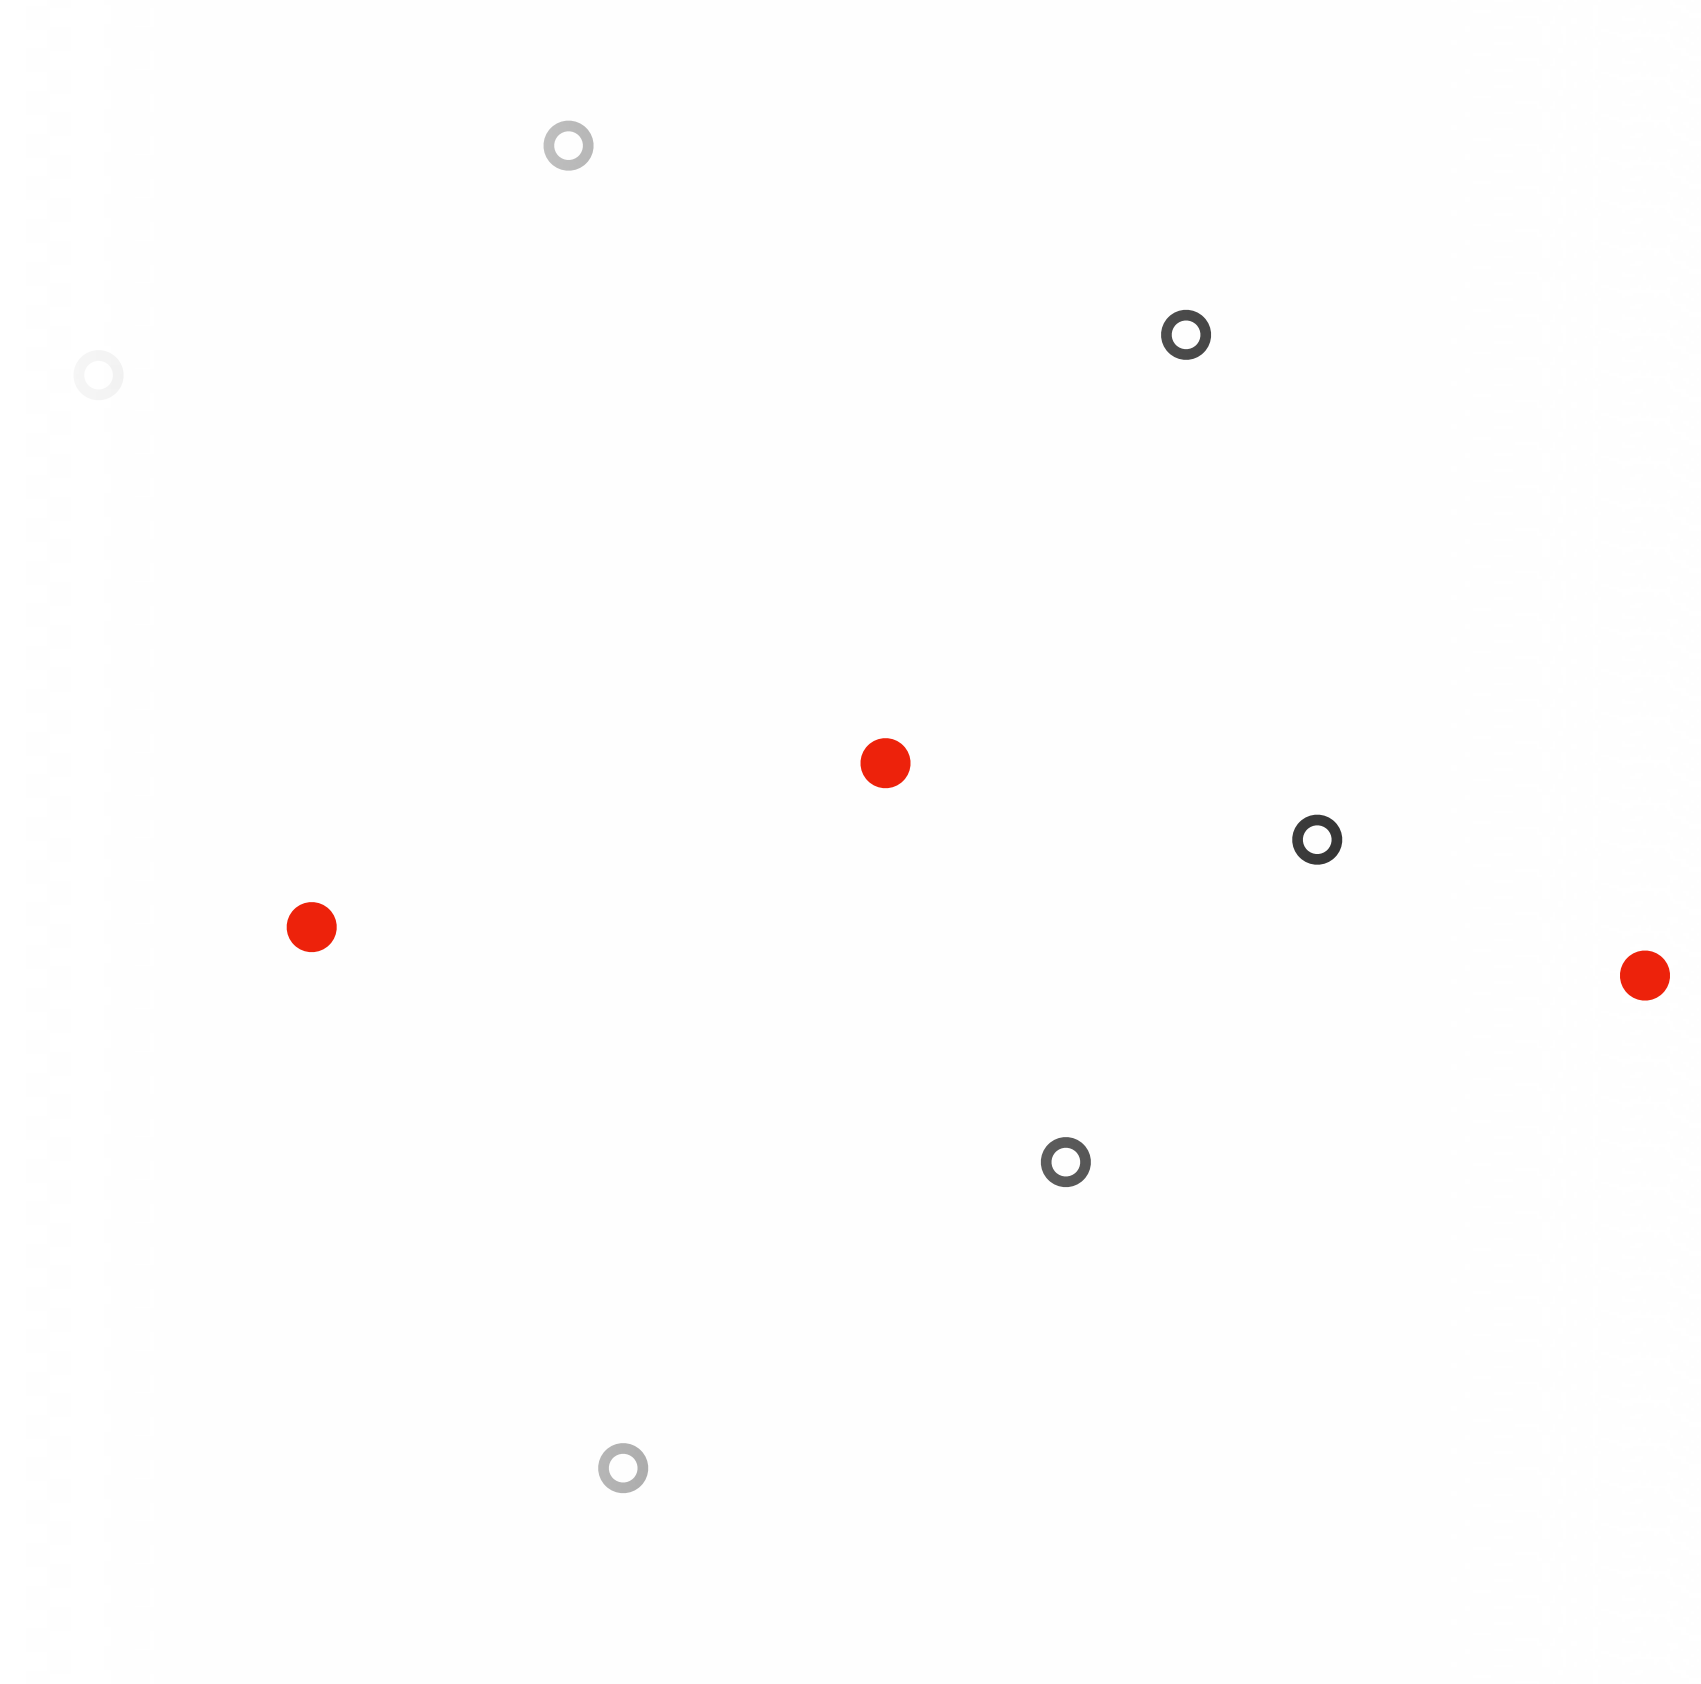
- Each particle does a random walk and compute a score
- No re-calibration based on observations
- Score decreases at each factor statement



HMM: Importance Sampling

Problem:

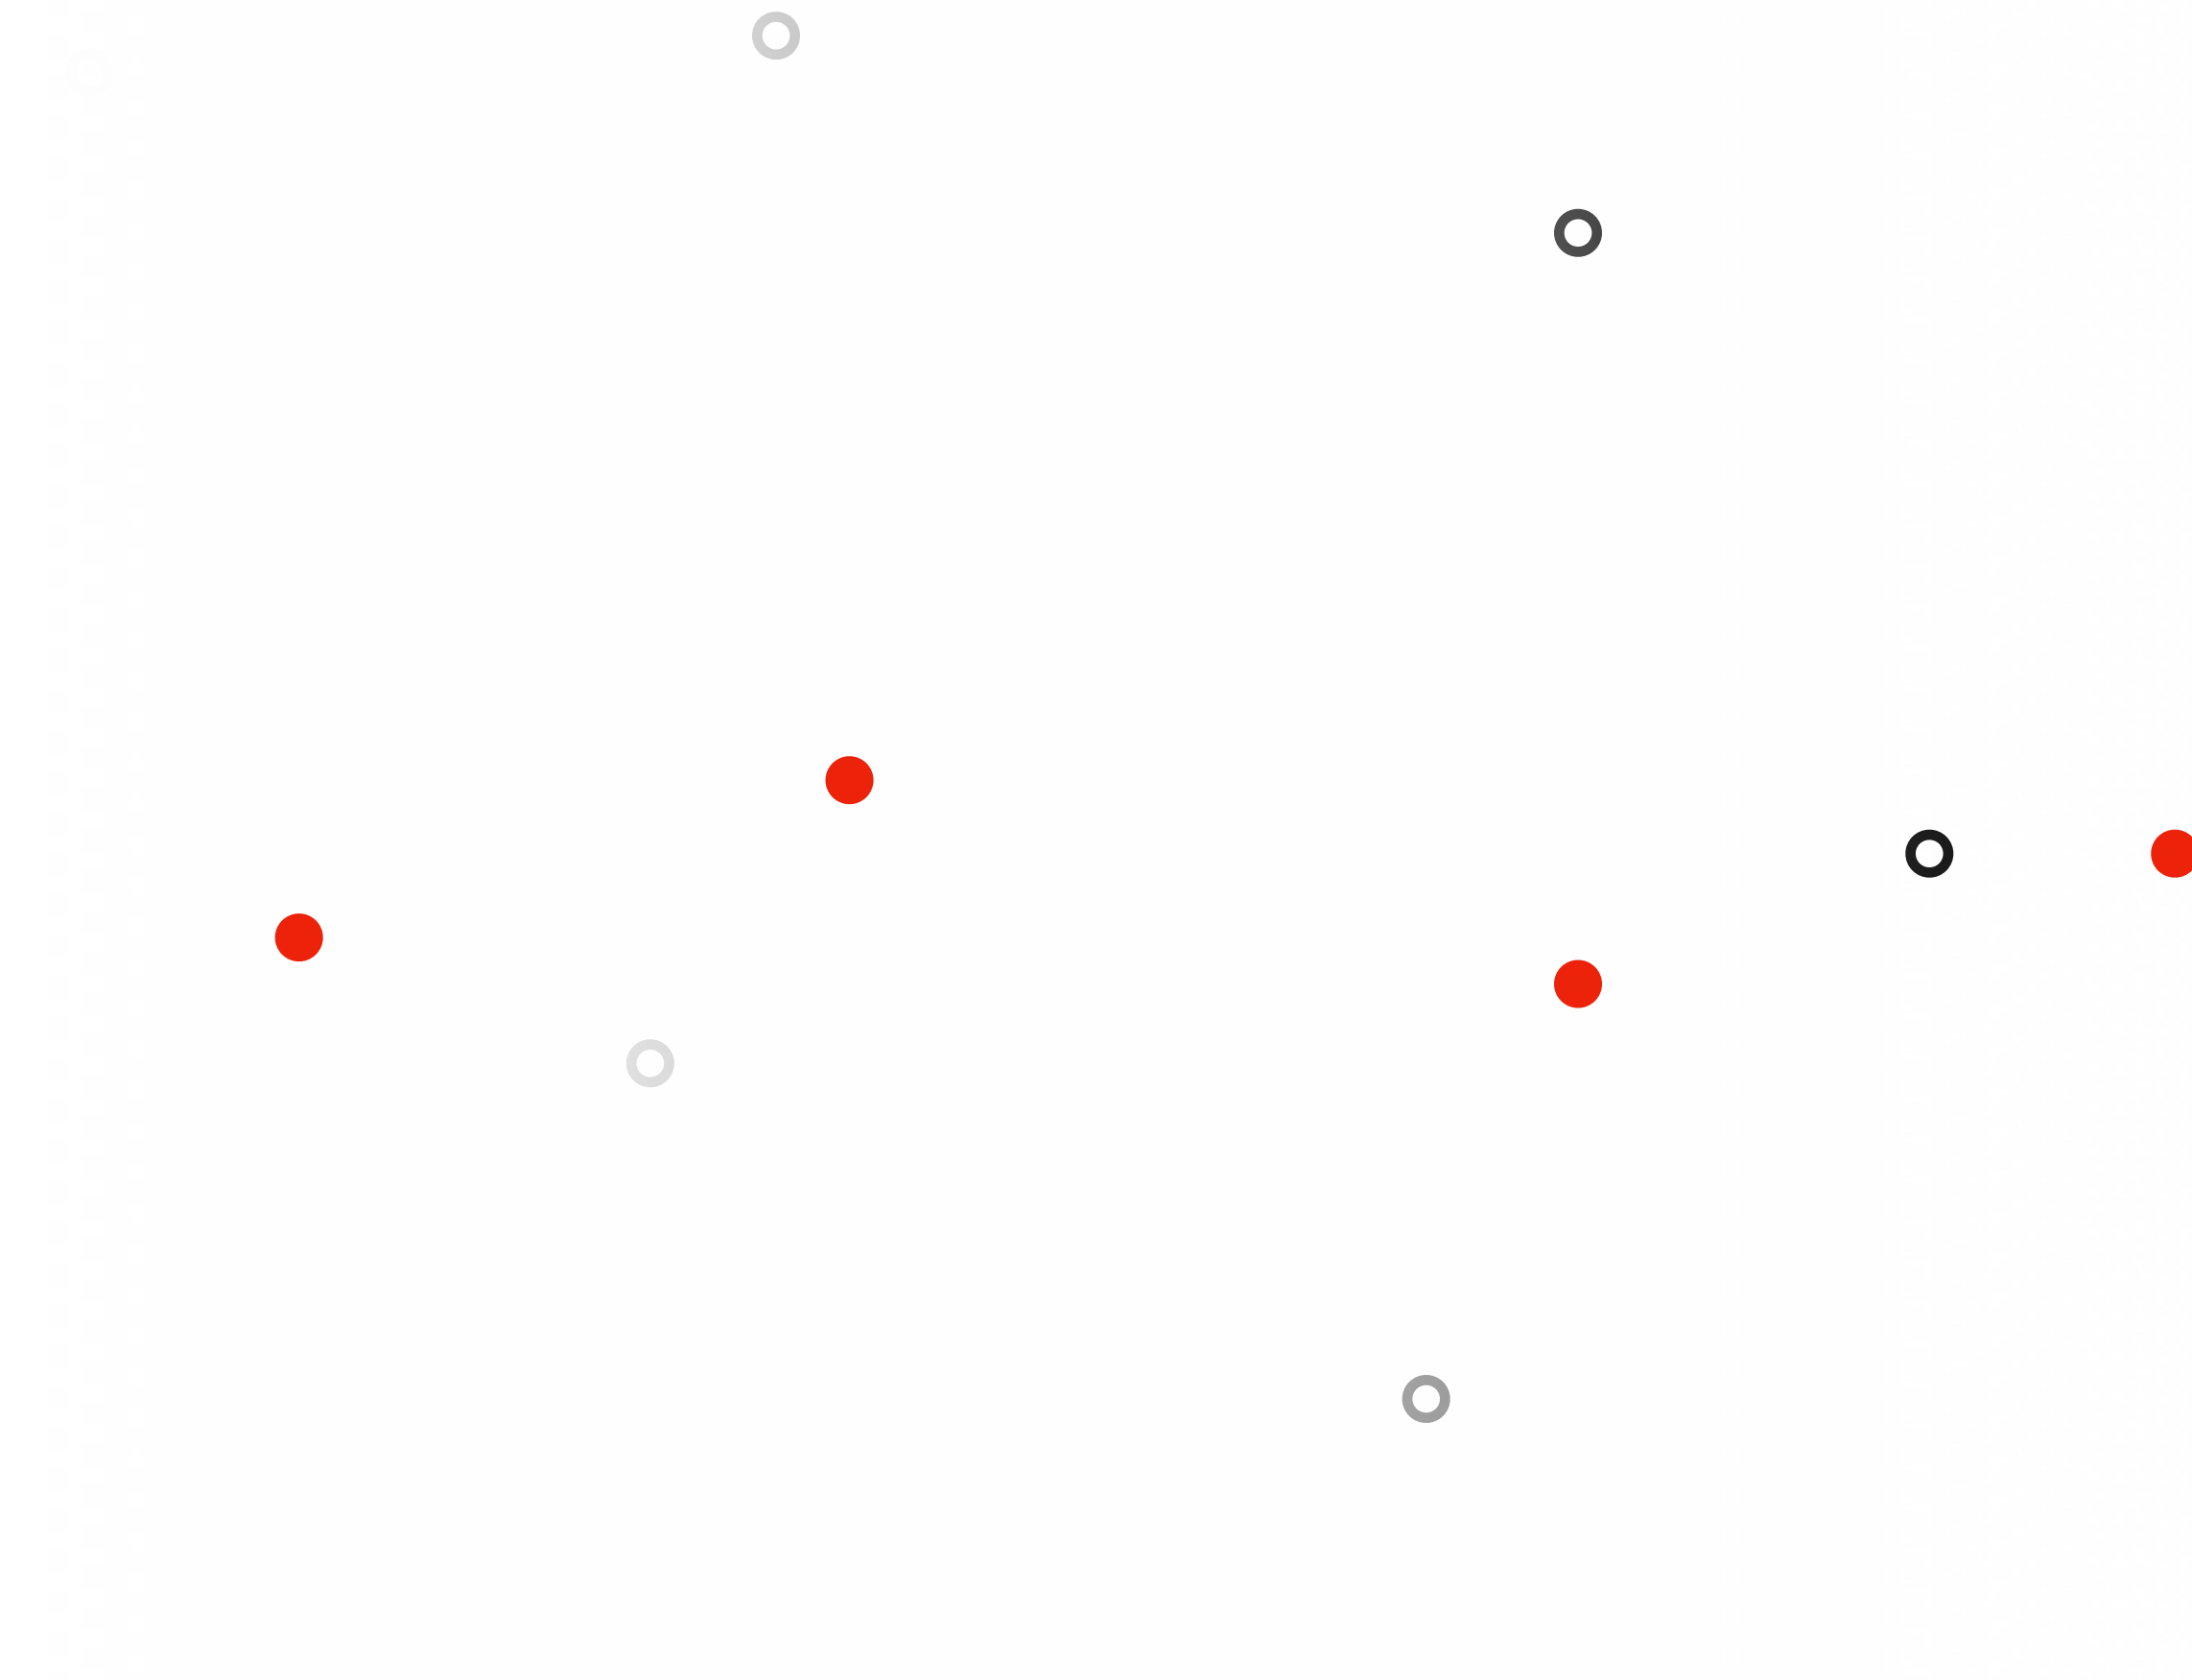
- Each particle does a random walk and compute a score
- No re-calibration based on observations
- Score decreases at each factor statement



HMM: Importance Sampling

Problem:

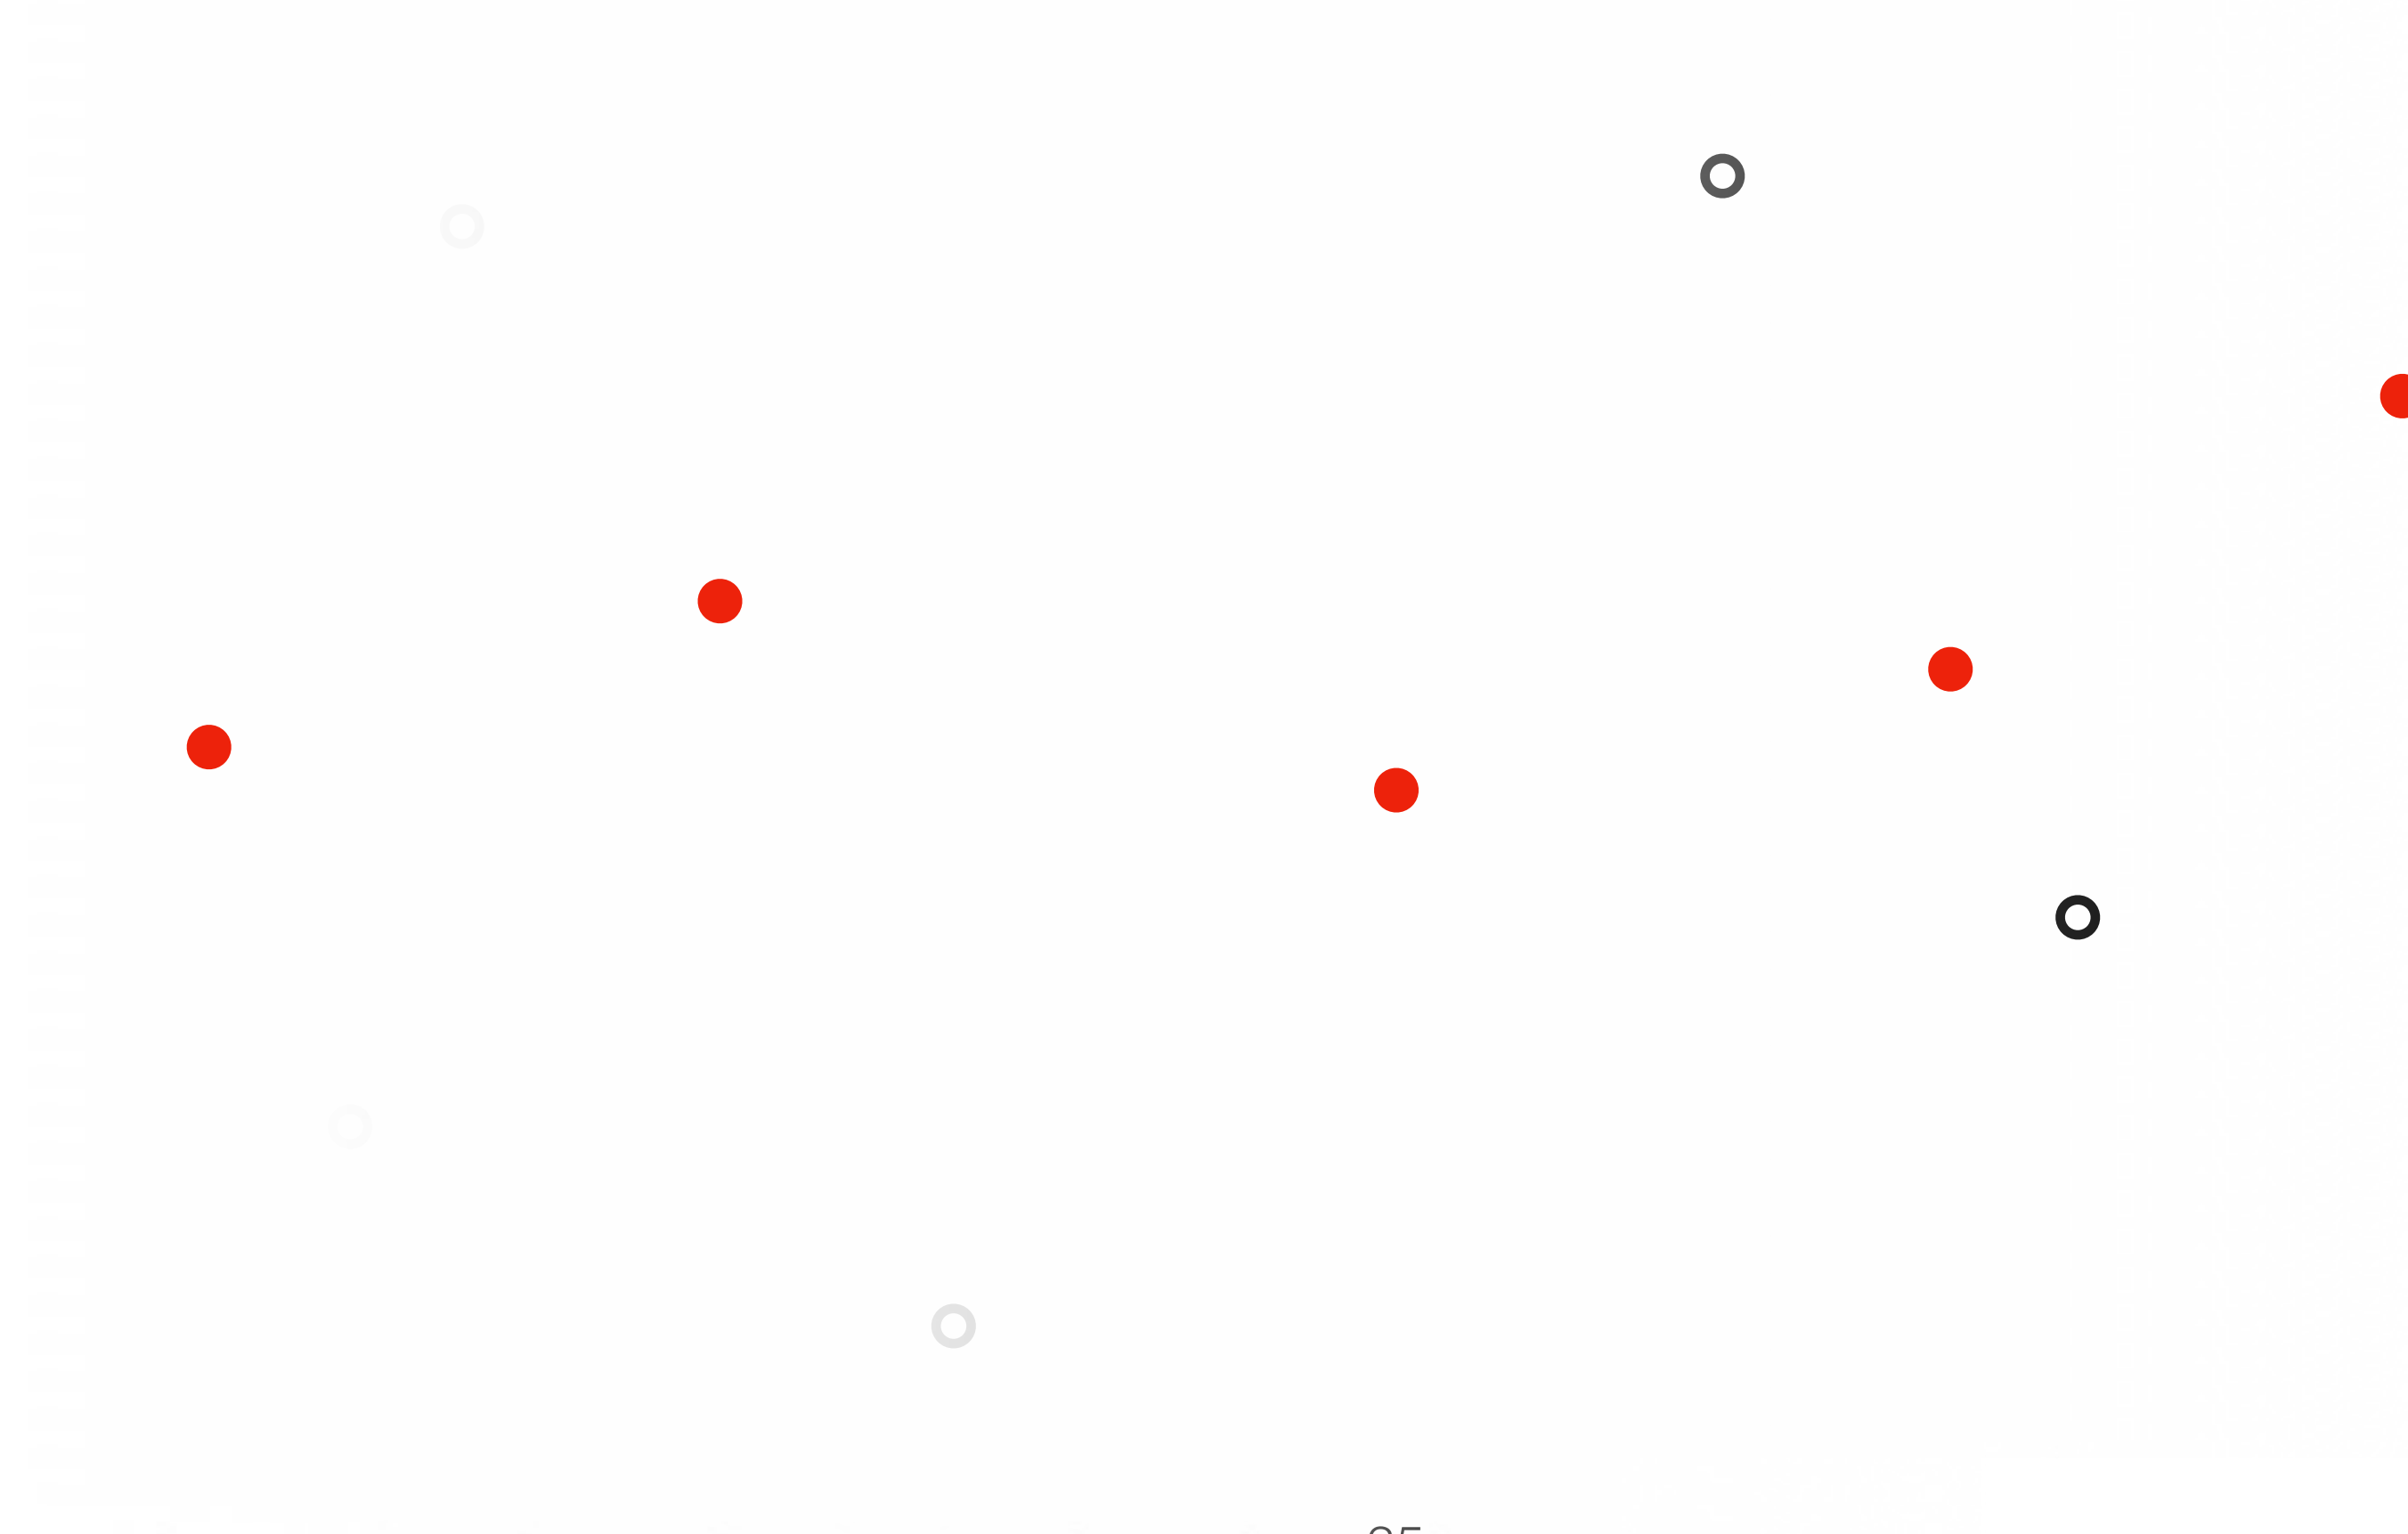
- Each particle does a random walk and compute a score
- No re-calibration based on observations
- Score decreases at each factor statement



HMM: Importance Sampling

Problem:

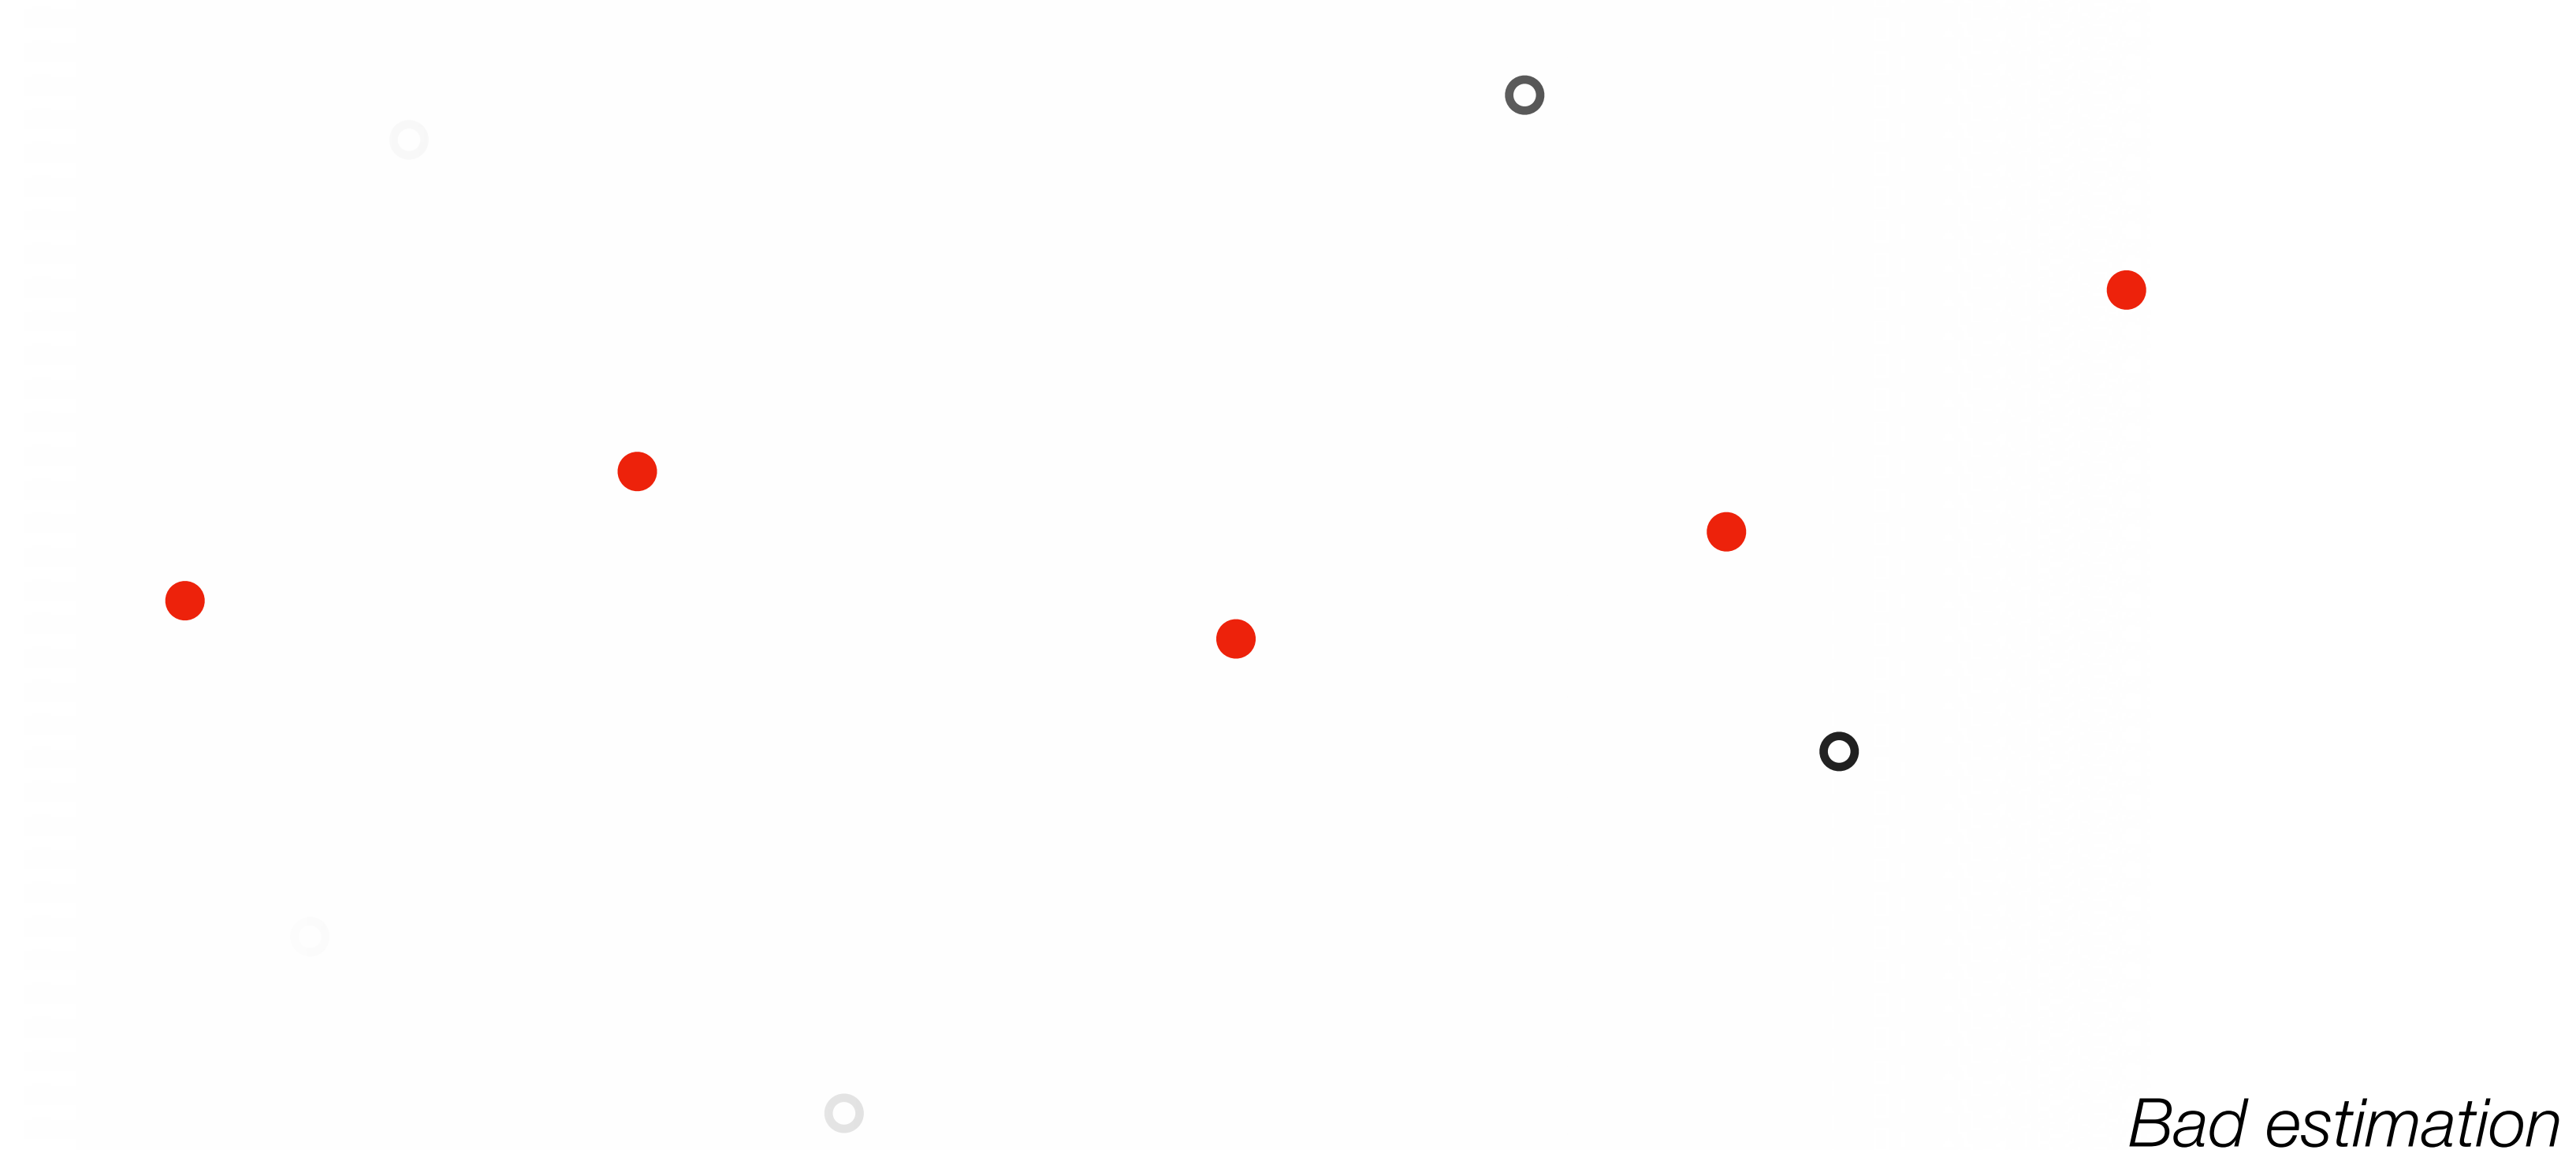
- Each particle does a random walk and compute a score
- No re-calibration based on observations
- Score decreases at each factor statement



HMM: Importance Sampling

Problem:

- Each particle does a random walk and compute a score
- No re-calibration based on observations
- Score decreases at each factor statement



HMM: Particle Filter

Add a resampling step

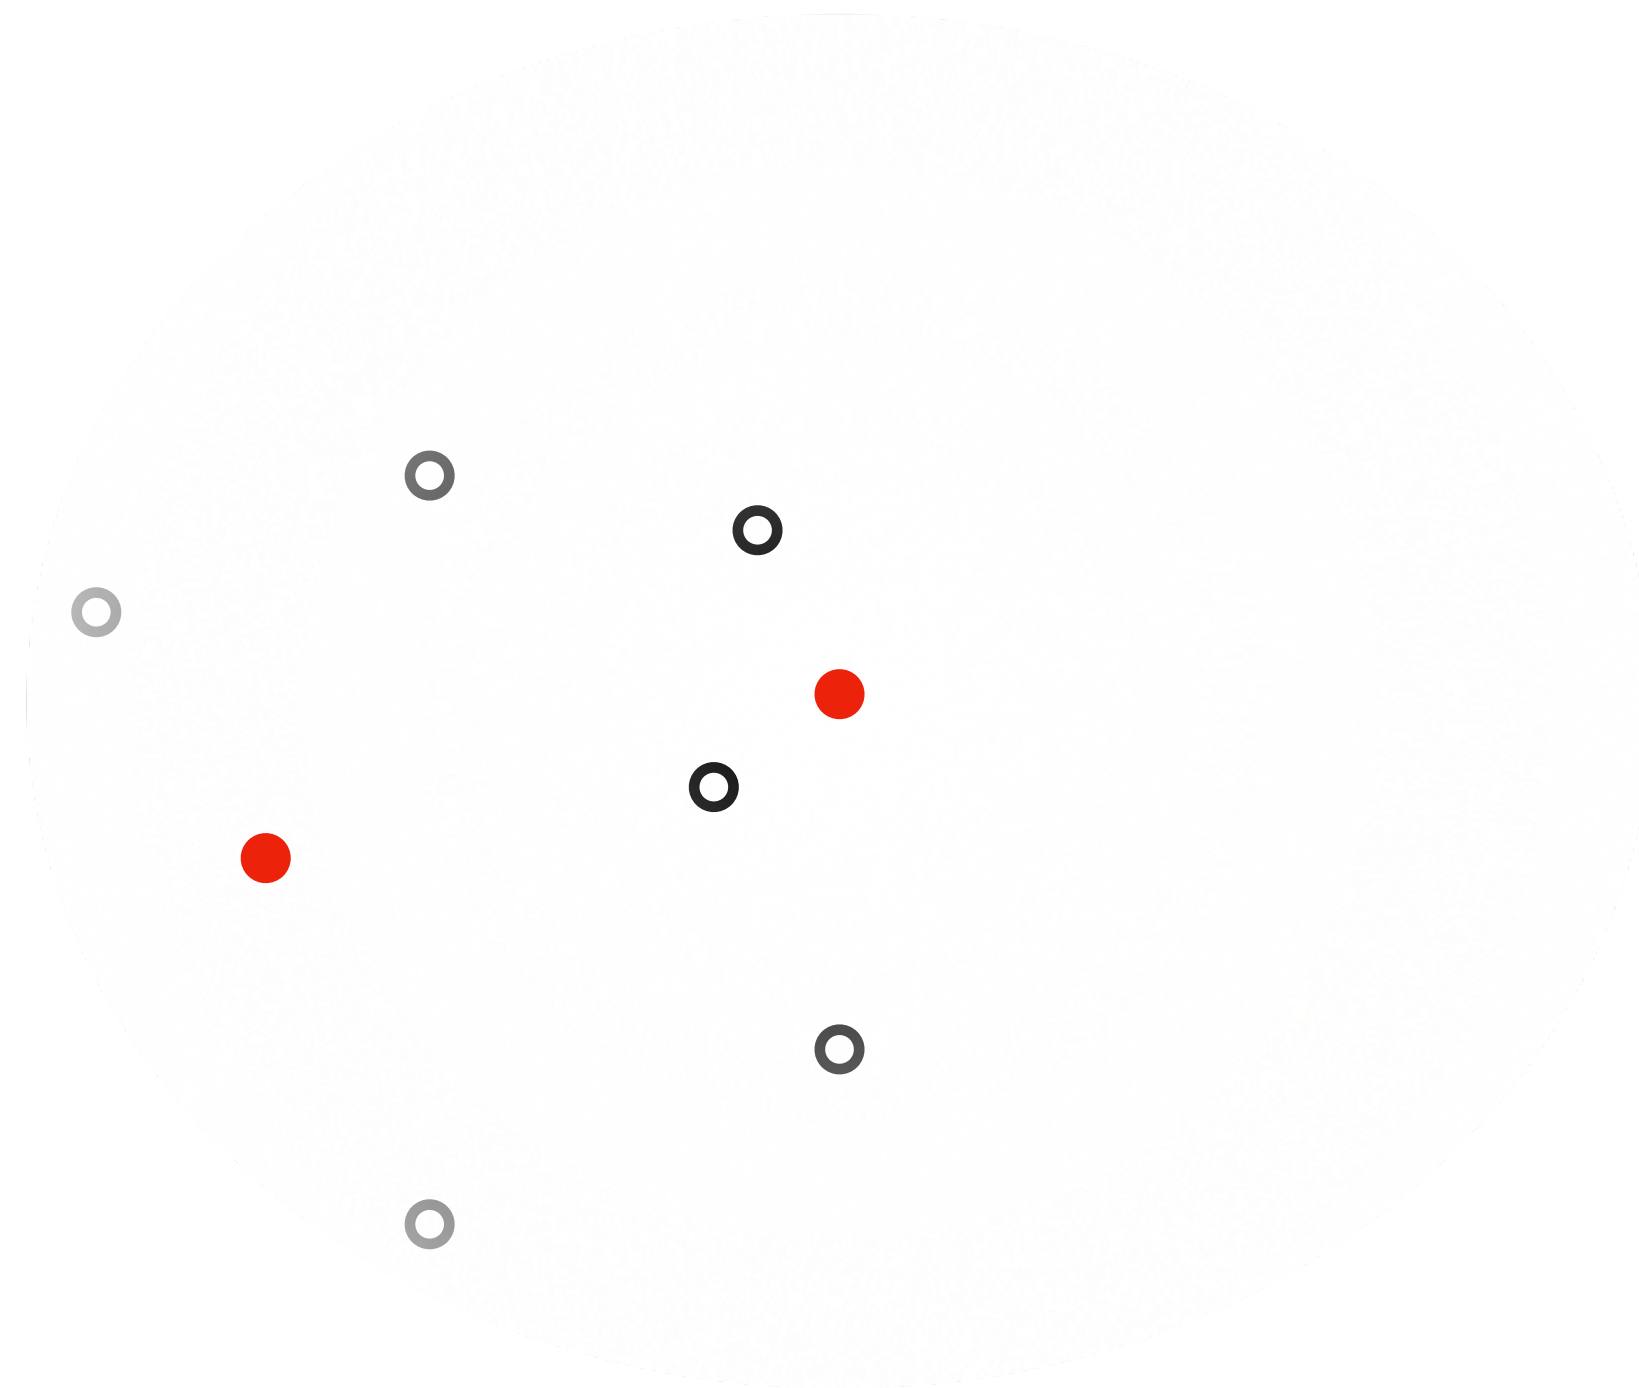
- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution



HMM: Particle Filter

Add a resampling step

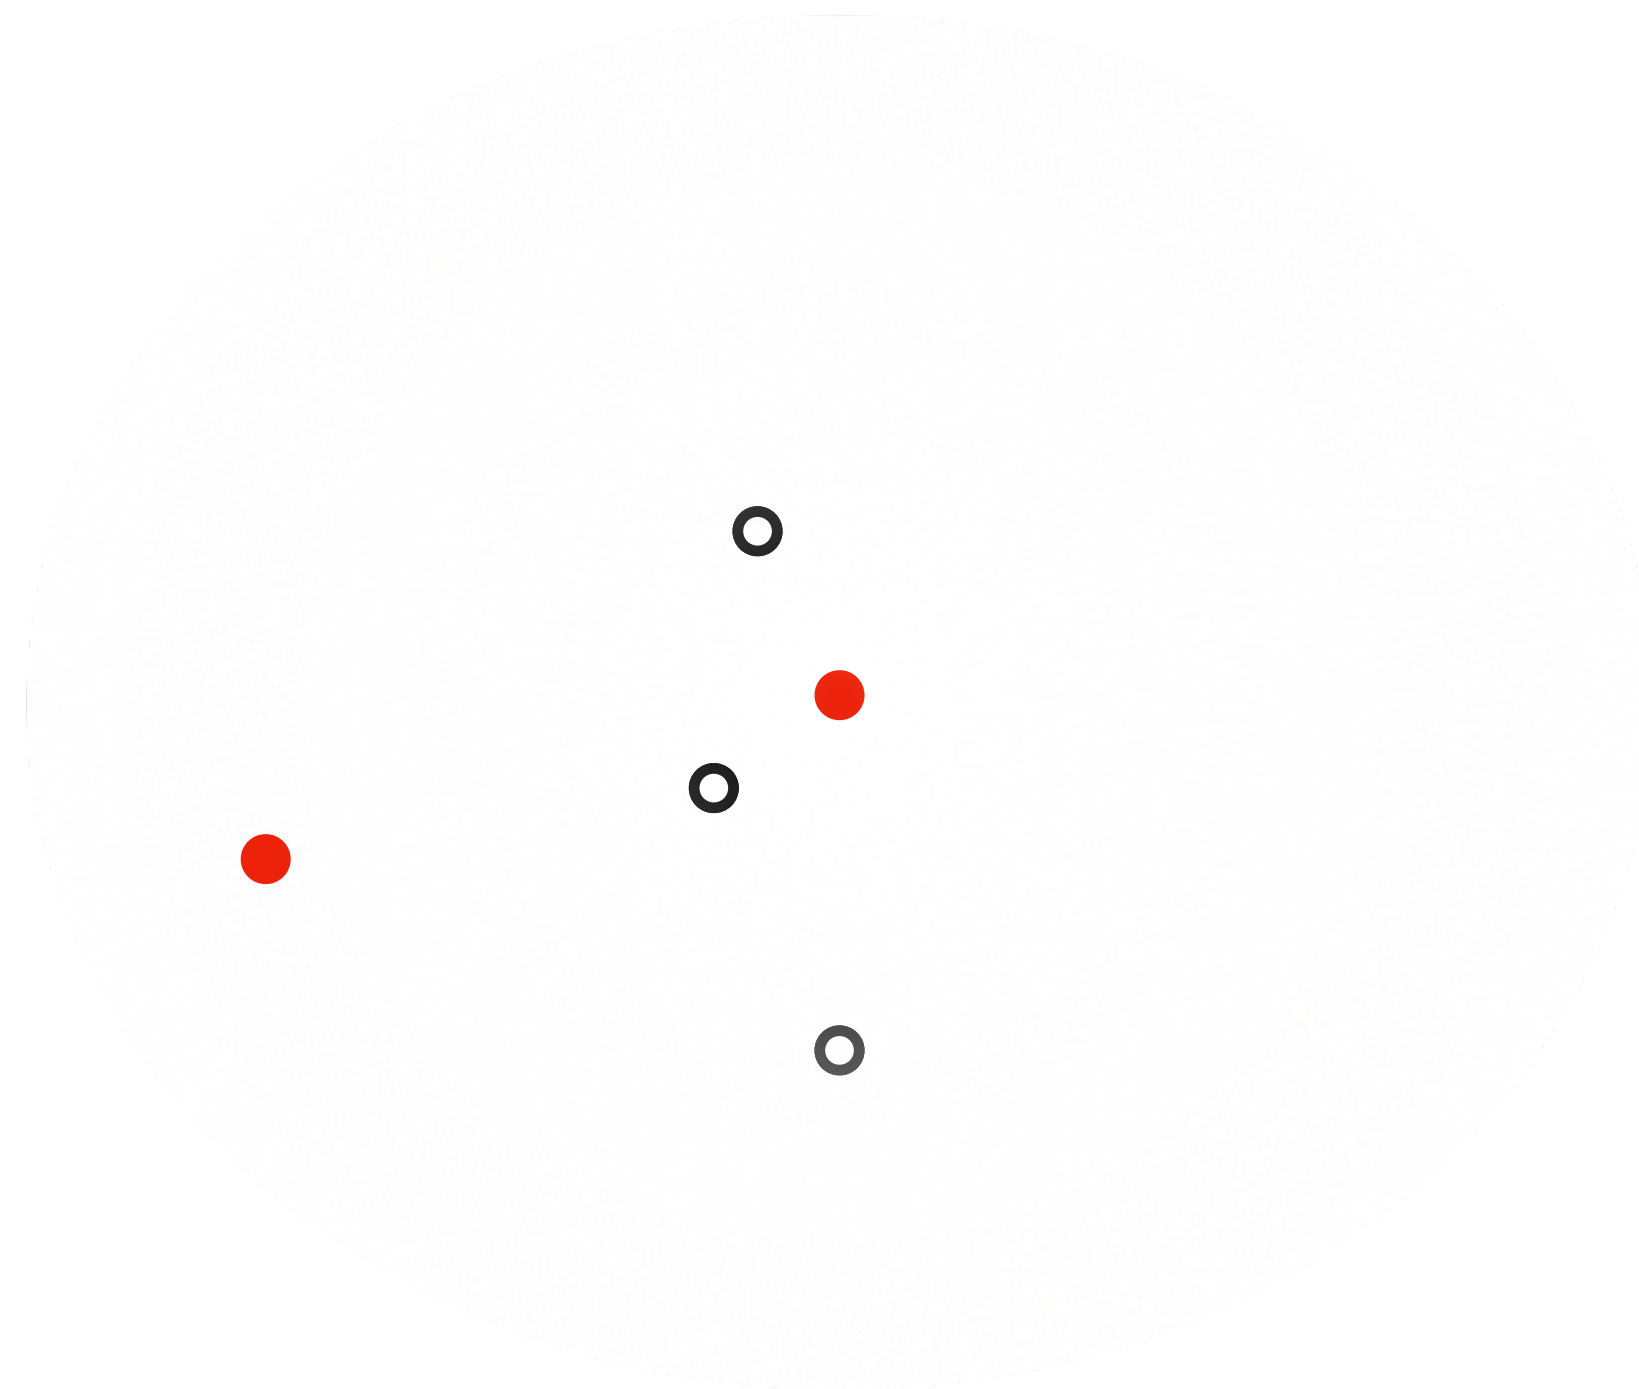
- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution



HMM: Particle Filter

Add a resampling step

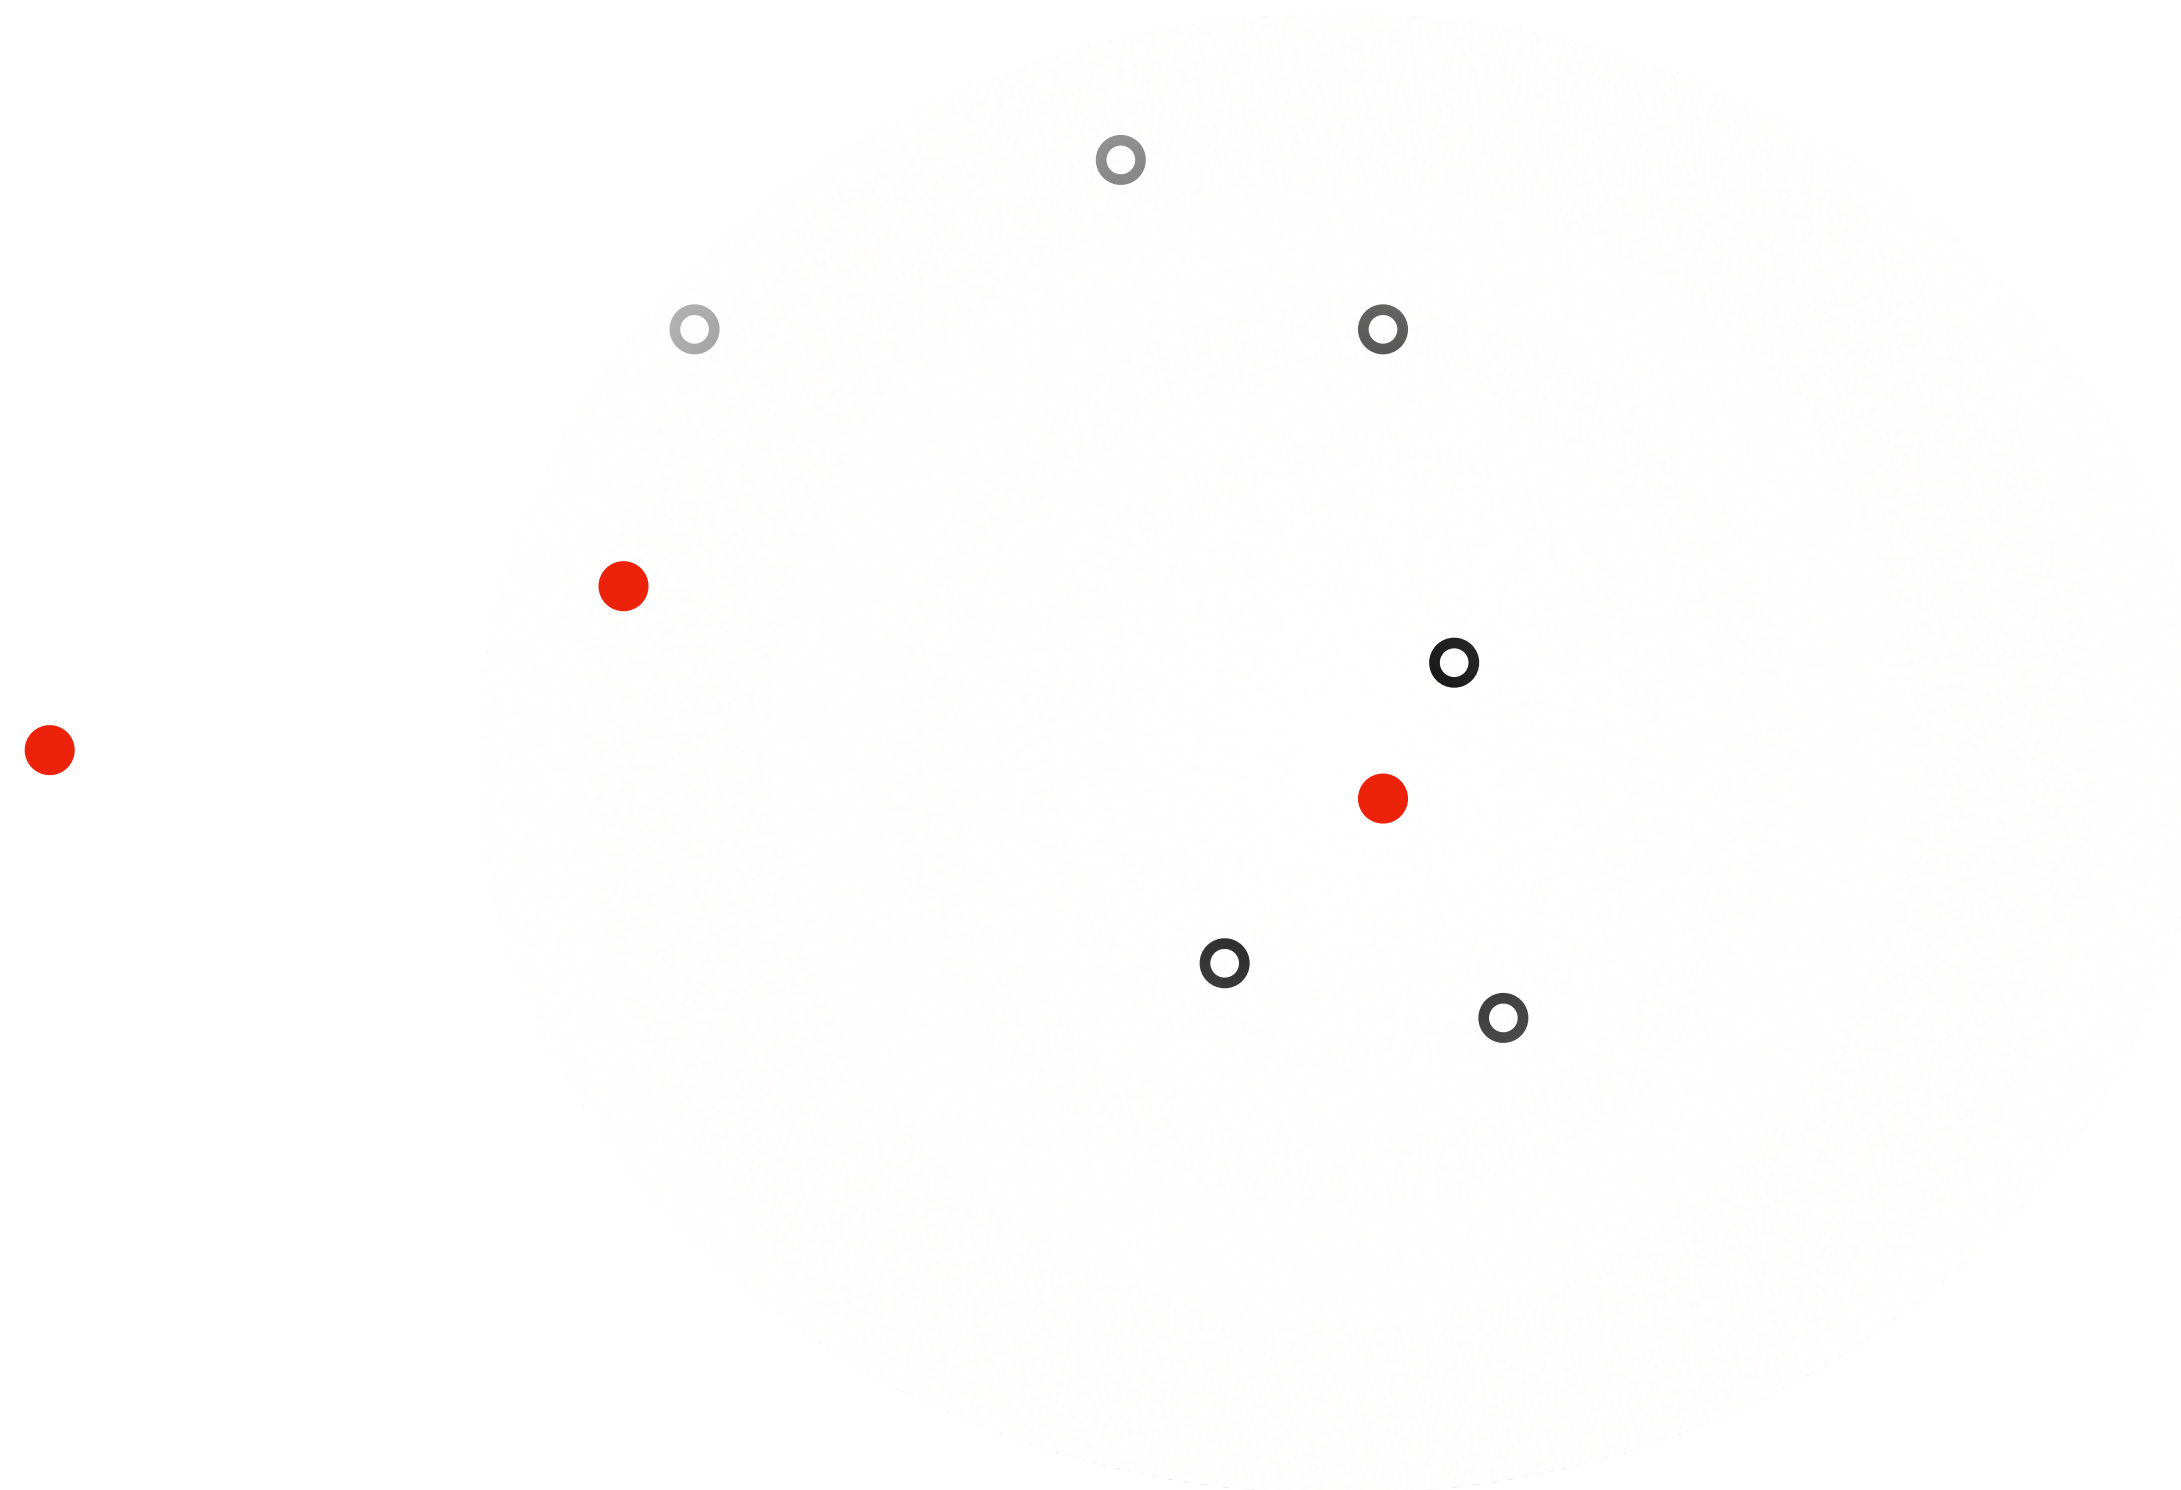
- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution



HMM: Particle Filter

Add a resampling step

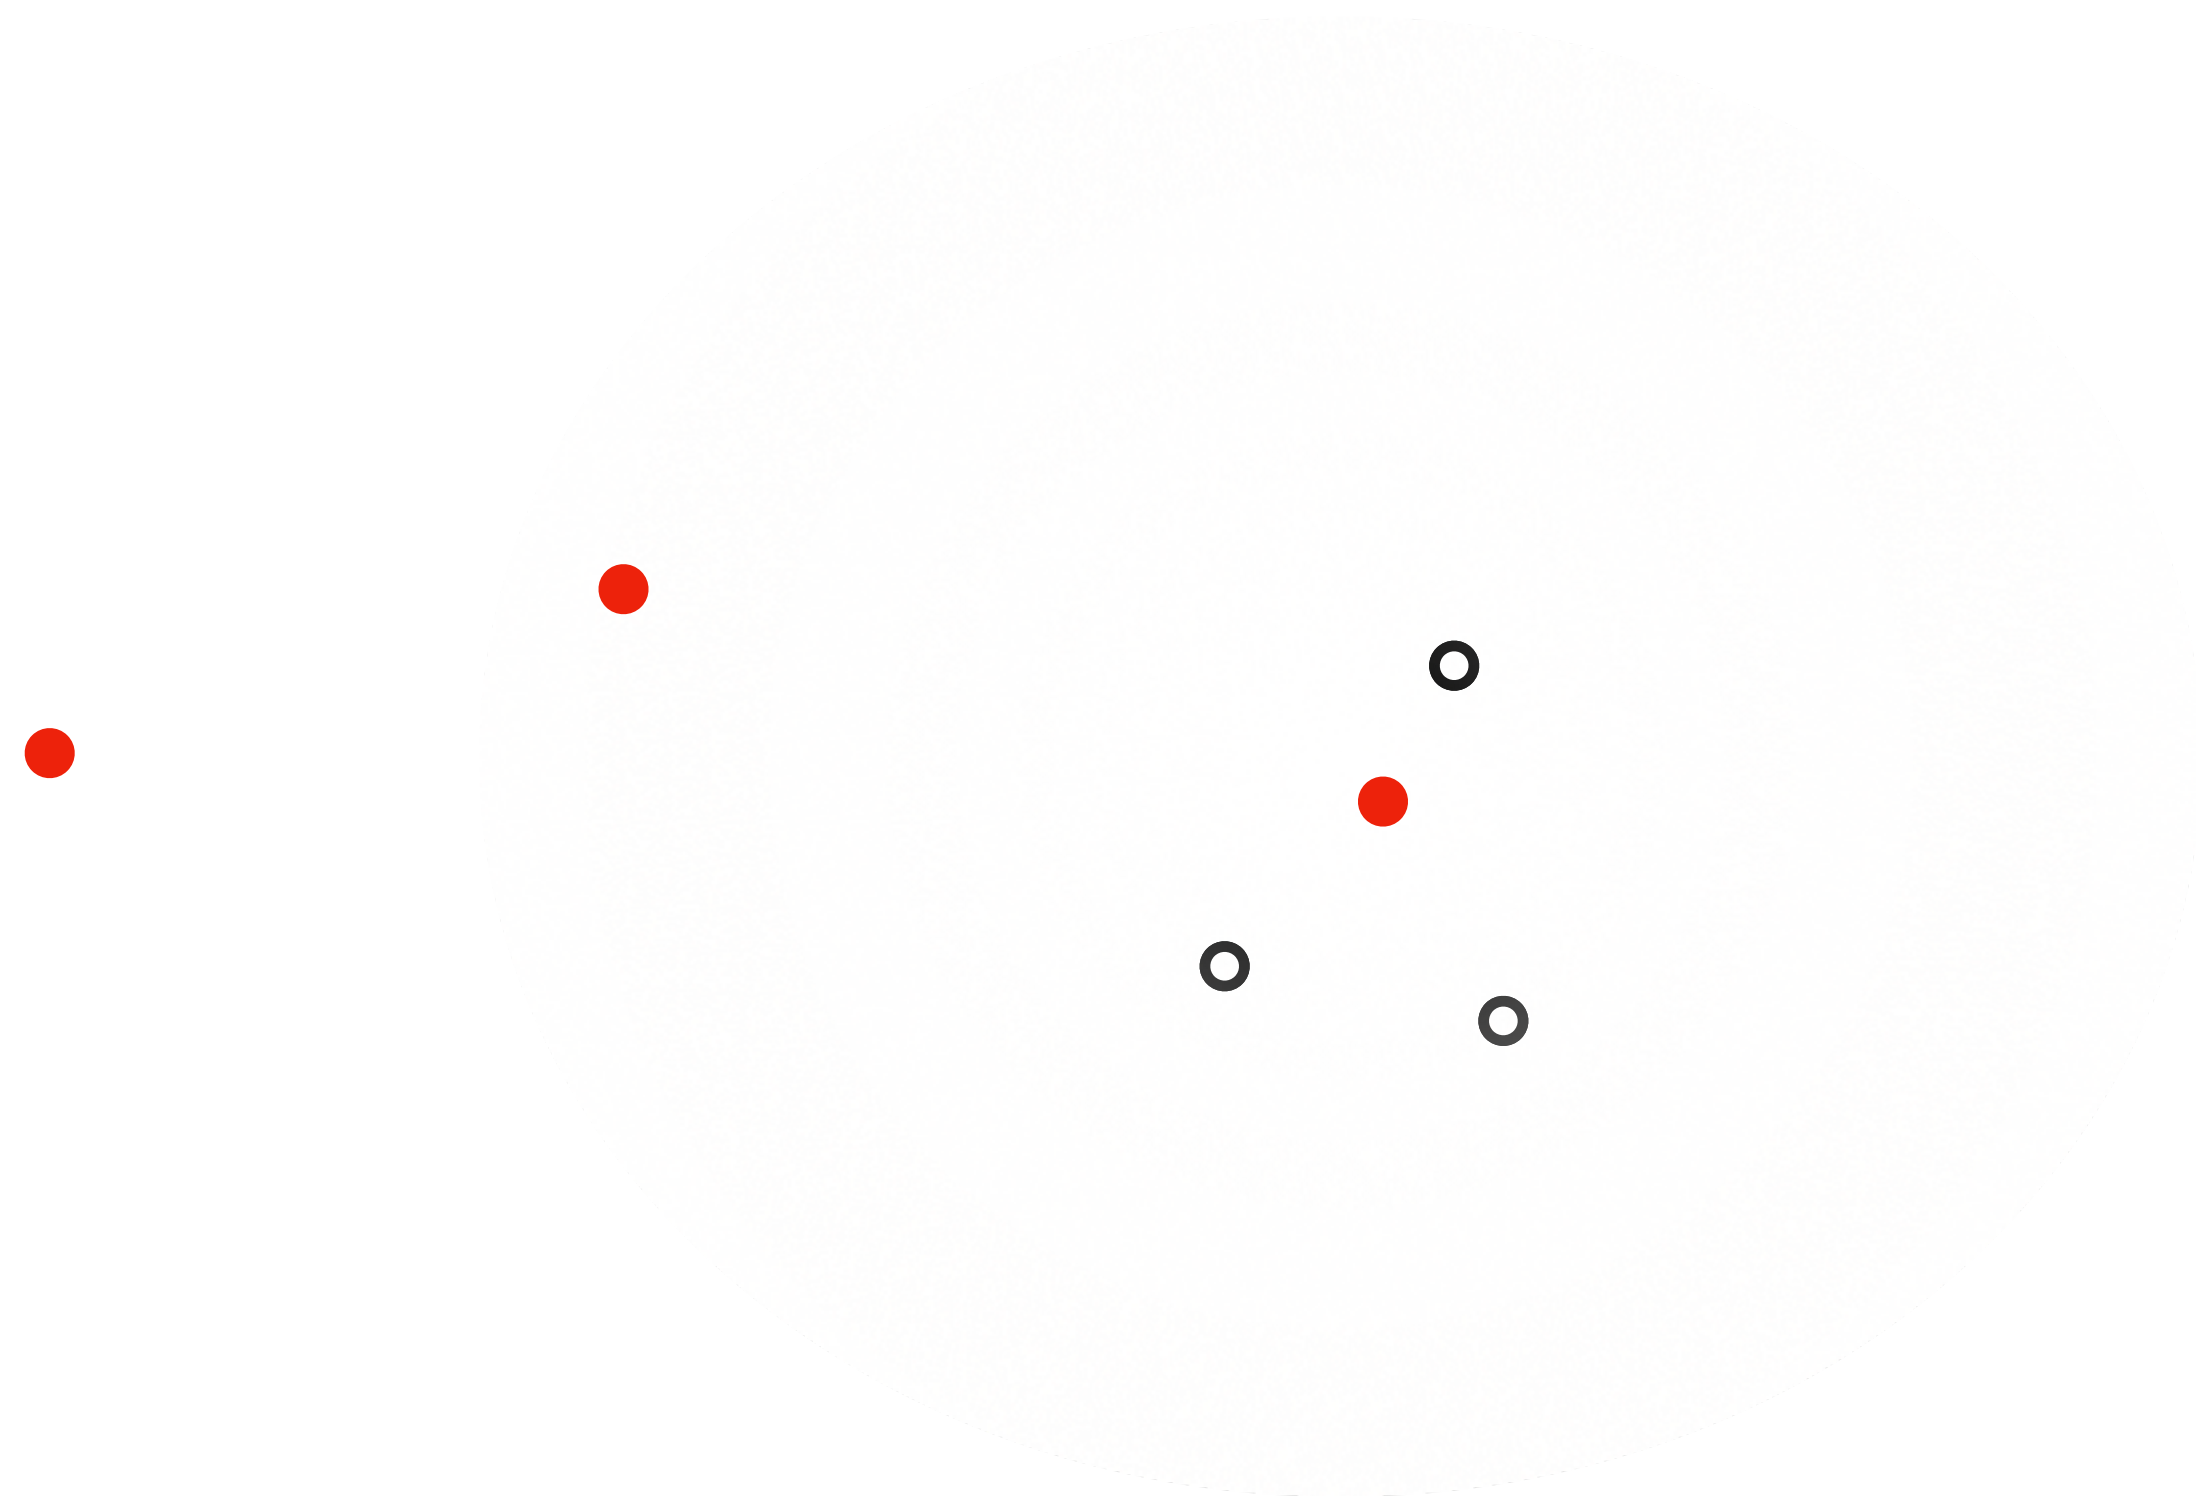
- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution



HMM: Particle Filter

Add a resampling step

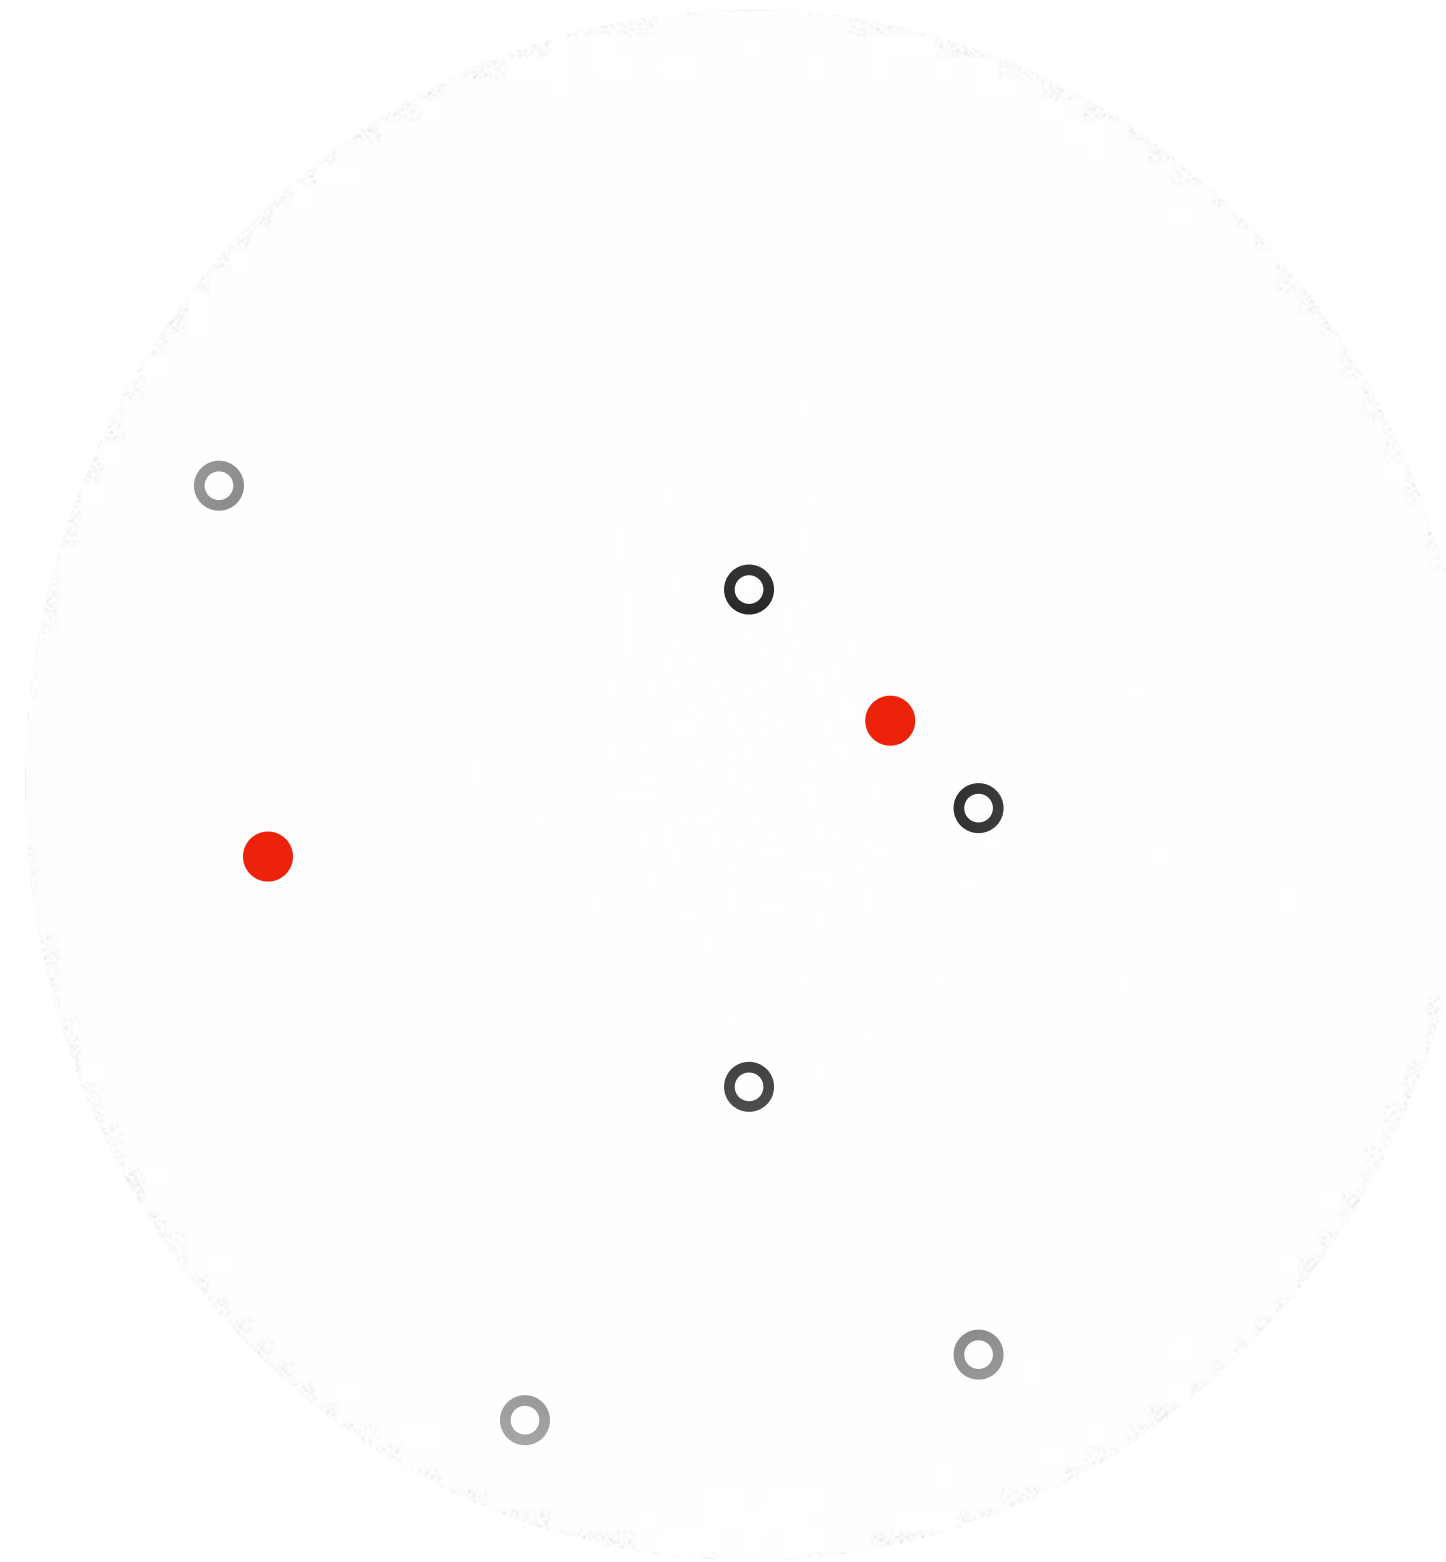
- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution



HMM: Particle Filter

Add a resampling step

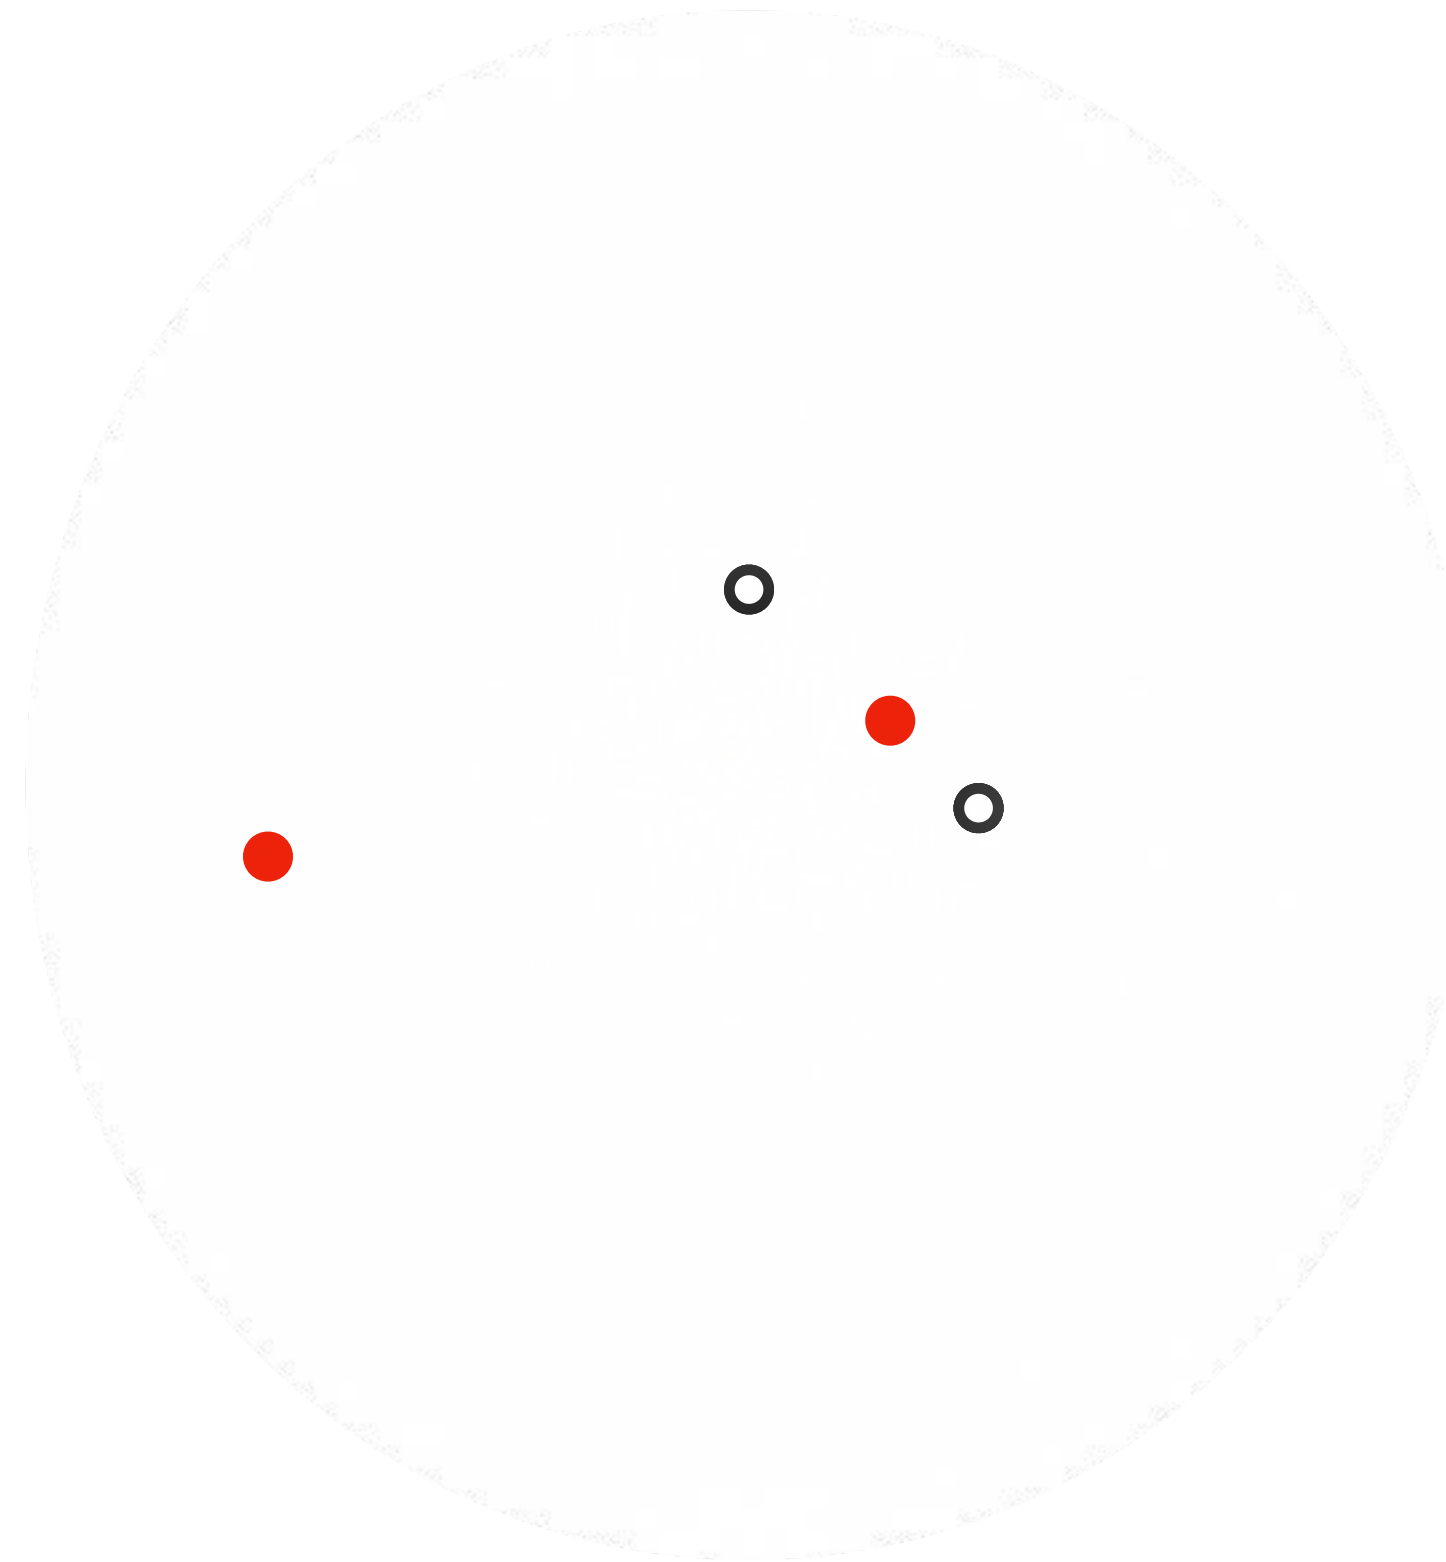
- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution



HMM: Particle Filter

Add a resampling step

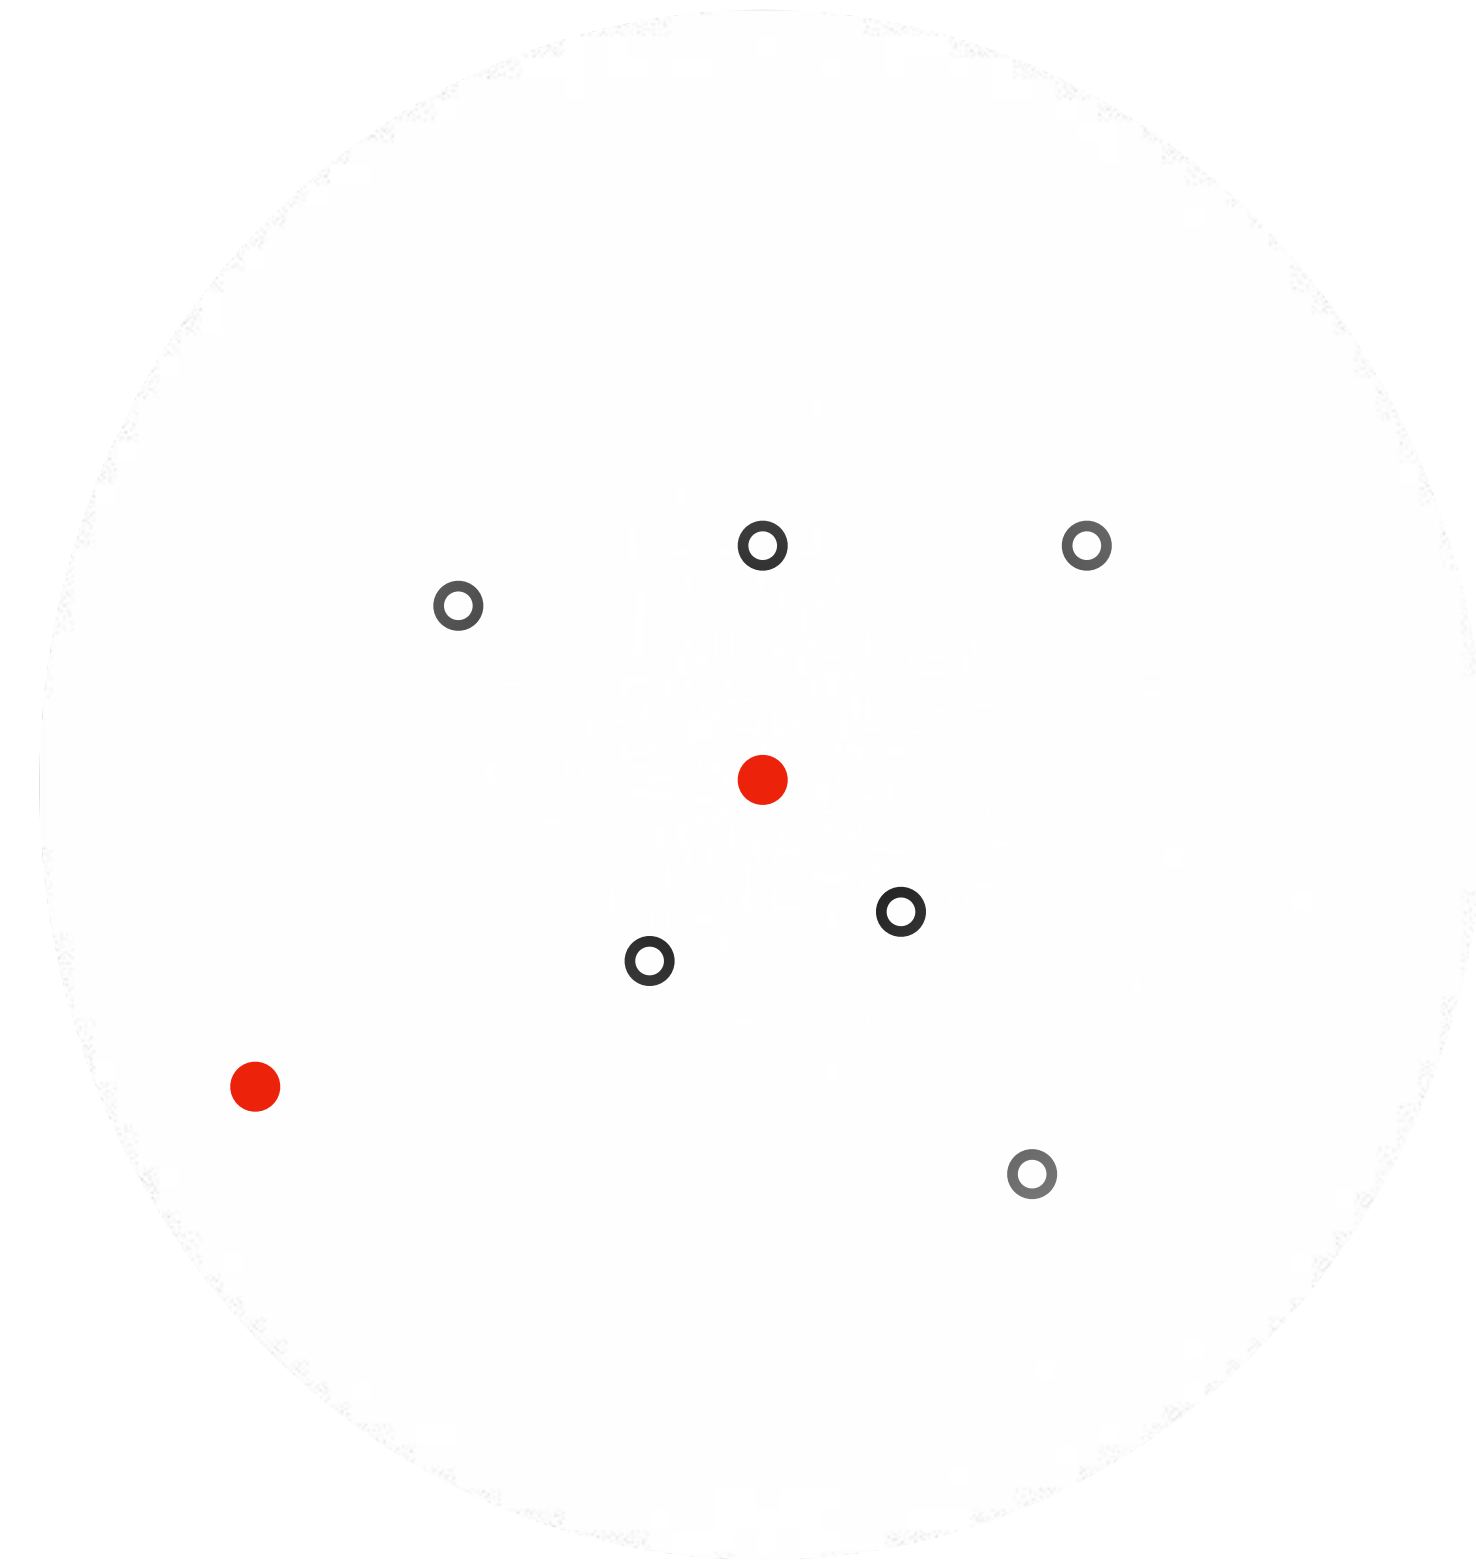
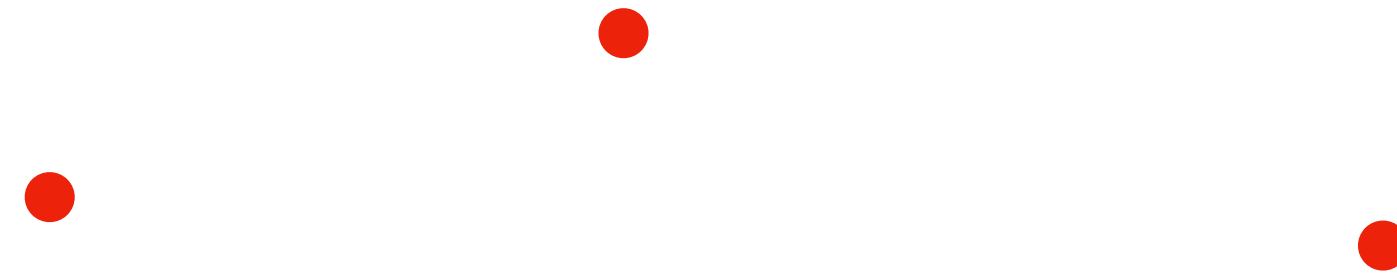
- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution



HMM: Particle Filter

Add a resampling step

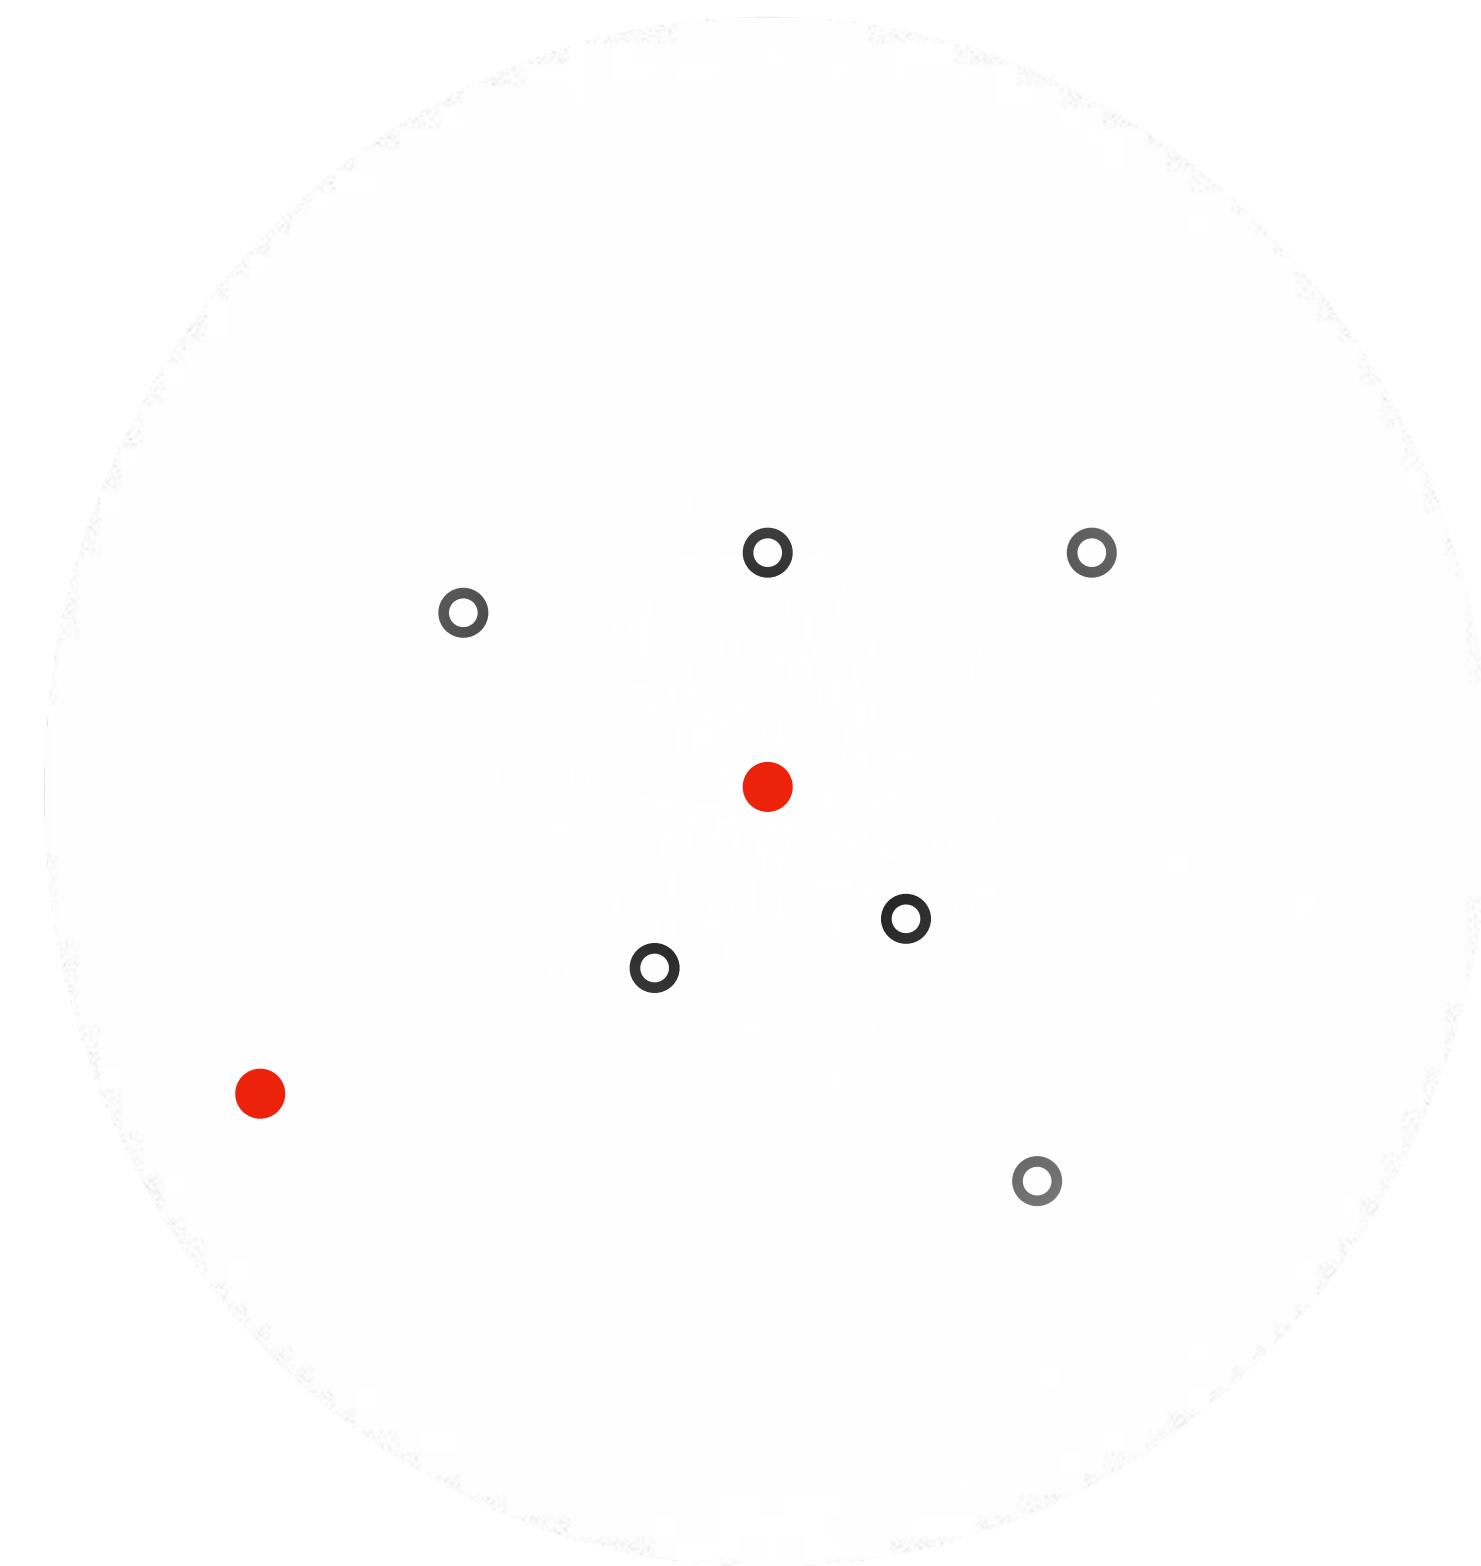
- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution



HMM: Particle Filter

Add a resampling step

- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution



Good estimation

Problem: Duplications

Problem: Duplications

How can we duplicate a particle during execution?

- Rerun the particle from the start?
- Force reuse sampled values?
- Clone the memory state?

Problem: Duplications

How can we duplicate a particle during execution?

- Rerun the particle from the start?
- Force reuse sampled values?
- Clone the memory state?

Continuation Passing Style

- Functions take an extra argument `k`: the continuation
- `k` implements what should be done with the result of the function
- In our context, we can use continuation to interrupt/restart the execution of a model

Continuation Passing Style (CPS)

BYO-PPL

Reminders: CPS

tree.ml

```
let rec tree_height t =  
  match t with  
  | Empty → 0  
  | Node (_, l, r) → 1 + max (tree_height l) (tree_height r)
```

Reminders: CPS

tree.ml

```
let rec tree_height t =  
  match t with  
  | Empty → 0  
  | Node (_, l, r) → 1 + max (tree_height l) (tree_height r)
```

```
let rec tree_height t =  
  match t with  
  | Empty → 0  
  | Node (_, l, r) →  
    let hl = tree_height l in  
    let hr = tree_height r in  
    (1 + max hl hr)
```

1. Add intermediate values

Reminders: CPS

tree.ml

```
let rec tree_height t =  
  match t with  
  | Empty → 0  
  | Node (_, l, r) → 1 + max (tree_height l) (tree_height r)
```

```
let rec tree_height t k =  
  match t with  
  | Empty → k 0  
  | Node (_, l, r) →  
    let hl = tree_height l in  
    let hr = tree_height r in  
    k (1 + max hl hr)
```

1. Add intermediate values
2. Add call to continuation

Reminders: CPS

tree.ml

```
let rec tree_height t =  
  match t with  
  | Empty → 0  
  | Node (_, l, r) → 1 + max (tree_height l) (tree_height r)
```

```
let rec tree_height t k =  
  match t with  
  | Empty → k 0  
  | Node (_, l, r) →  
    tree_height l (fun hl →  
      tree_height r (fun hr →  
        k (1 + max hl hr)))
```

1. Add intermediate values
2. Add call to continuation
3. Turn let/in into nested function call

Funny Bernoulli CPS

funny_bernoulli.ml

```
let funny_bernoulli () =  
  let a = sample (bernoulli ~p:0.5) in  
  let b = sample (bernoulli ~p:0.5) in  
  let c = sample (bernoulli ~p:0.5) in  
  let () = assume (a = 1 || b = 1) in  
  a + b + c
```

1. Add intermediate values
2. Add call to continuation
3. Turn let/in into nested function call

Funny Bernoulli CPS

funny_bernoulli.ml

```
let funny_bernoulli () =  
  let a = sample (bernoulli ~p:0.5) in  
  let b = sample (bernoulli ~p:0.5) in  
  let c = sample (bernoulli ~p:0.5) in  
  let () = assume (a = 1 || b = 1) in  
  a + b + c
```

```
let funny_bernoulli () k =  
  sample (bernoulli ~p:0.5) (fun a →  
    sample (bernoulli ~p:0.5) (fun b →  
      sample (bernoulli ~p:0.5) (fun c →  
        assume (a = 1 || b = 1) (fun () →  
          k (a + b + c))))
```

1. Add intermediate values
2. Add call to continuation
3. Turn let/in into nested function call

CPS Monadic Operators

cps_operators.ml

```
let return e k = k e
```

```
let ( let* ) e f k = e (fun x → f x k)    (* let* x = e in f(x) *)
```

CPS Monadic Operators

cps_operators.ml

```
let return e k = k e
```

```
let ( let* ) e f k = e (fun x → f x k)    (* let* x = e in f(x) *)
```

```
let funny_bernoulli () k =  
  sample (bernoulli ~p:0.5) (fun a →  
    sample (bernoulli ~p:0.5) (fun b →  
      sample (bernoulli ~p:0.5) (fun c →  
        assume (a = 1 || b = 1) (fun () →  
          k (a + b + c))))
```

```
let funny_bernoulli () =  
  let* a = sample (bernoulli ~p:0.5) in  
  let* b = sample (bernoulli ~p:0.5) in  
  let* c = sample (bernoulli ~p:0.5) in  
  let* () = assume (a = 1 || b = 1) in  
  return (a + b + c)
```

Sample Generation (CPS)

BYO-PPL

CPS Models

infer.ml

```
module Gen : sig
  type 'a prob
  and 'a next = 'a prob → 'a prob
  and ('a, 'b) model = 'a → ('b → 'b next) → 'b next

  val sample : 'a Distribution.t → ('a → 'b next) → 'b next
  val factor : float → (unit → 'b next) → 'b next
  val draw: ('a, 'b) model → 'a → 'b
end = struct ... end
```

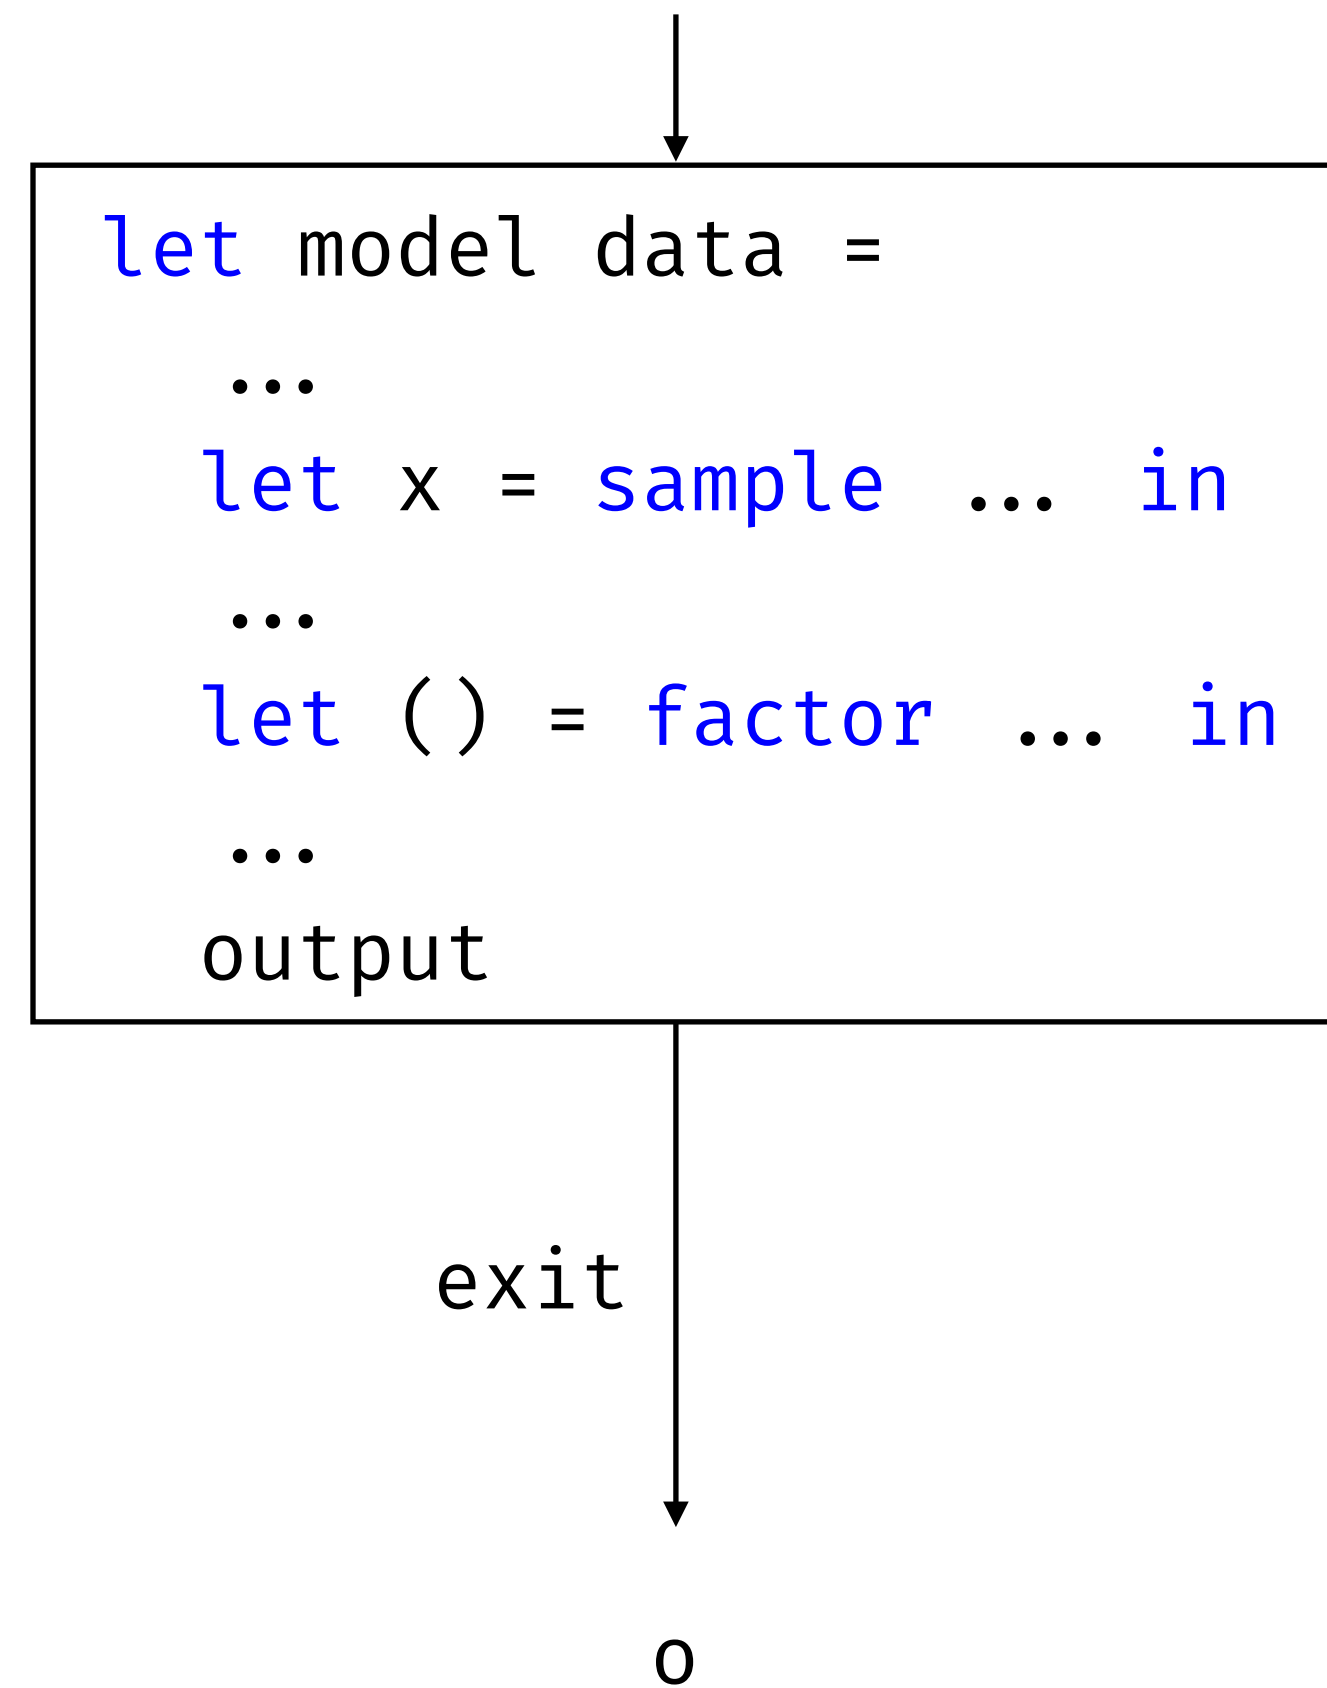
Type 'a prob

- Store all information required for inference (e.g., particles array)
- Type ('a, 'b) model capture input/output types

Models and probabilistic constructs are CPS functions

- Two arguments: input 'a and a continuation on the return value ('b → 'b next).
- The return value is a continuation 'a next that updates a probabilistic state of type 'a prob.

Sample Generation



Sample Generation

infer.ml

```
module Gen = struct
  type 'a prob = 'a option
  and 'a next = 'a prob → 'a prob
  and ('a, 'b) model = 'a → ('b → 'b next) → 'b next

  let exit v _prob = Some v

  let sample d k prob =
    let v = Distribution.draw d in
    k v prob

  let factor _s k prob = k () prob

  let draw m data =
    let v = (m data) exit None in
    Option.get v
end
```

Funny Bernoulli

funny_bernoulli.ml

```
open Infer.Gen
```

```
let funny_bernoulli () =  
  let* a = sample (bernoulli ~p:0.5) in  
  let* b = sample (bernoulli ~p:0.5) in  
  let* c = sample (bernoulli ~p:0.5) in  
  let* () = assume (a = 1 || b = 1) in  
  return (a + b + c)
```

```
let _ =  
  for _ = 1 to 10 do  
    let v = draw funny_bernoulli () in  
    Format.printf "%d " v  
  done
```

```
> dune exec ./examples/funny_bernoulli.exe
```

```
1 1 2 2 2 2 2 1 3 2
```

Importance Sampling (CPS)

BYO-PPL

Importance Sampling

infer.ml

```
module Importance_sampling : sig
  type 'a prob
  and 'a next = 'a prob → 'a prob
  and ('a, 'b) model = 'a → ('b → 'b next) → 'b next

  val sample : 'a Distribution.t → ('a → 'b next) → 'b next
  val factor : float → (unit → 'b next) → 'b next
  val infer : ('a, 'b) model → 'a → 'b Distribution.t
end = struct ... end
```

Inference algorithm

- Run a set of n independent executions
- **sample**: draw a sample from a distribution
- **factor**: associate a score to the current execution
- Gather output values and score to approximate the posterior distribution

Importance Sampling

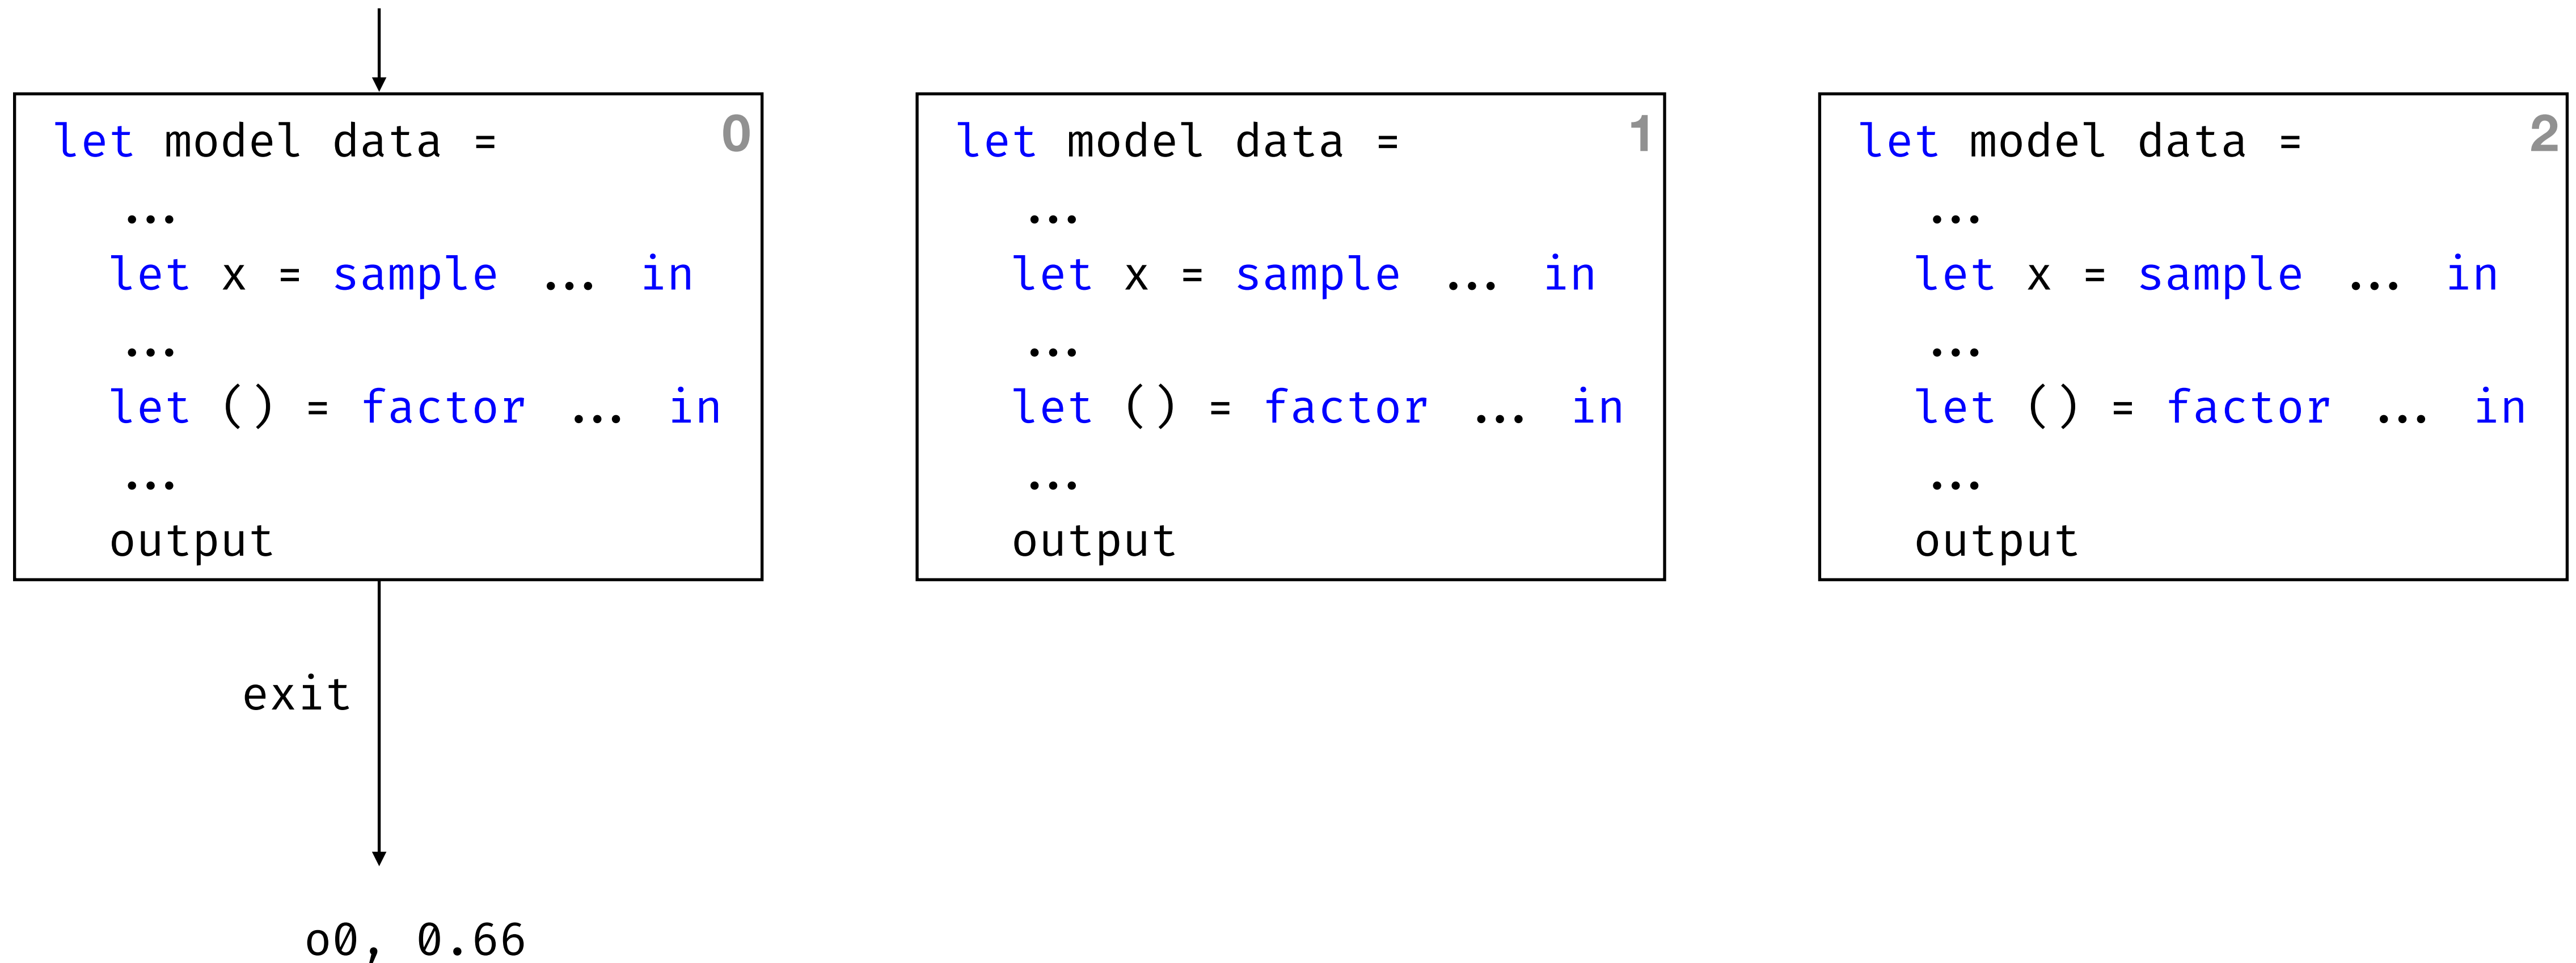


```
let model data = 0  
  ...  
  let x = sample ... in  
  ...  
  let () = factor ... in  
  ...  
  output
```

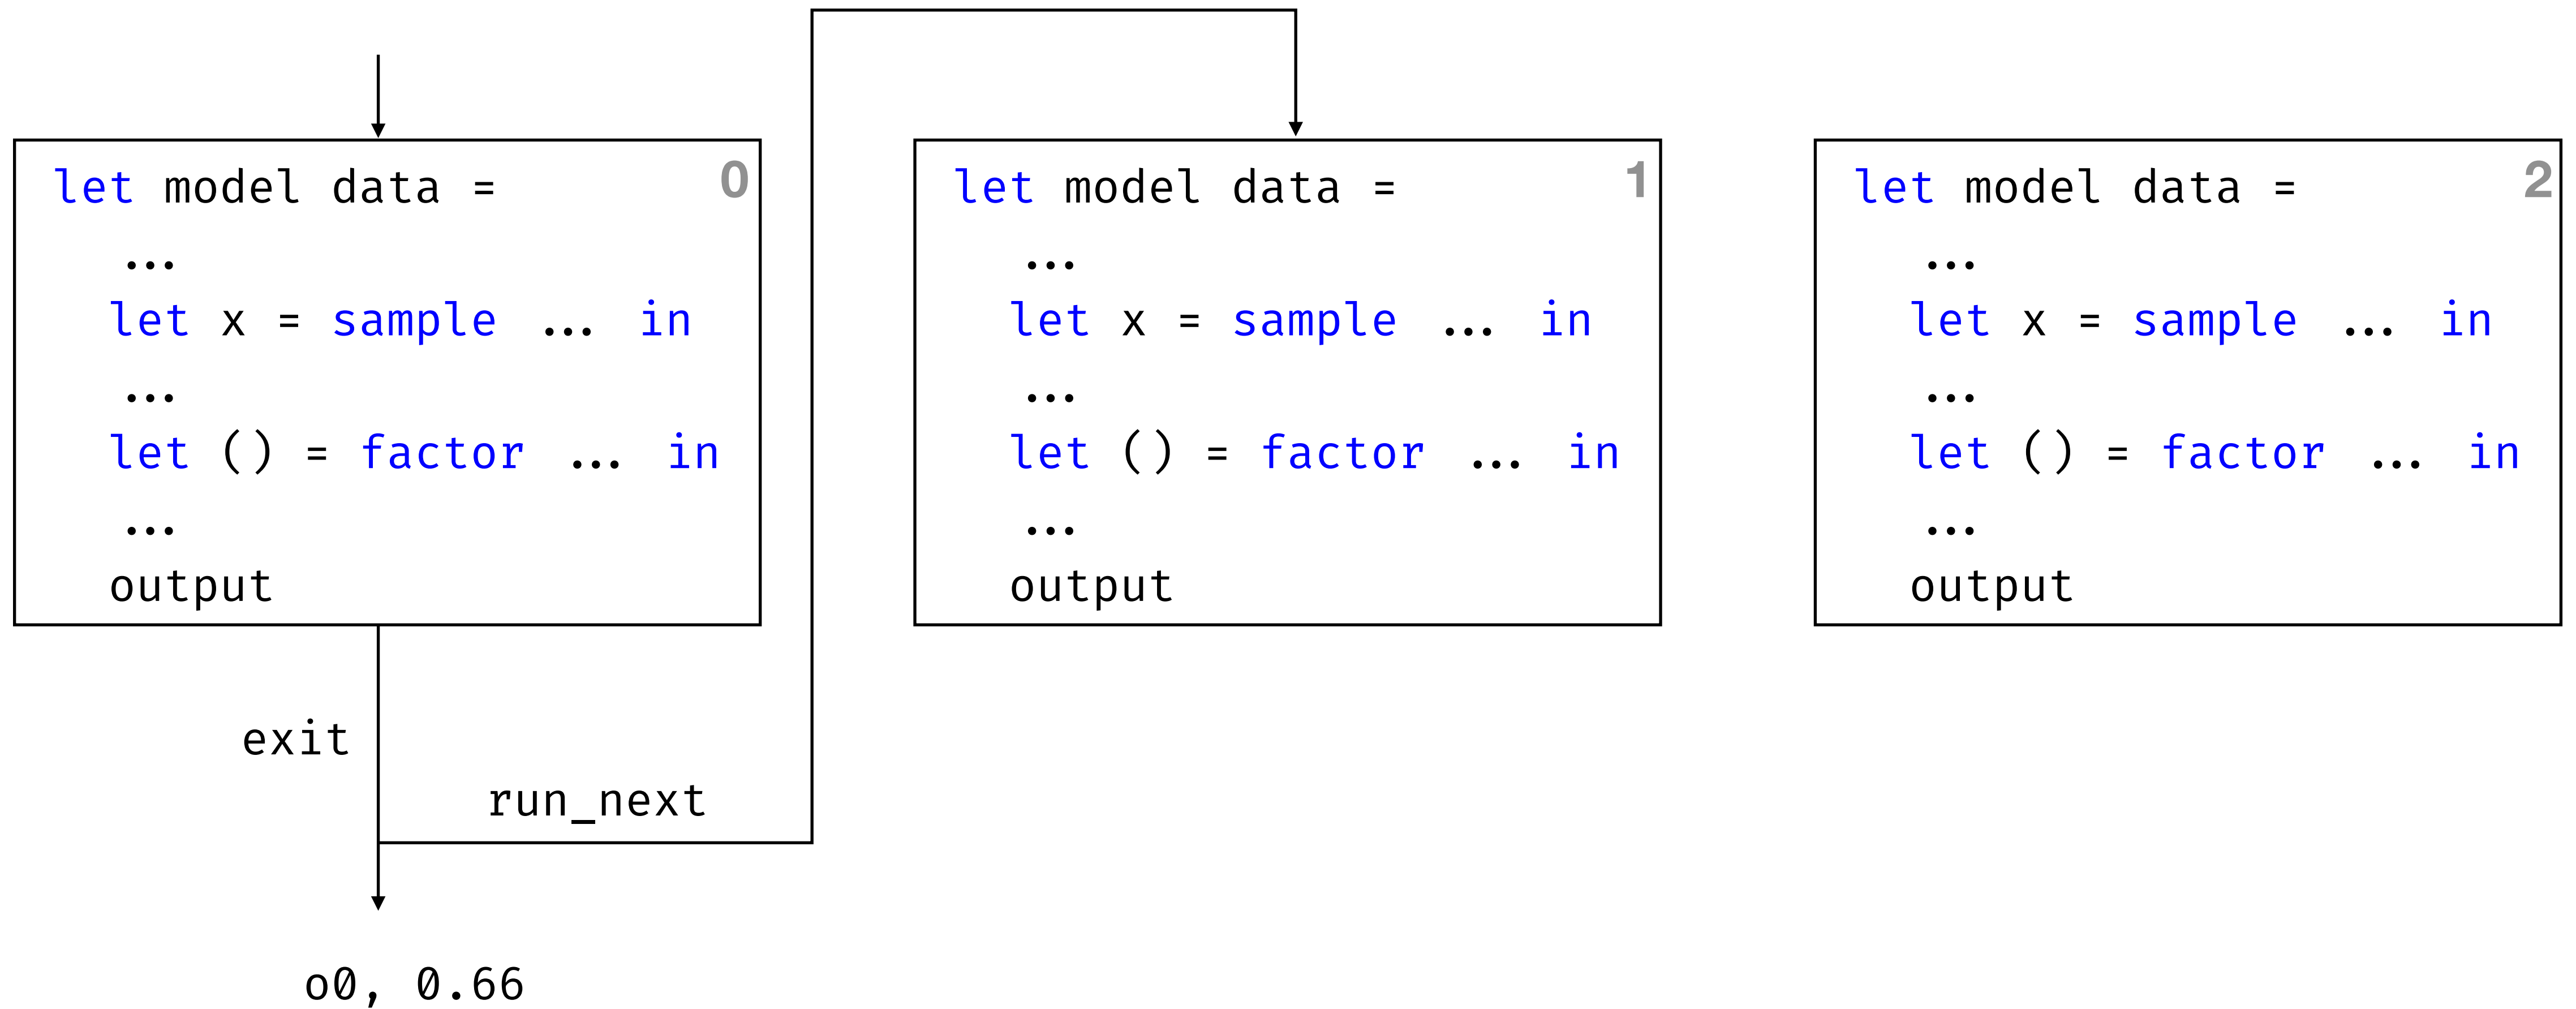
```
let model data = 1  
  ...  
  let x = sample ... in  
  ...  
  let () = factor ... in  
  ...  
  output
```

```
let model data = 2  
  ...  
  let x = sample ... in  
  ...  
  let () = factor ... in  
  ...  
  output
```

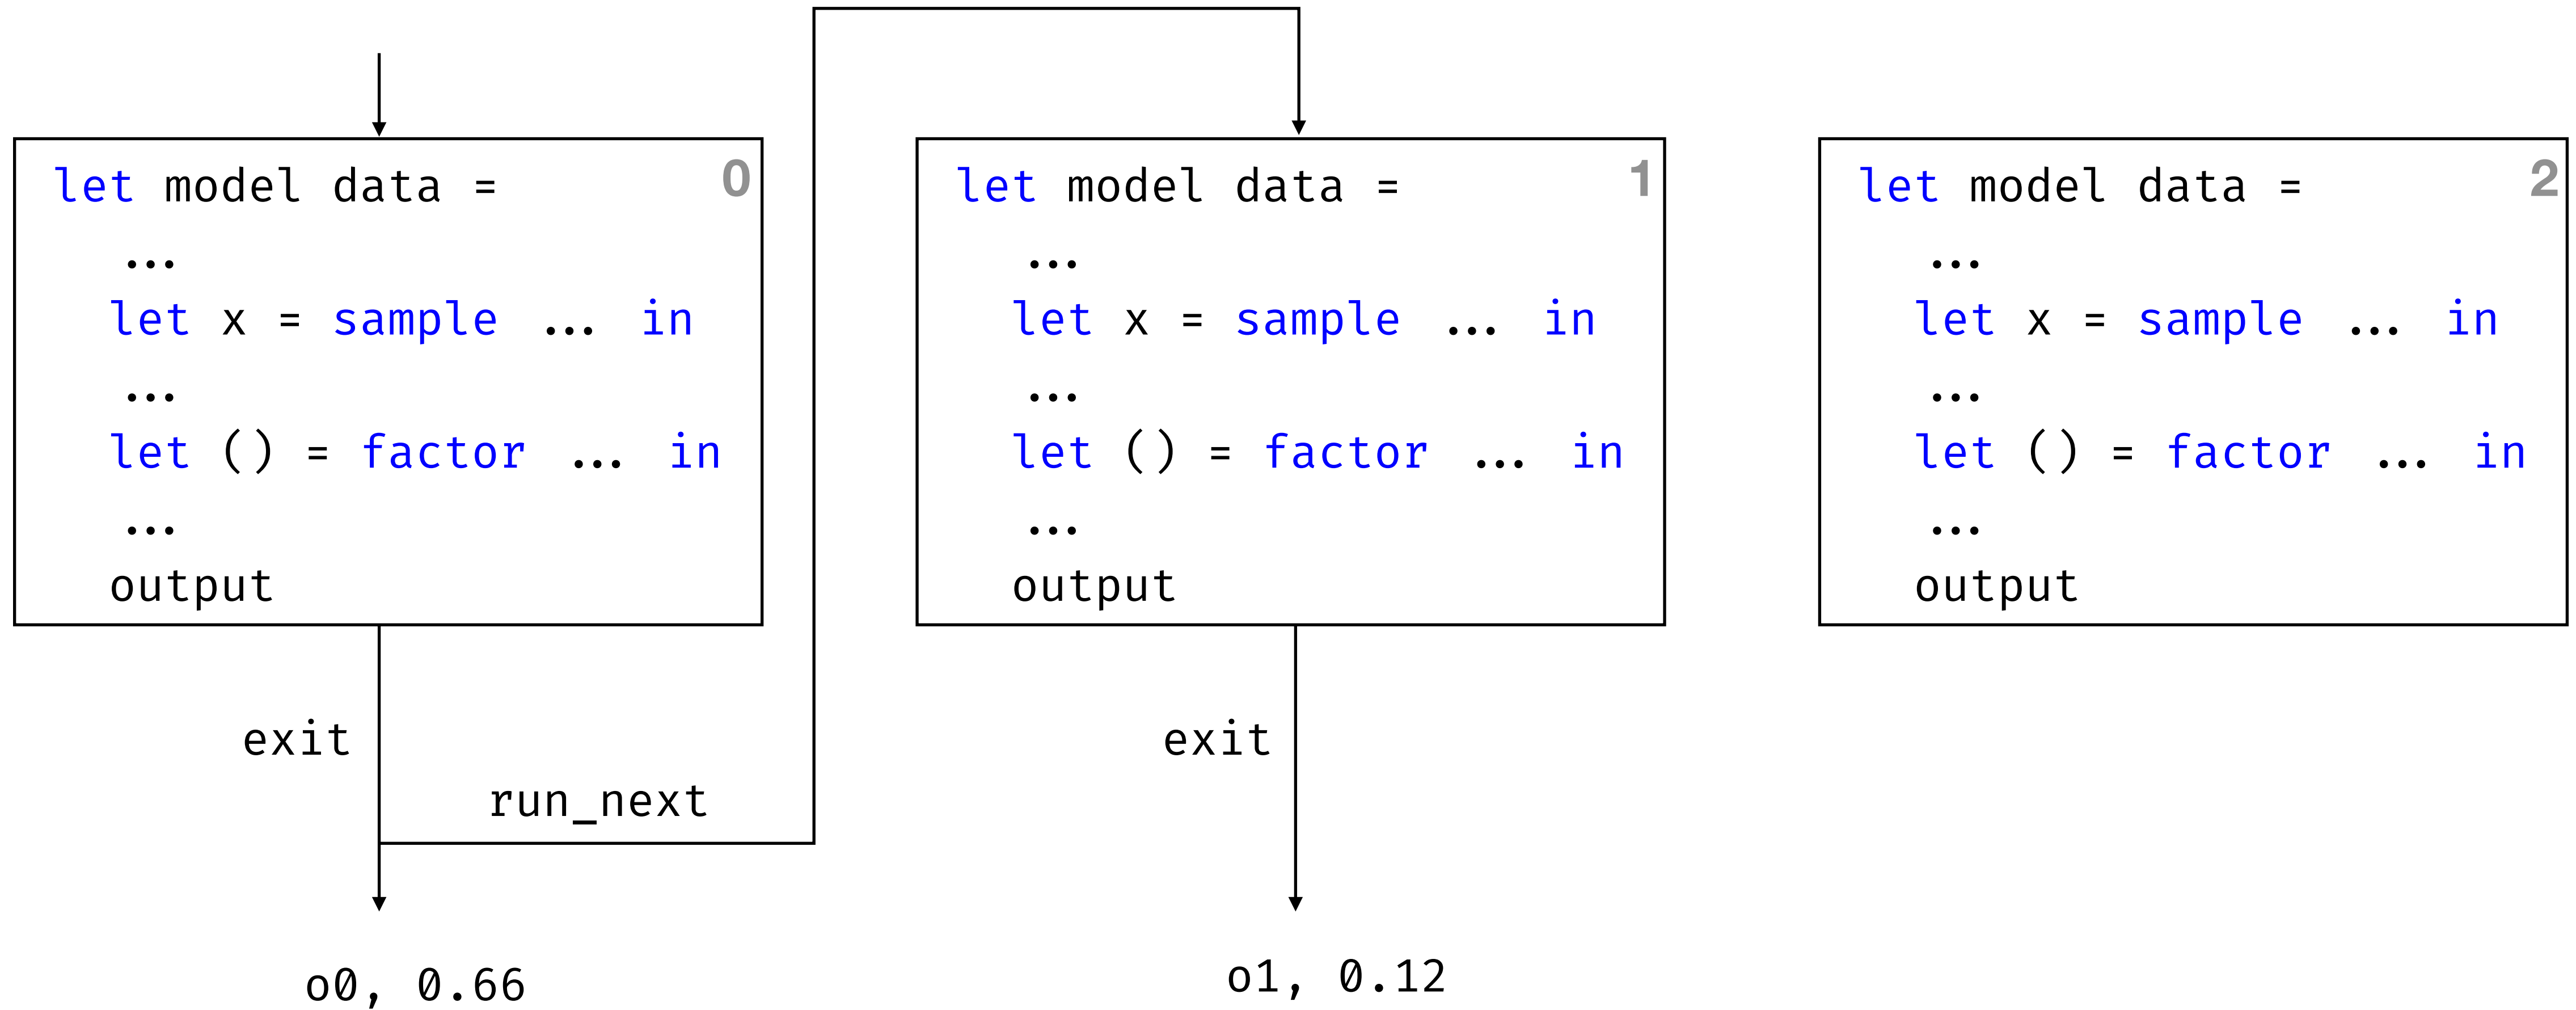
Importance Sampling



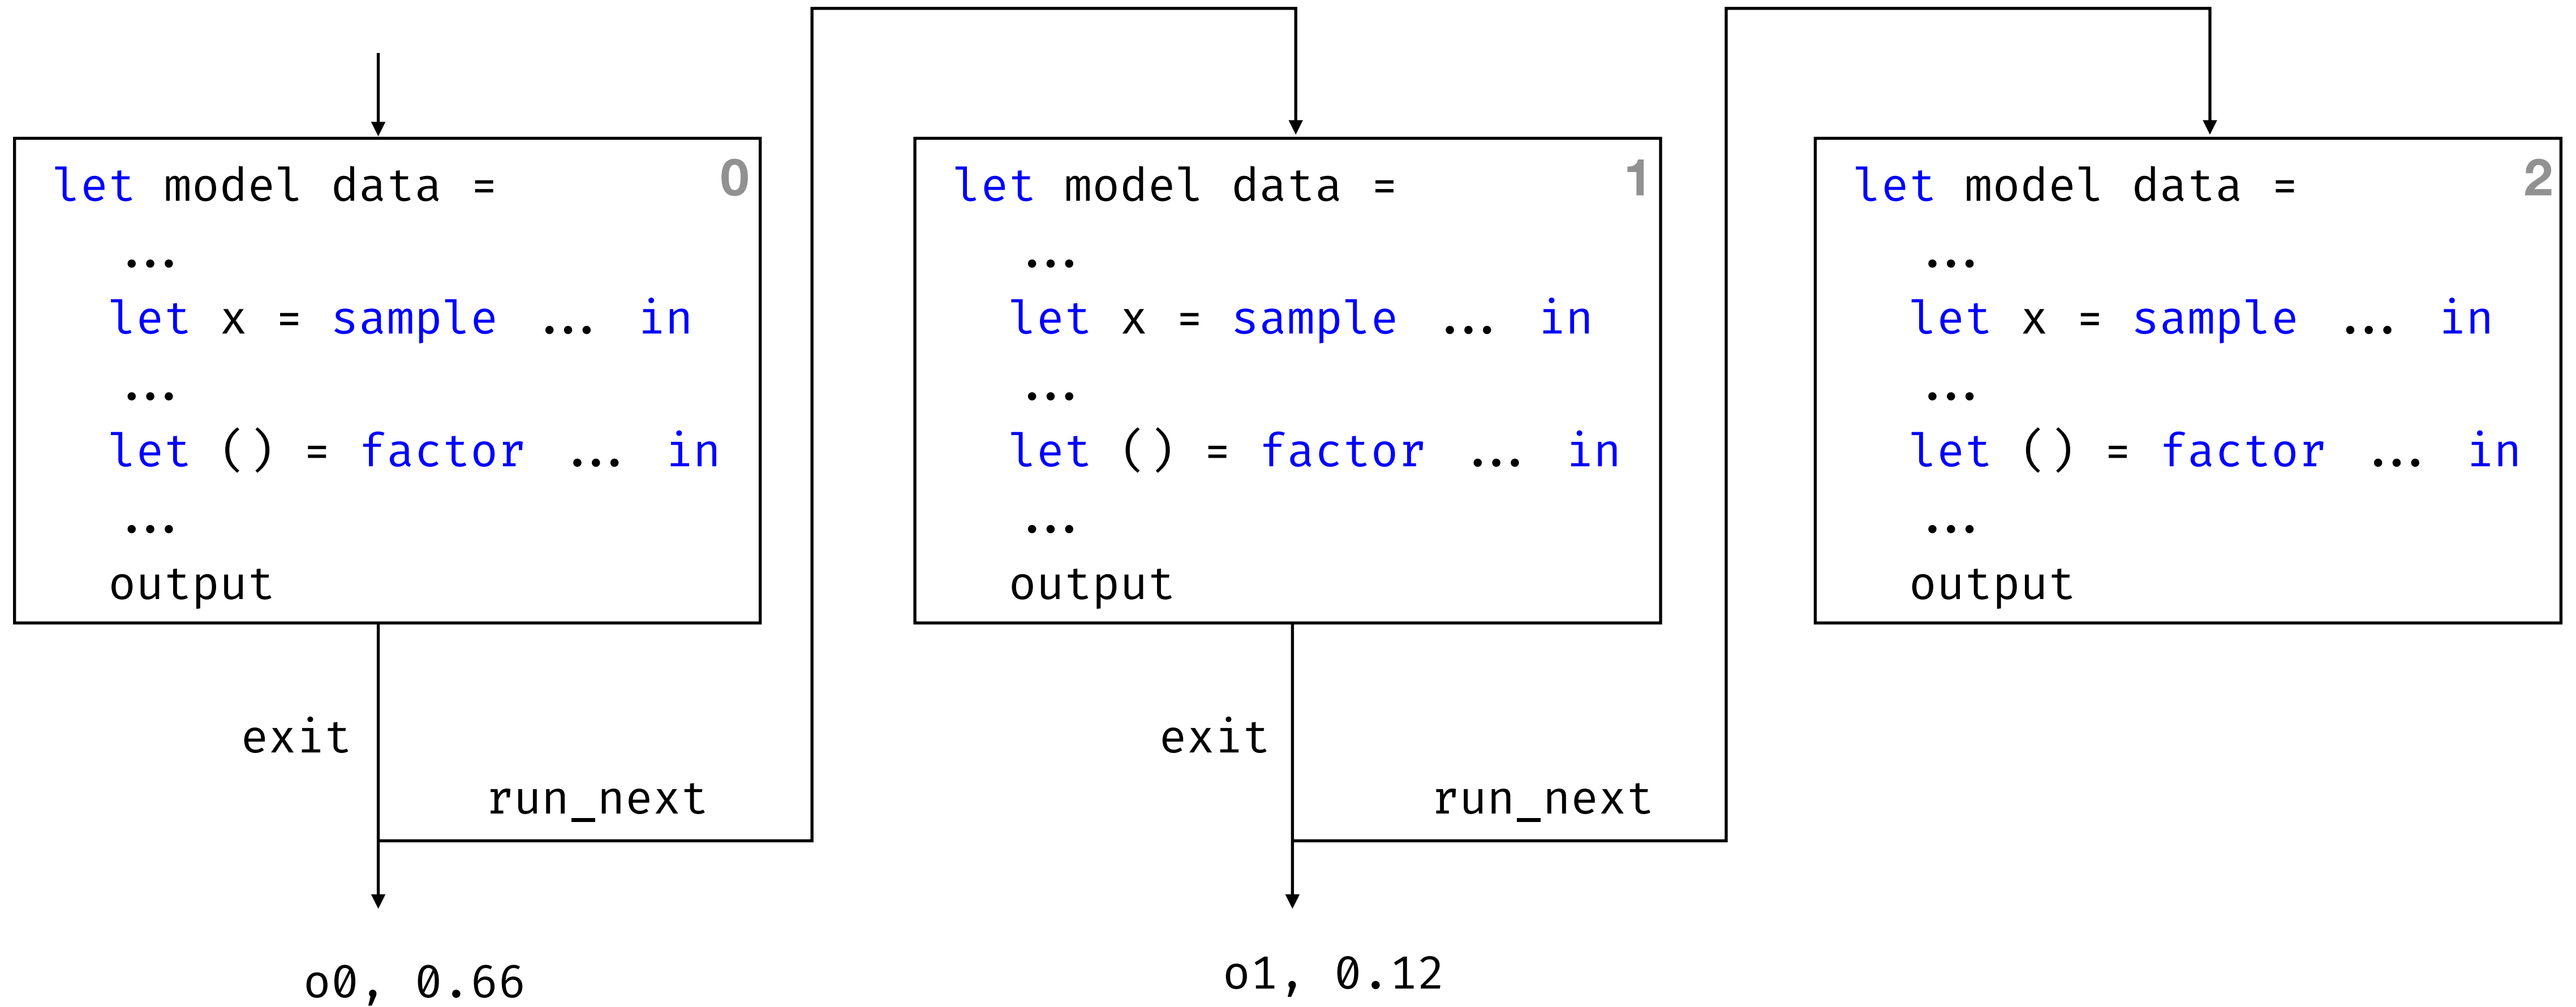
Importance Sampling



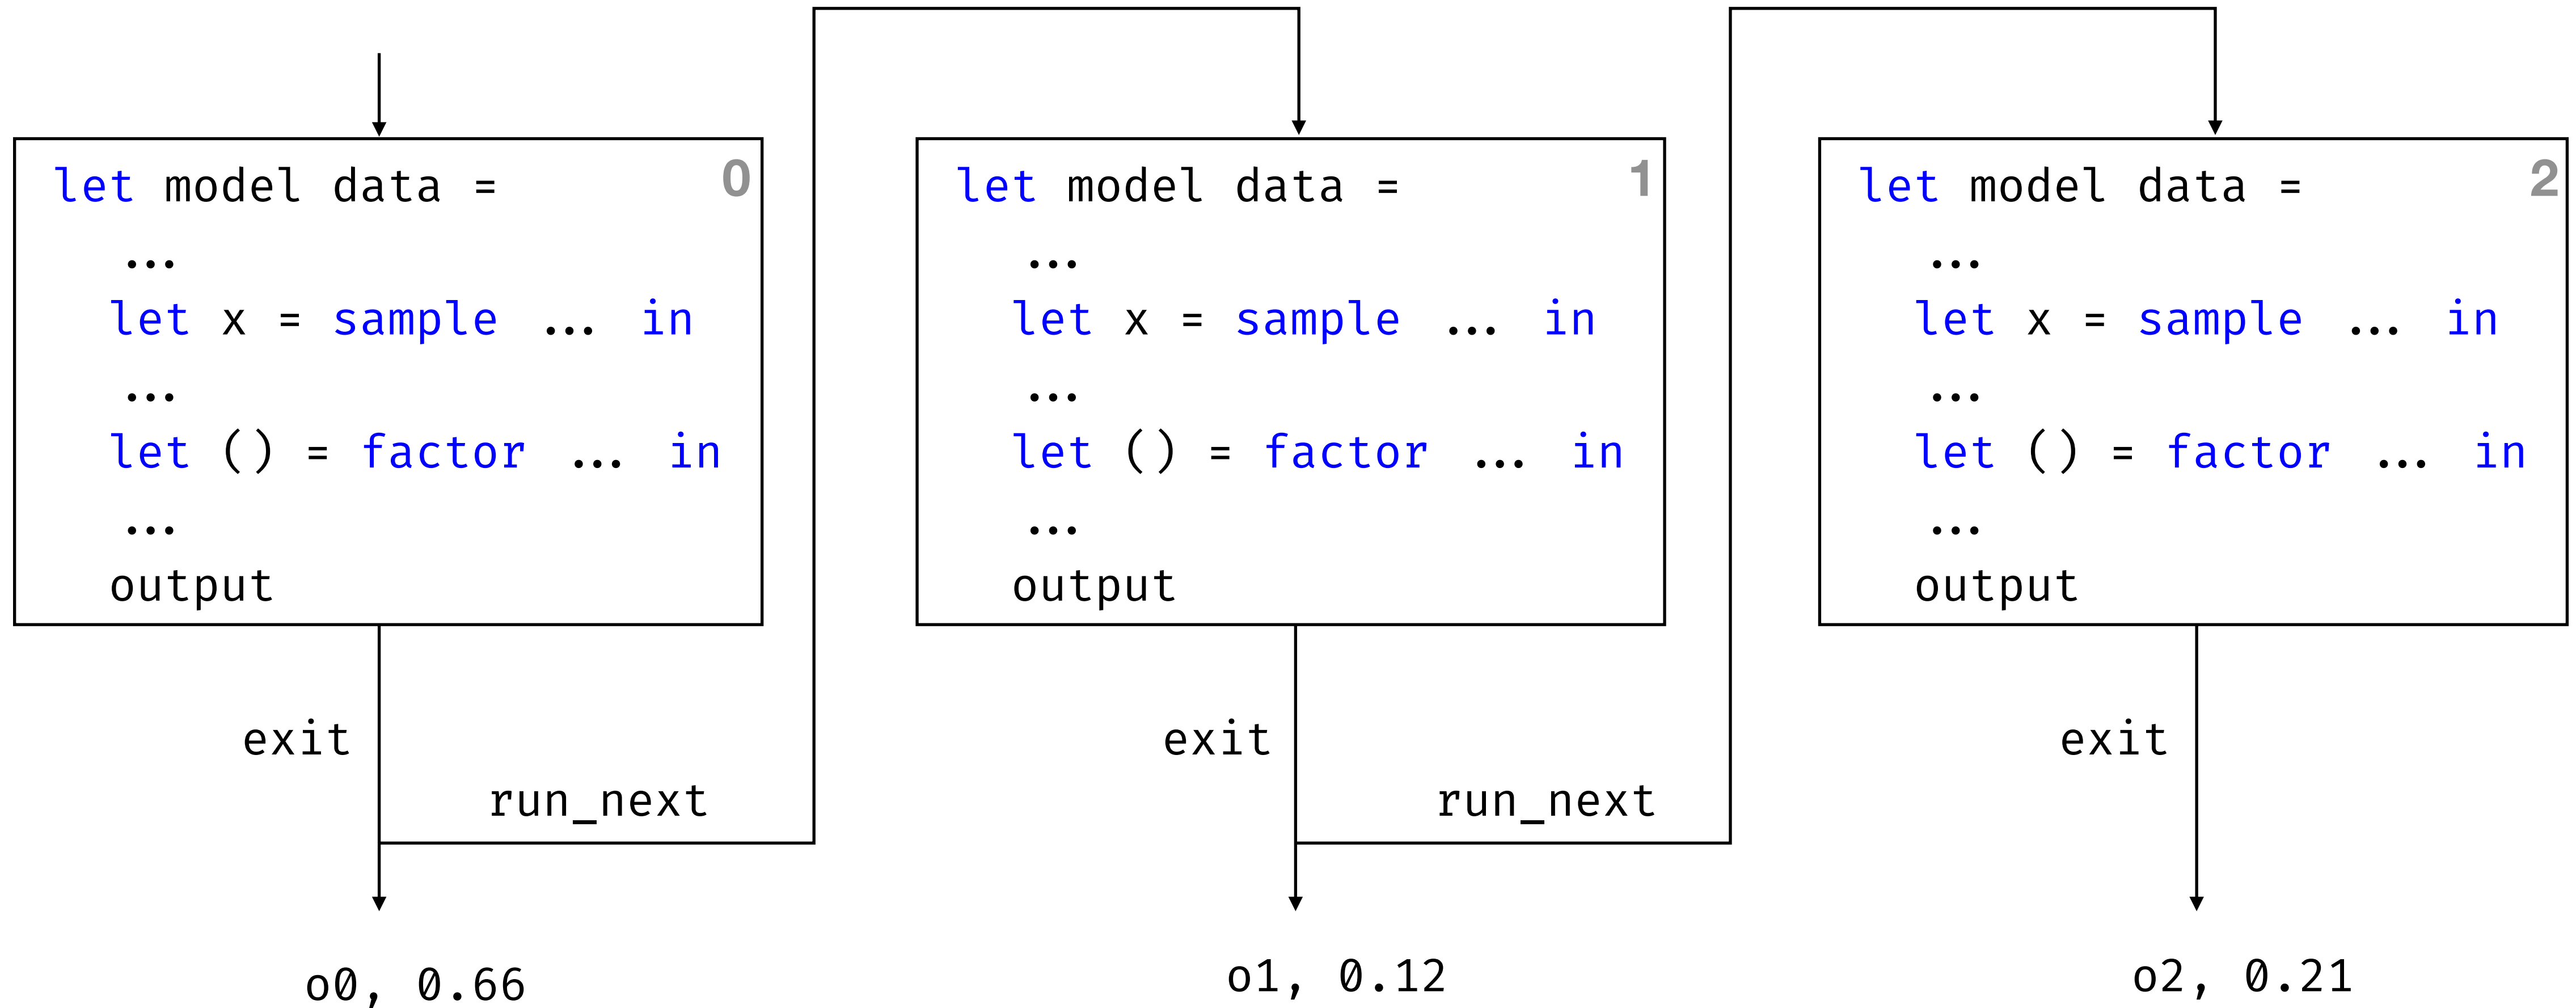
Importance Sampling



Importance Sampling



Importance Sampling



Importance Sampling

infer.ml

```
module Importance_sampling = struct
  type 'a prob = ...

  let sample d k prob = assert false
  let factor s k prob = assert false

  let infer ?(n = 1000) m data = assert false
end
```

Importance Sampling

infer.ml

```
module Importance_sampling = struct
  type 'a prob = { id : int; particles : 'a particle array }
  and 'a particle = { k : 'a next; value : 'a option; score : float }
  ...

  let sample d k prob =
    let v = Distribution.draw d in
    k v prob

  let factor s k prob =
    let particle = prob.particles.(prob.id) in
    prob.particles.(prob.id) ← { particle with score = s +. particle.score };
    k () prob
  ...
end
```

Importance Sampling

infer.ml

```
module Importance_sampling = struct
  type 'a prob = { id : int; particles : 'a particle array }
  and 'a particle = { k : 'a next; value : 'a option; score : float }
  ...

  (* Call the continuation of the next particle *)
  let run_next prob =
    if prob.id < Array.length prob.particles - 1 then
      let k = prob.particles.(prob.id + 1).k in
      k { prob with id = prob.id + 1 }
    else prob

  let exit v prob =
    let particle = prob.particles.(prob.id) in
    prob.particles.(prob.id) ← { particle with value = Some v };
    run_next prob
end
```

Importance Sampling

infer.ml

```
module Importance_sampling = struct
  type 'a prob = { id : int; particles : 'a particle array }
  and 'a particle = { k : 'a next; value : 'a option; score : float }
  ...

  let infer ?(n = 1000) m data =
    let init_particle = { k = (m data) exit; value = None; score = 0. } in
    let prob = { id = -1; particles = Array.init n (fun _ → init_particle) } in
    let prob = run_next prob in
    let values = Array.map (fun x → Option.get x.value) prob.particles in
    let logits = Array.map (fun x → x.score) prob.particles in
    Distribution.support ~values ~logits
end
```

Coin

coin.ml

```
open Infer.Importance_sampling
```

```
let coin x =  
  let* z = sample (uniform ~a:0. ~b:1.) in  
  let* () = Cps_list.iter (observe (bernoulli ~p:z)) x in  
  return z
```

```
let _ =  
  let dist = infer coin [1; 1; 0; 0; 0; 0; 0; 0; 0; 0] in  
  let m, s = Distribution.stats dist in  
  Format.printf "Coin bias, mean:%f, std:%f@." m s
```

```
> dune exec ./examples/coin.exe
```

```
Coin bias, mean:0.247876, std:0.118921  
Beta(2+1, 8+1), mean:0.250000, std:0.120096
```

Particle Filter (CPS)

BYO-PPL

Particle Filter

basic.ml

```
module Particle_filter = struct
  include Importance_sampling

  let resample particles = assert false
  let factor s k prob = assert false
end
```

Inference algorithm : importance sampling, but...

- Add a resampling step at each **factor**
- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution

Particle Filter

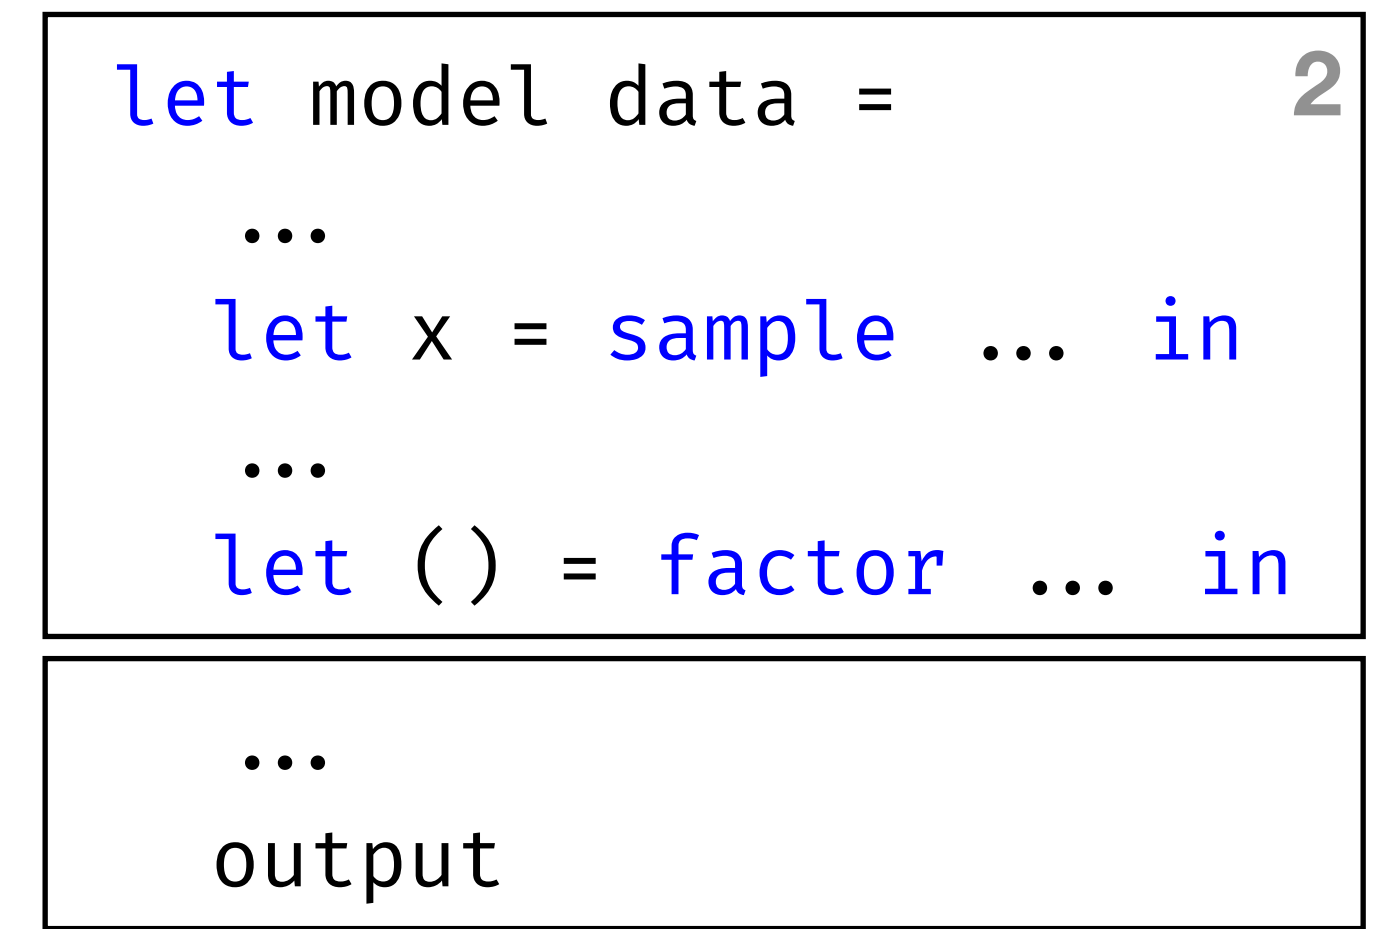
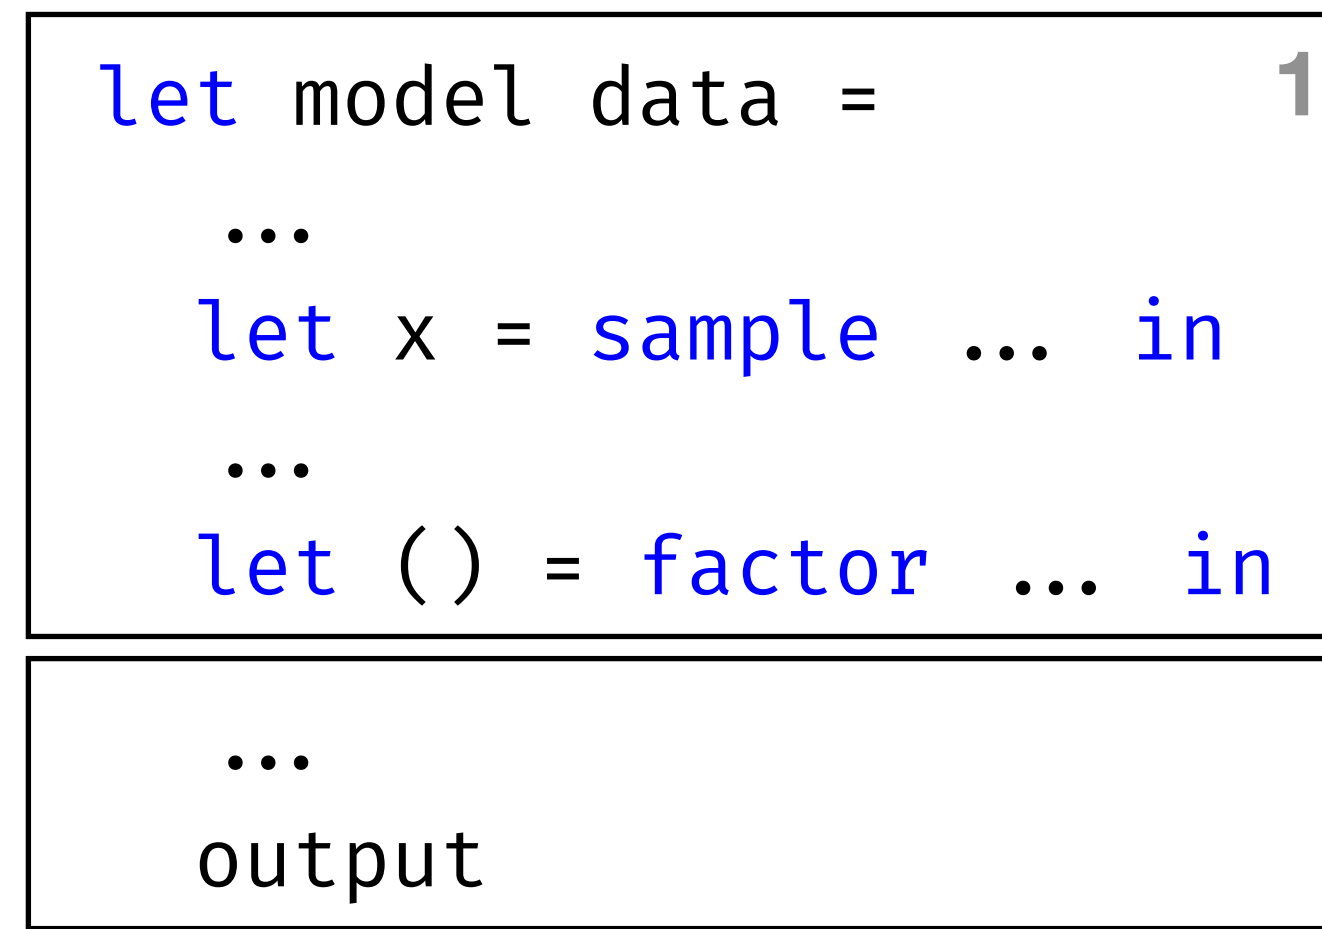
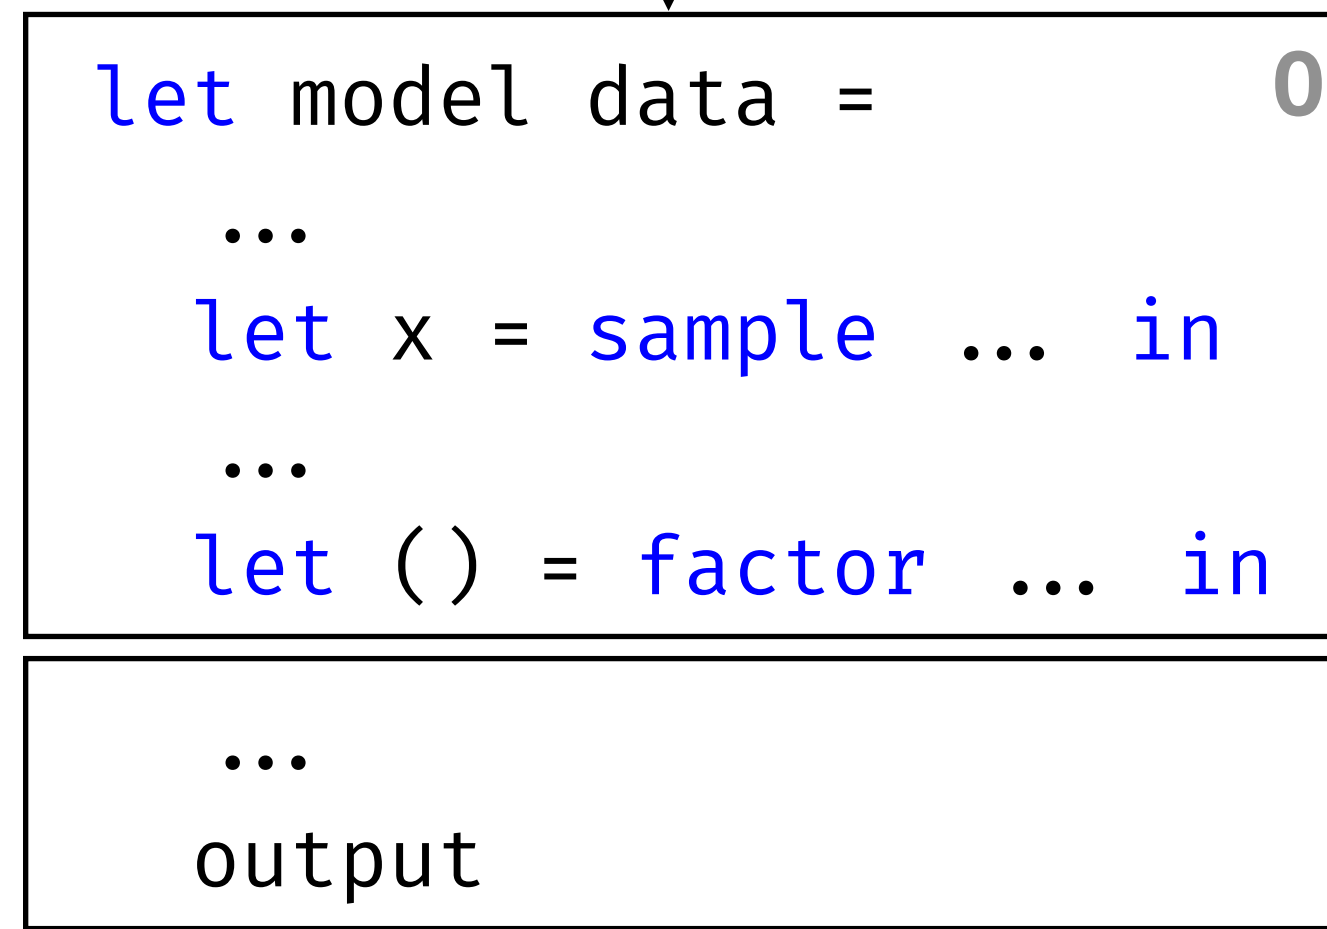


```
let model data = 0
  ...
  let x = sample ... in
  ...
  let () = factor ... in
  ...
  output
```

```
let model data = 1
  ...
  let x = sample ... in
  ...
  let () = factor ... in
  ...
  output
```

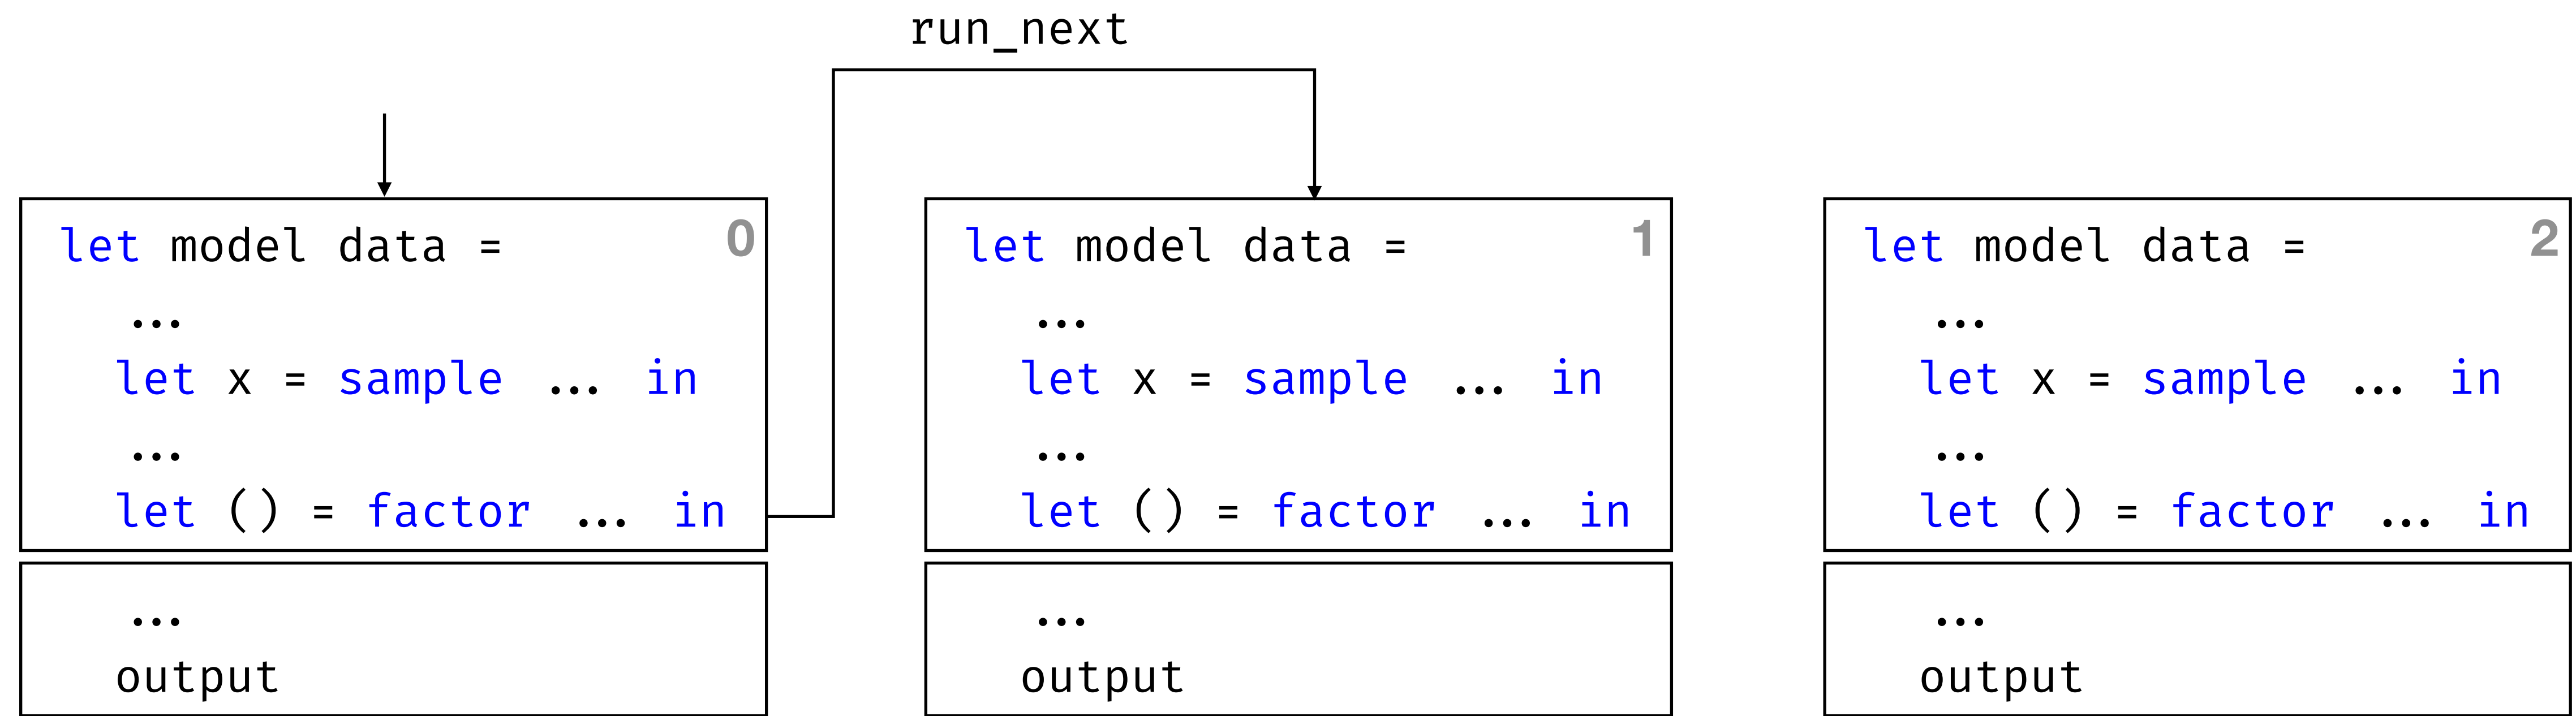
```
let model data = 2
  ...
  let x = sample ... in
  ...
  let () = factor ... in
  ...
  output
```

Particle Filter



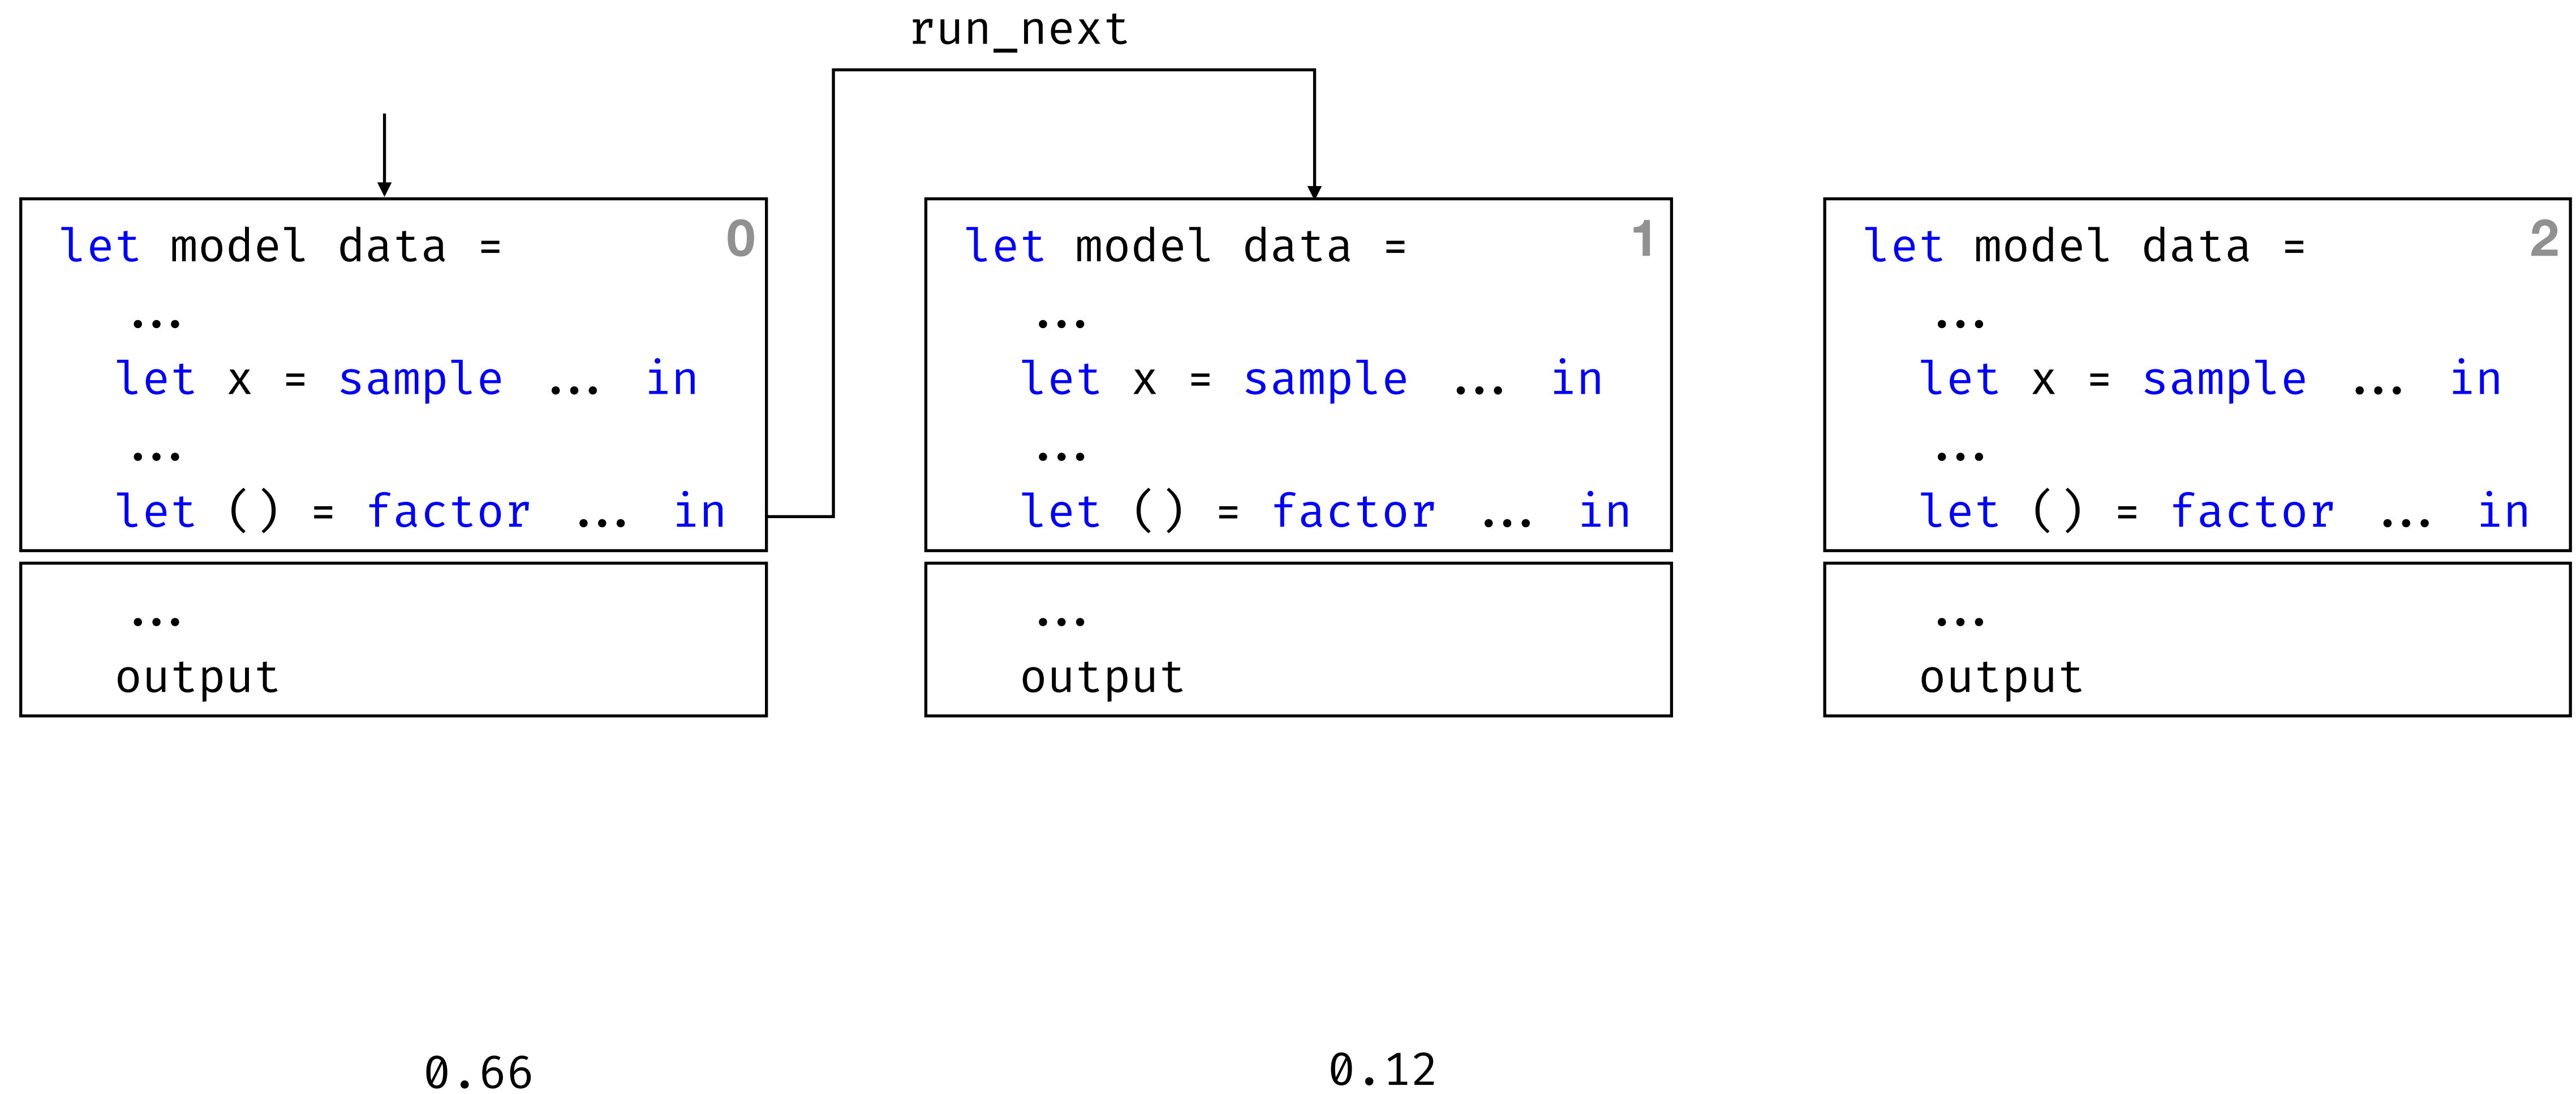
0.66

Particle Filter

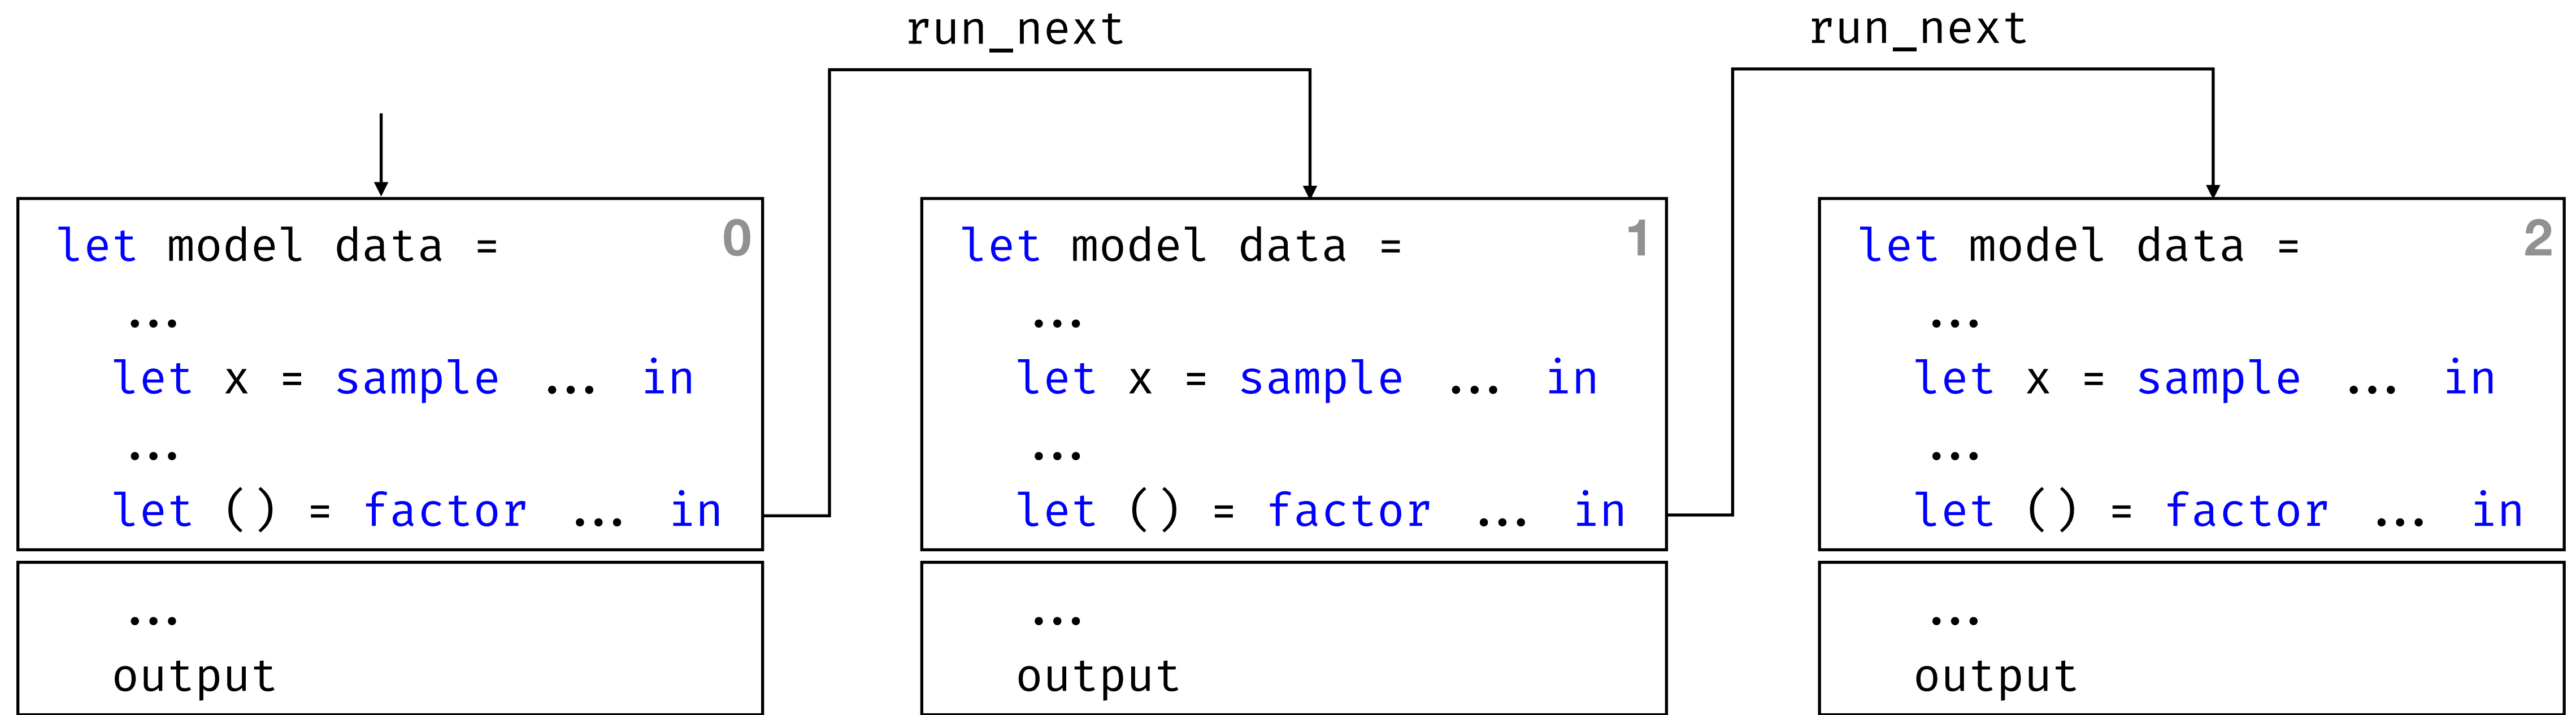


0.66

Particle Filter



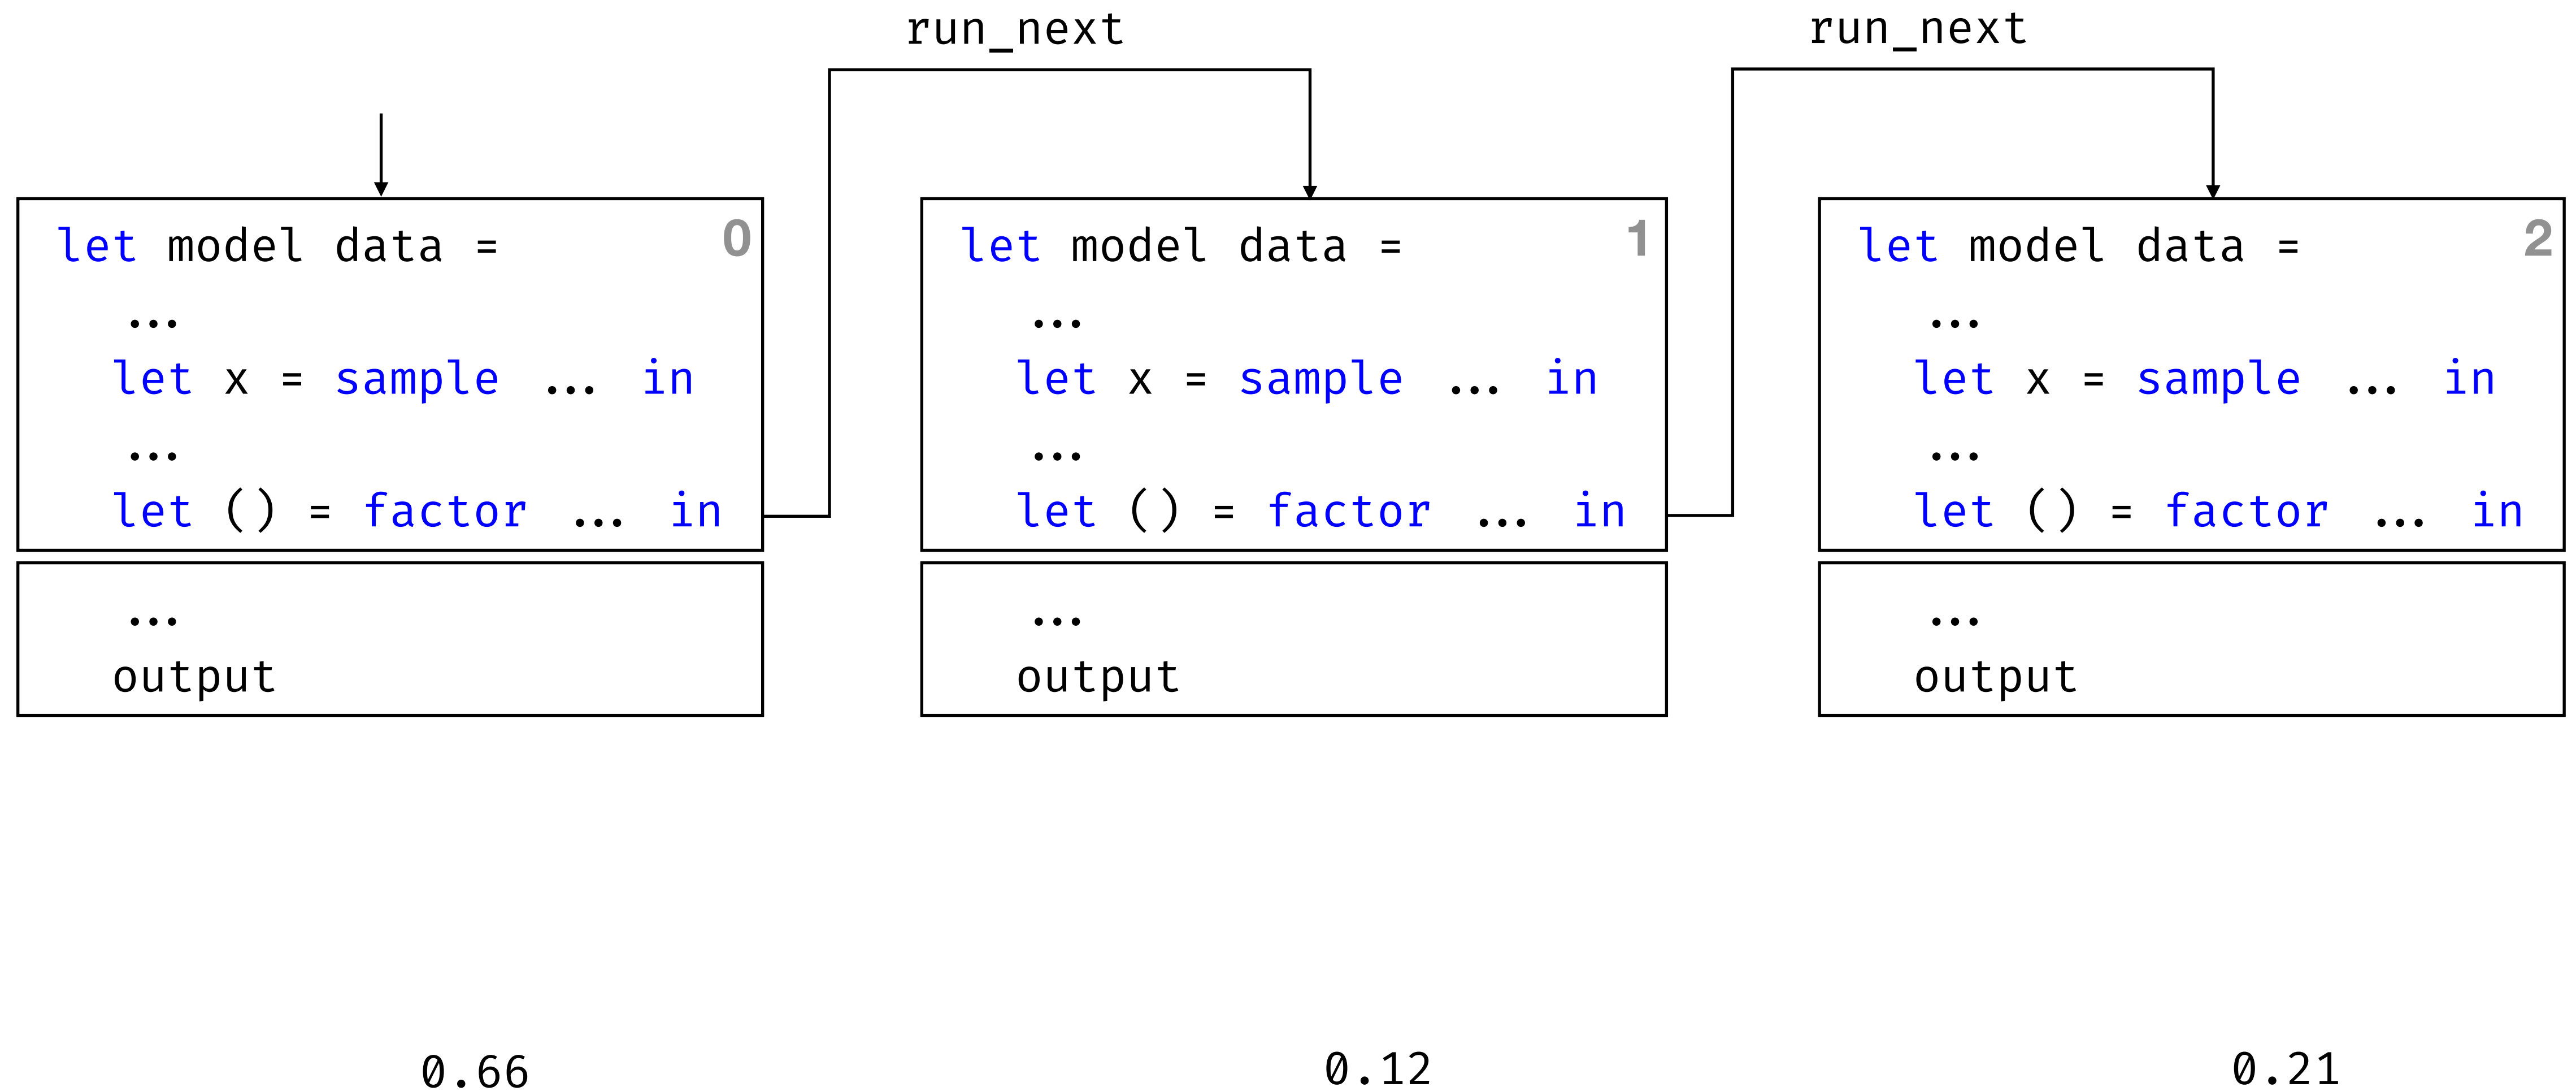
Particle Filter



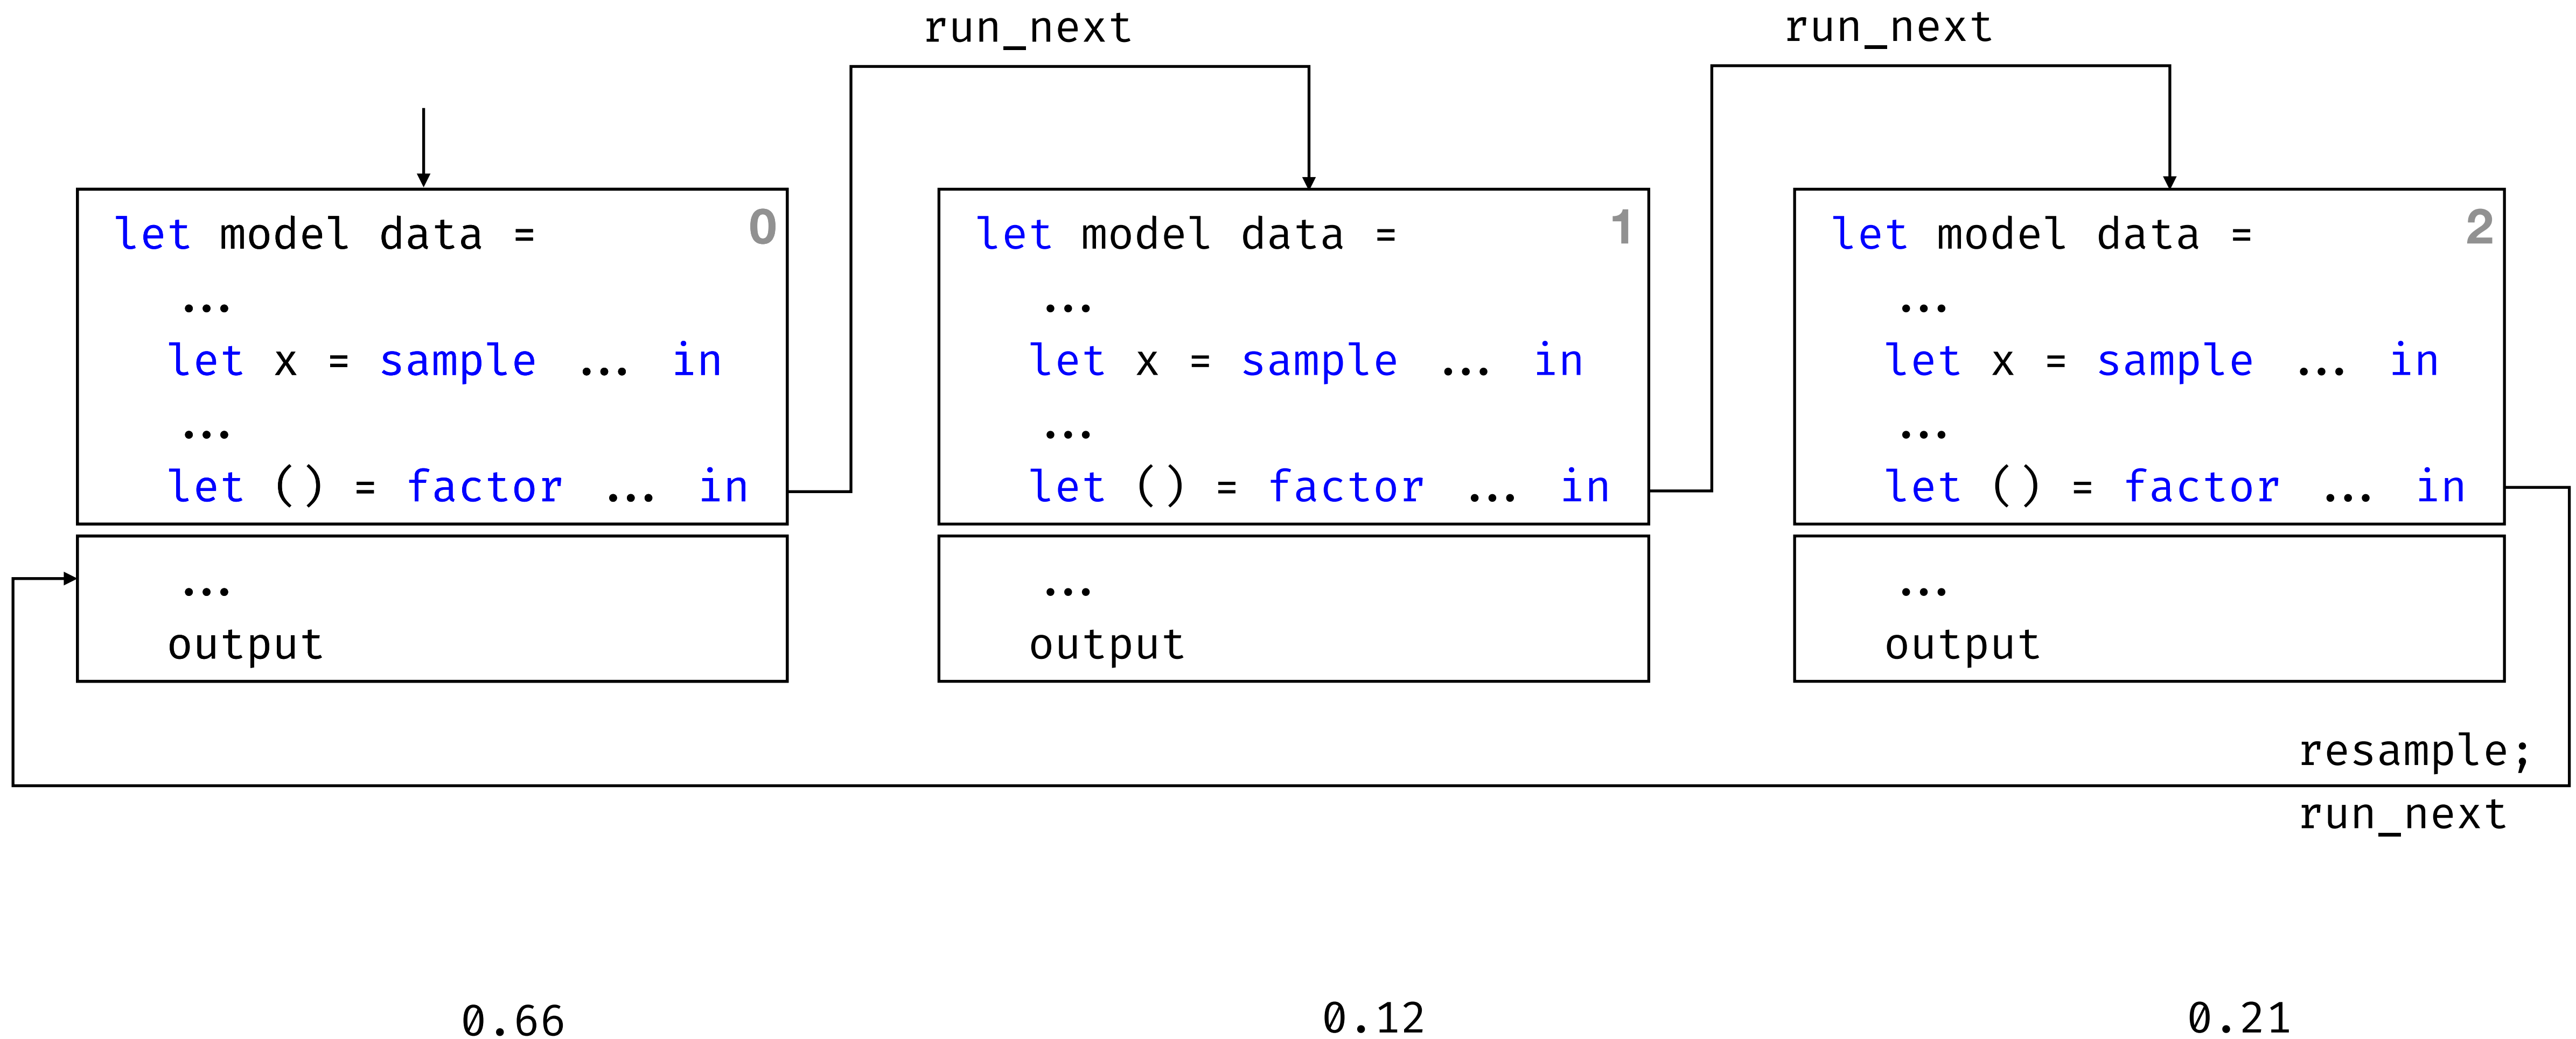
0.66

0.12

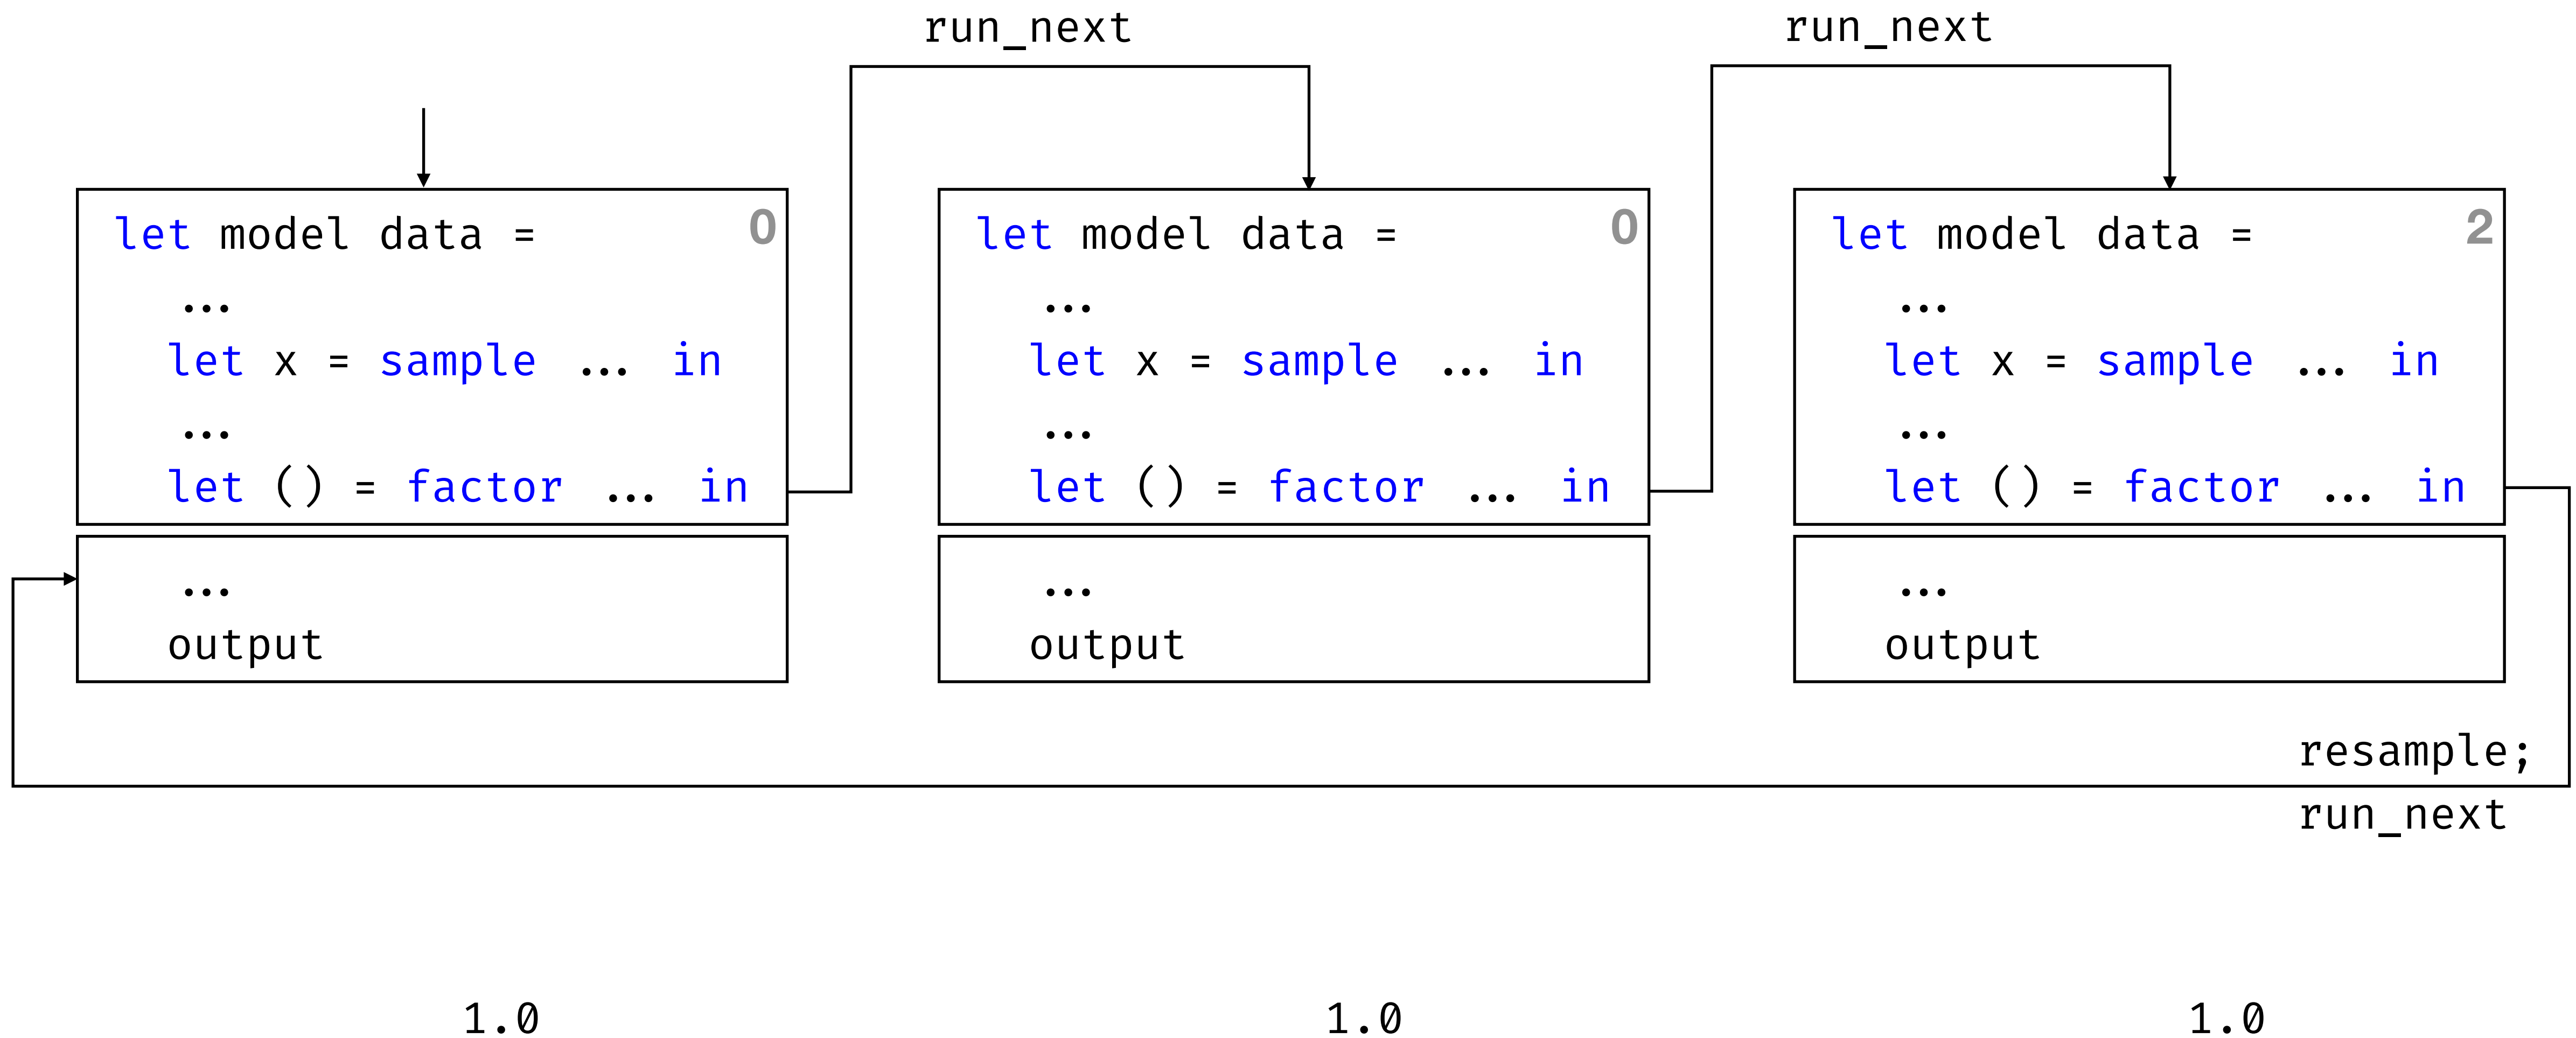
Particle Filter



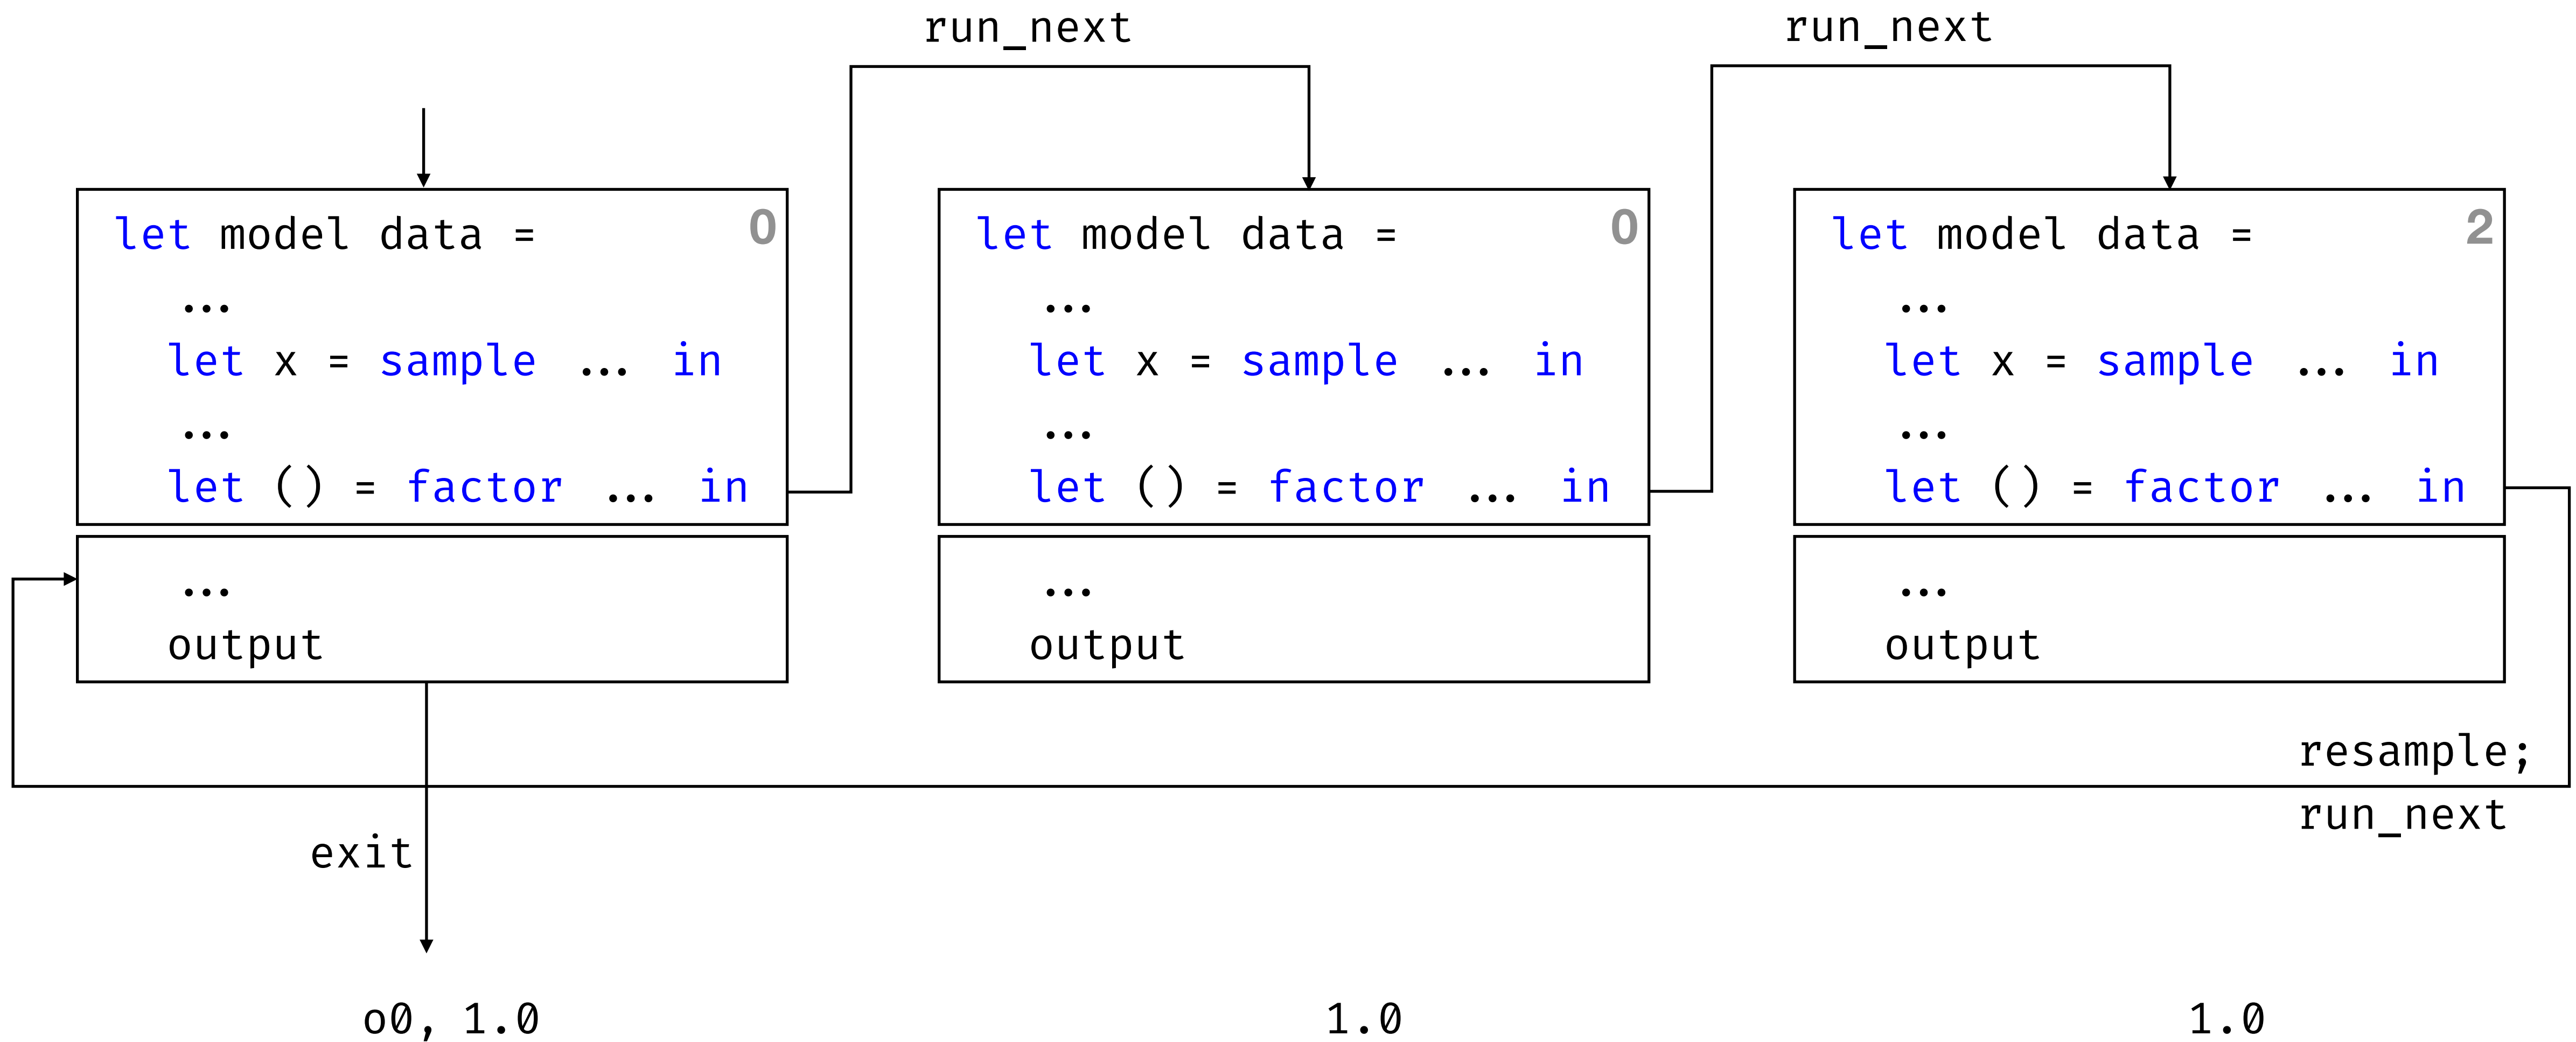
Particle Filter



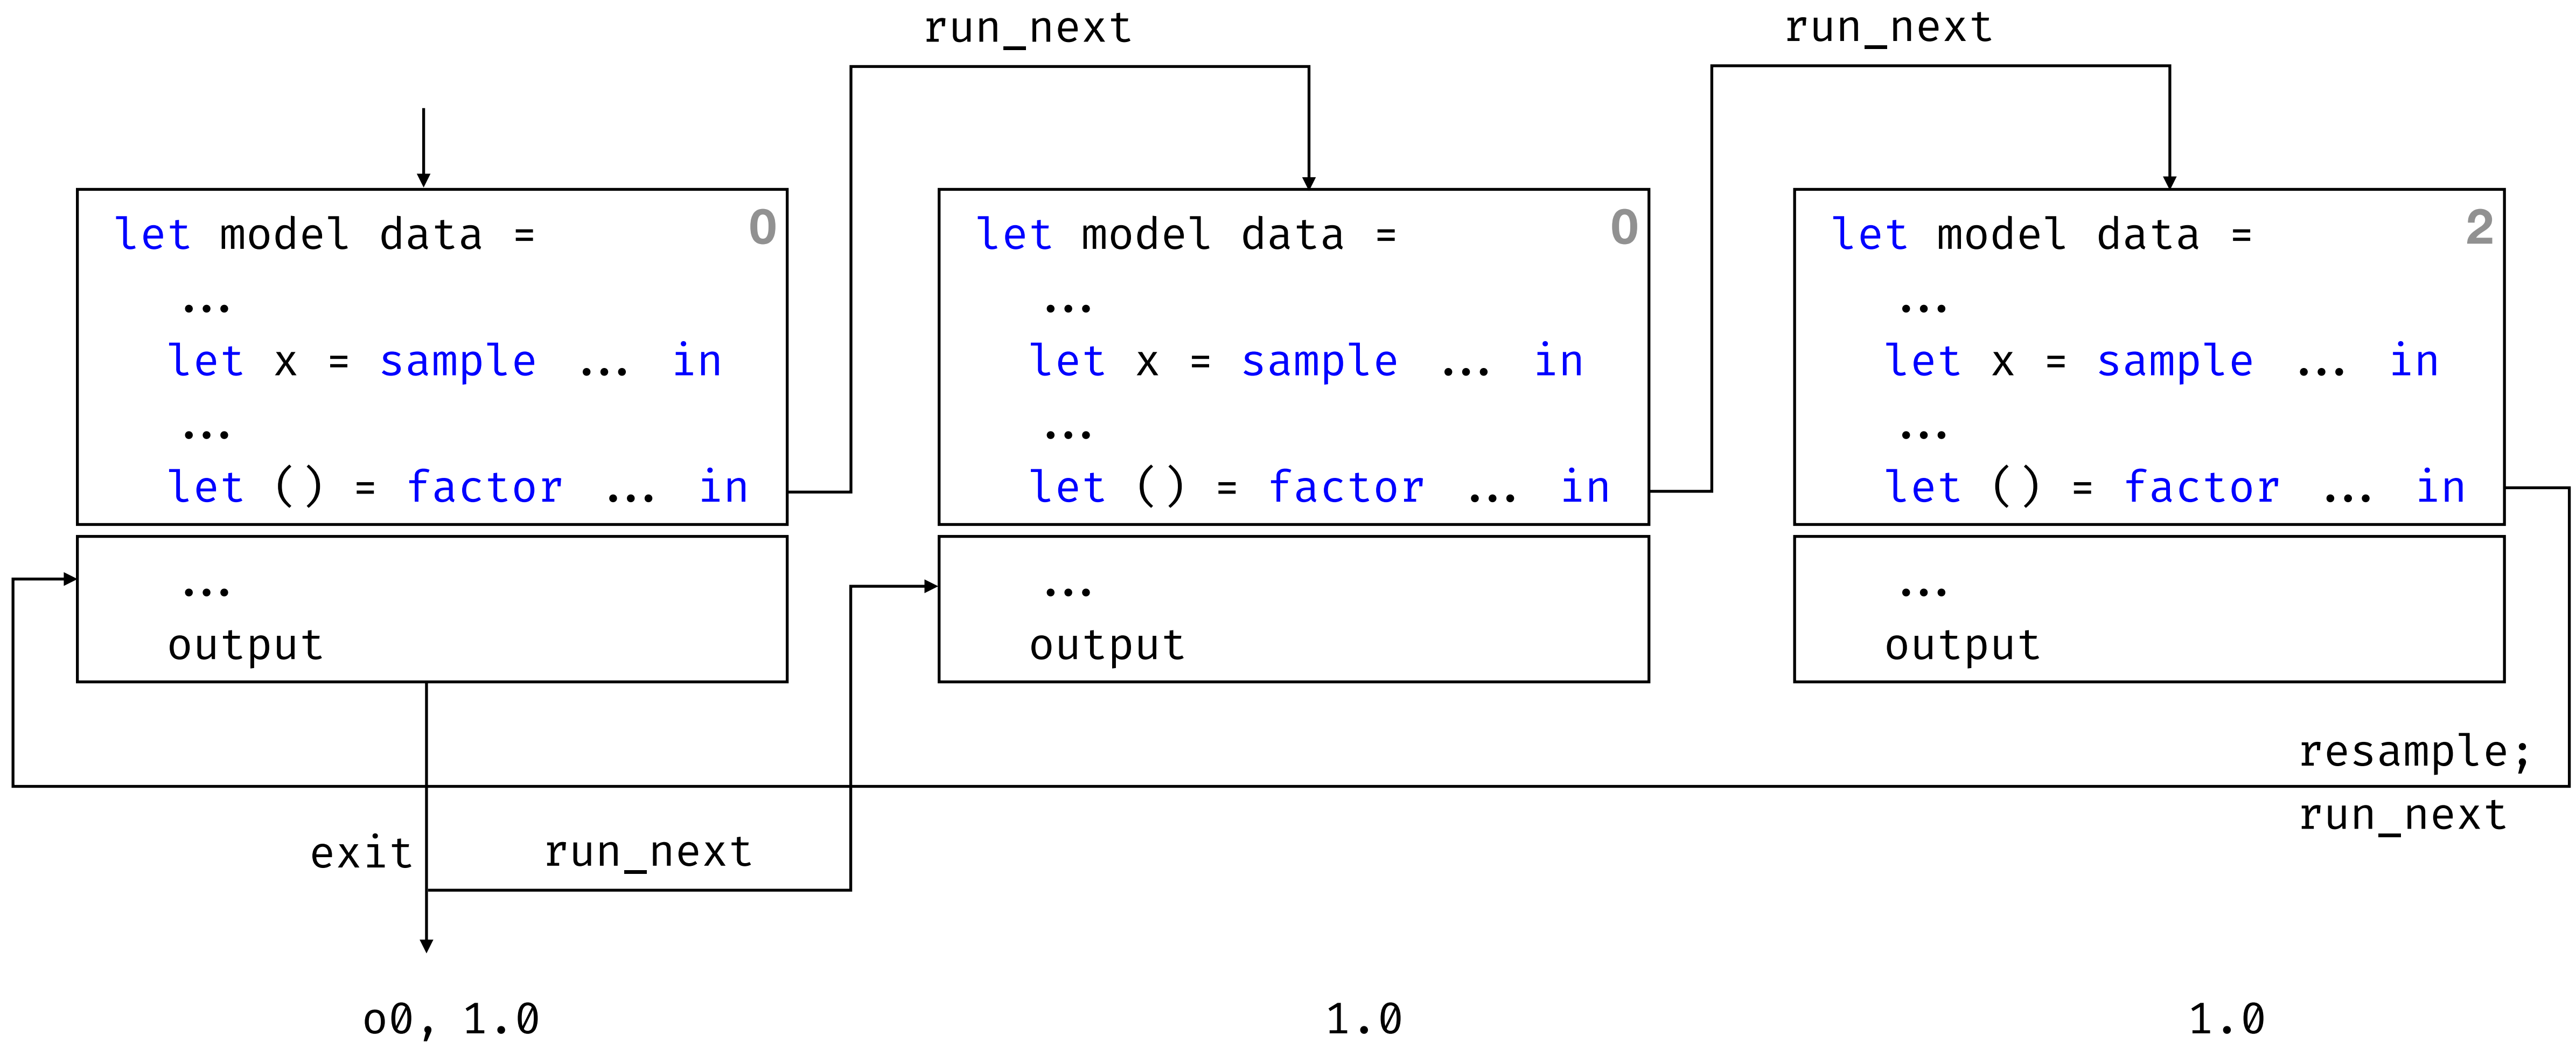
Particle Filter



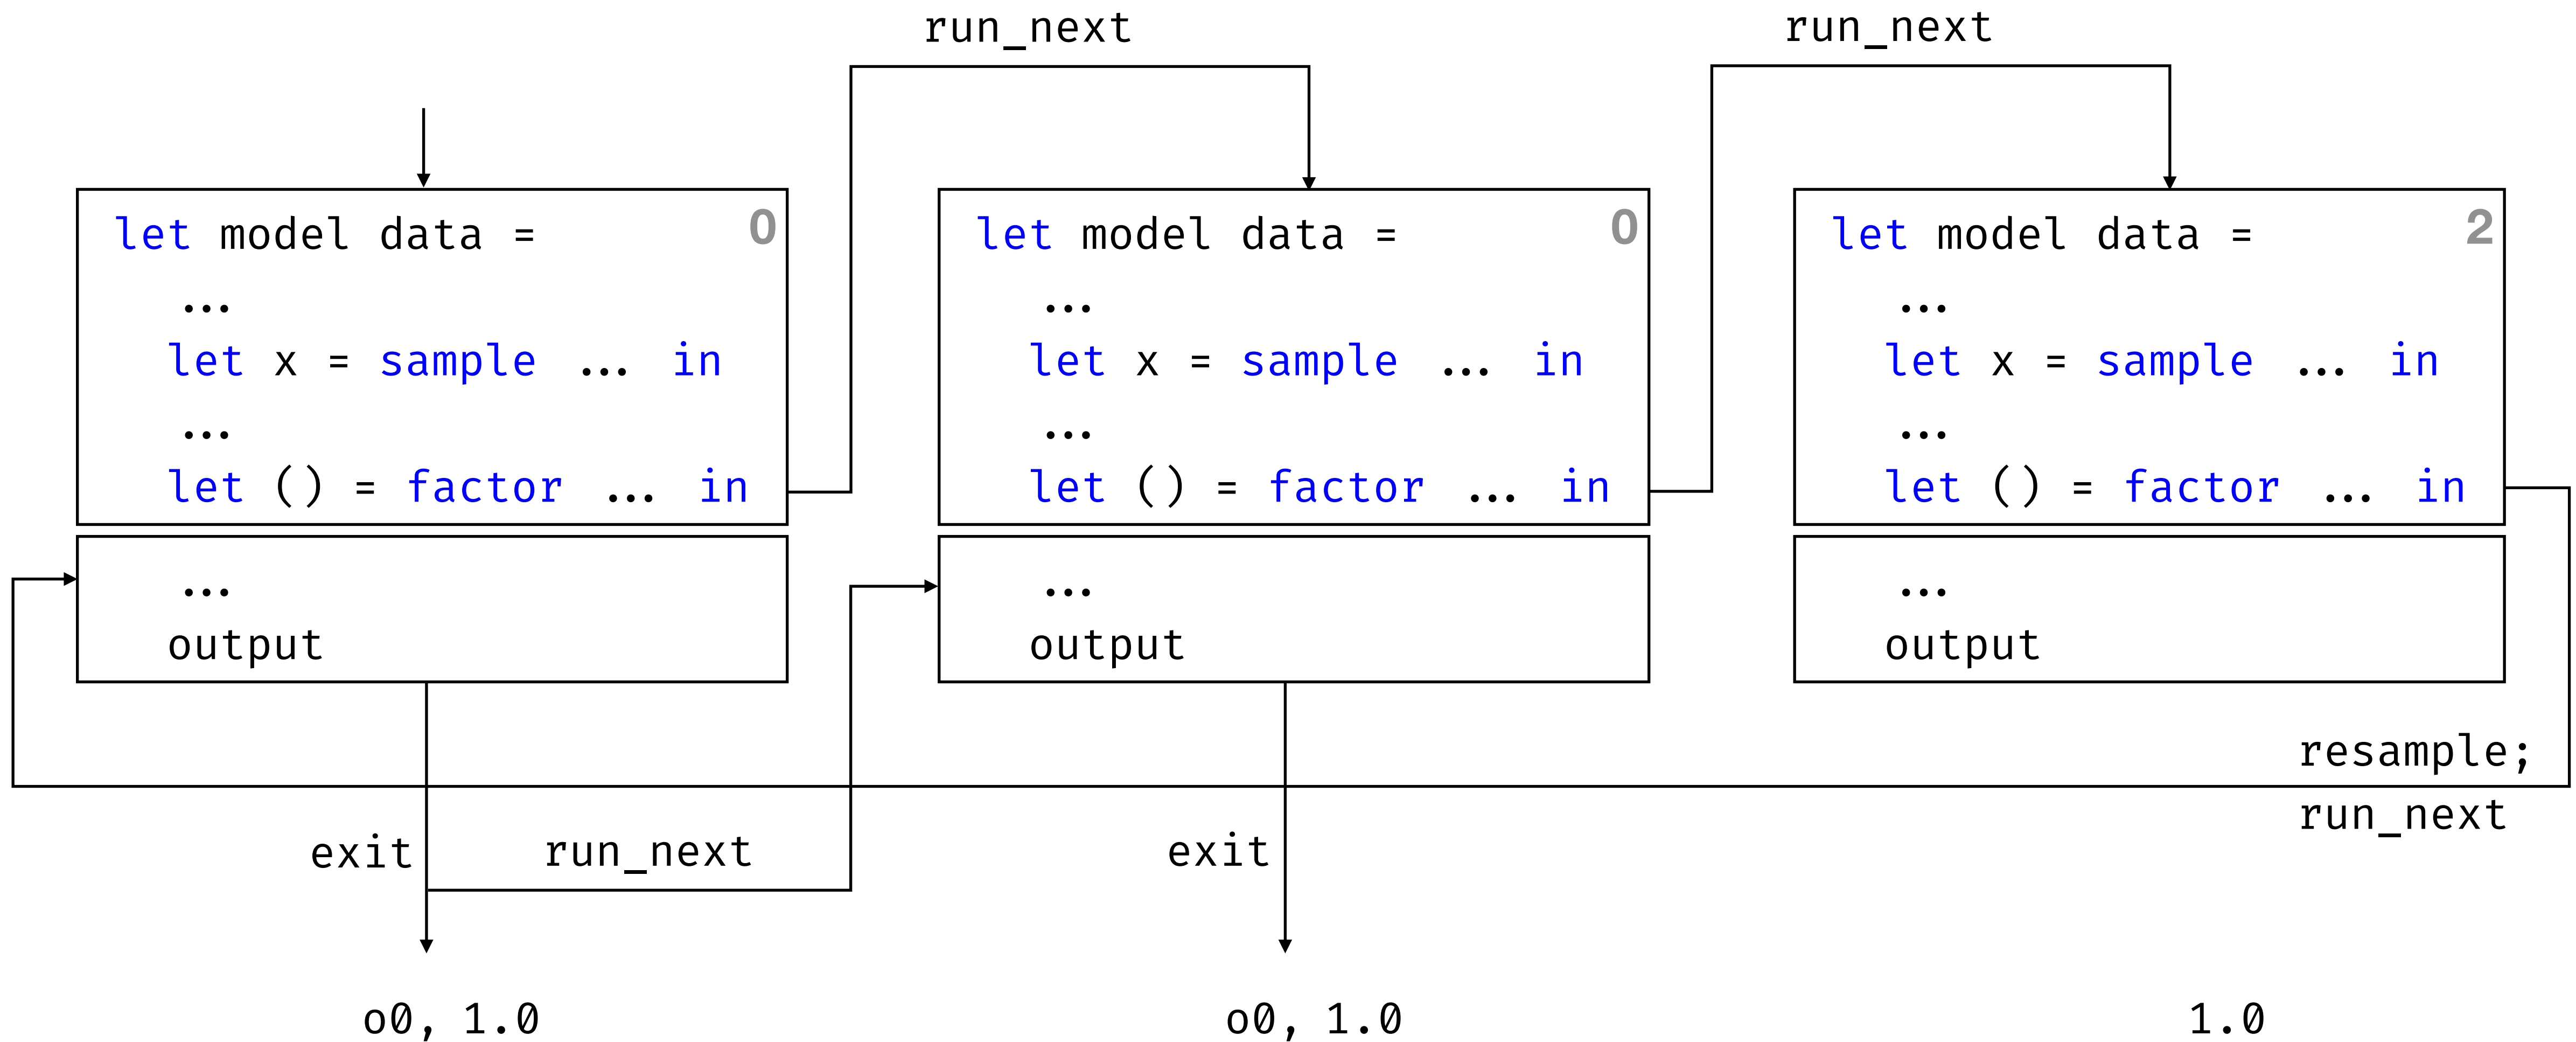
Particle Filter



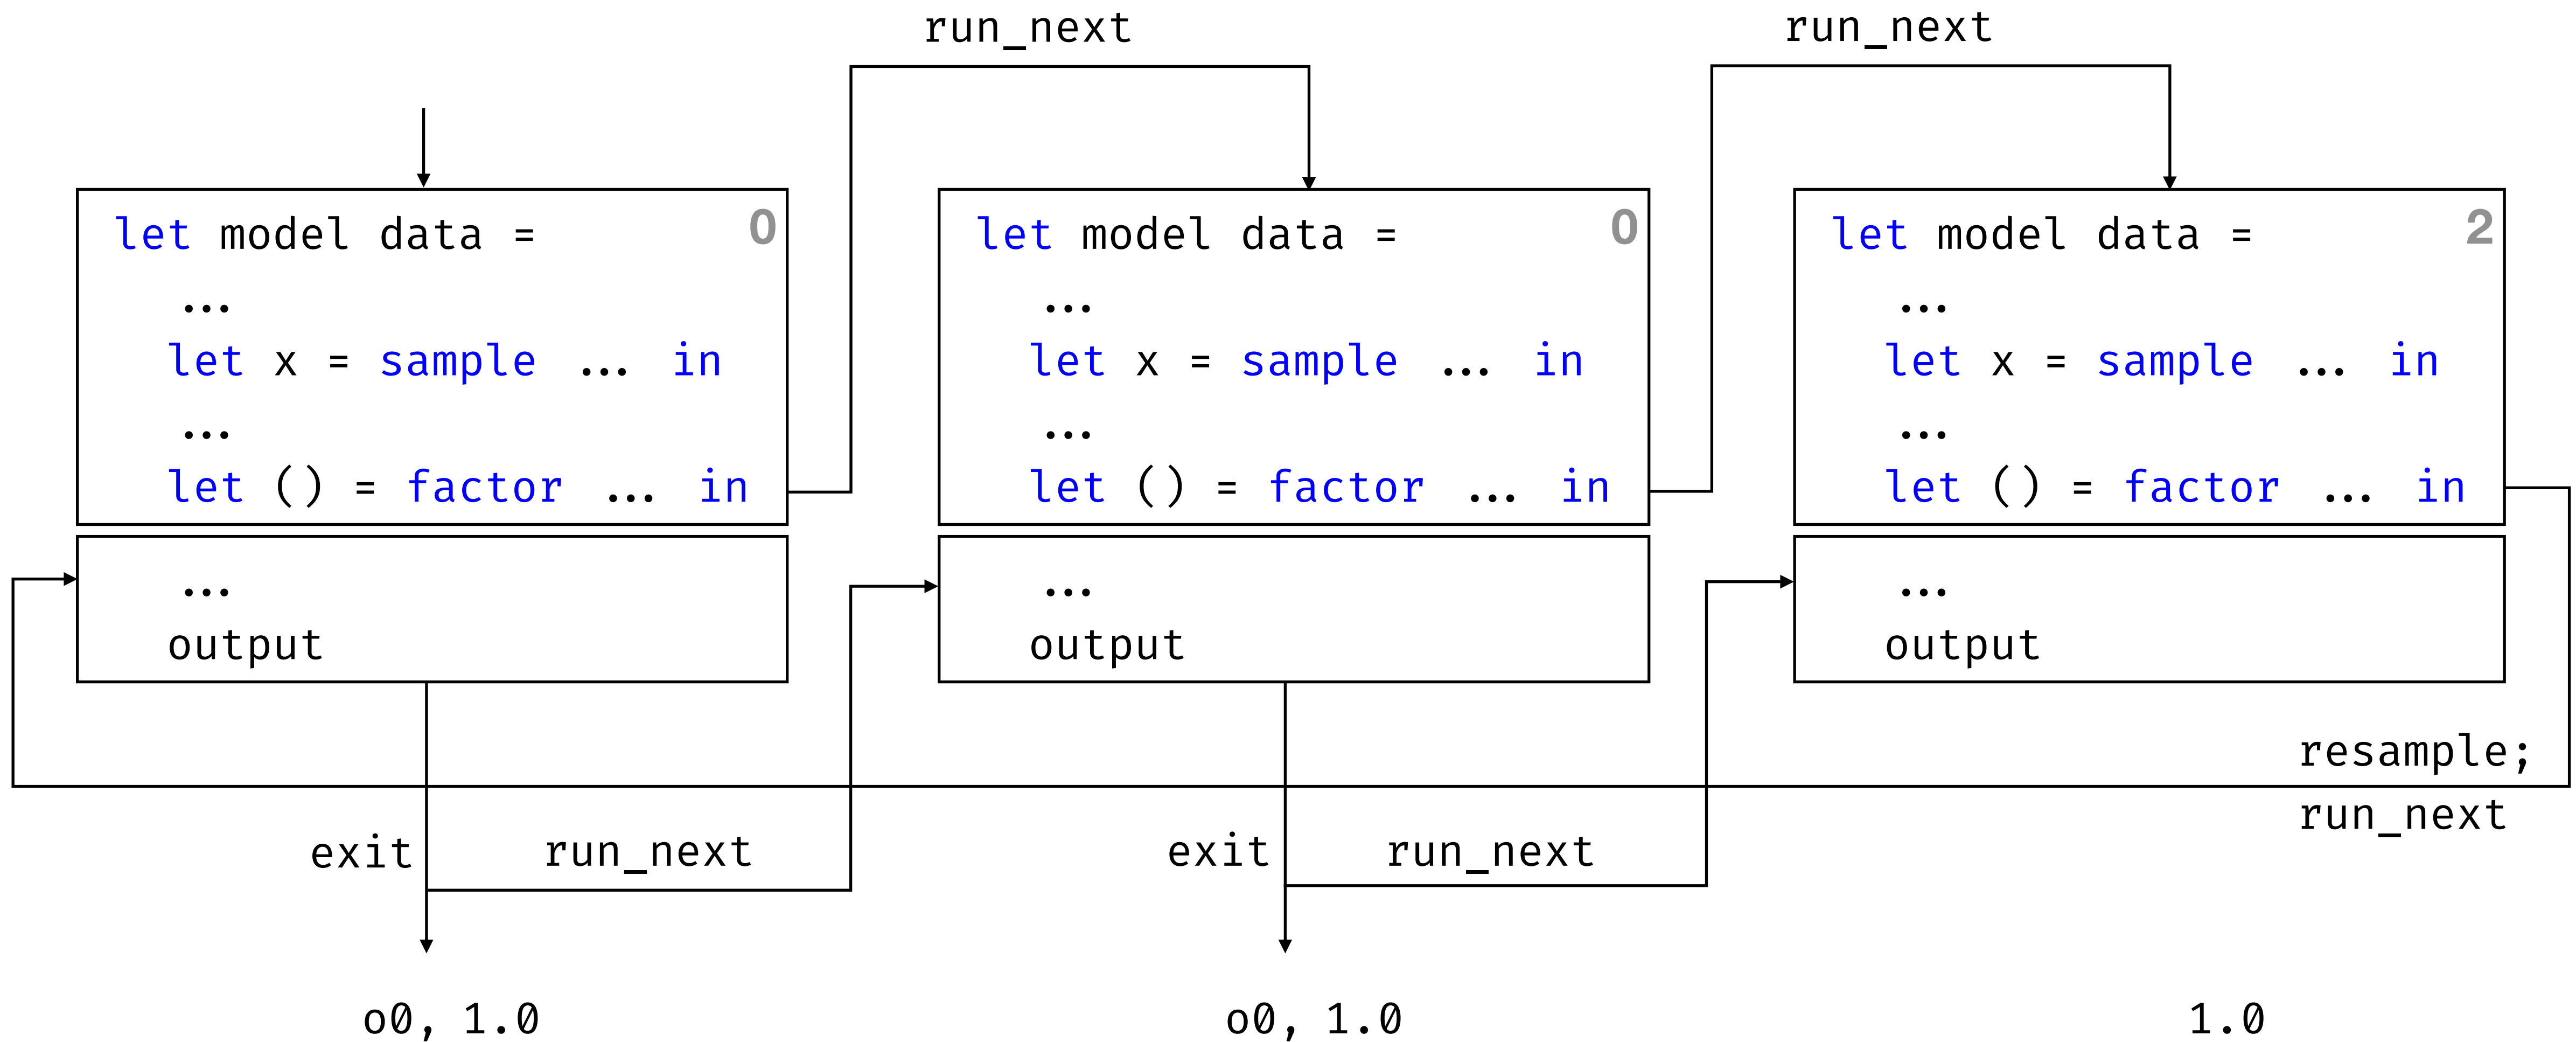
Particle Filter



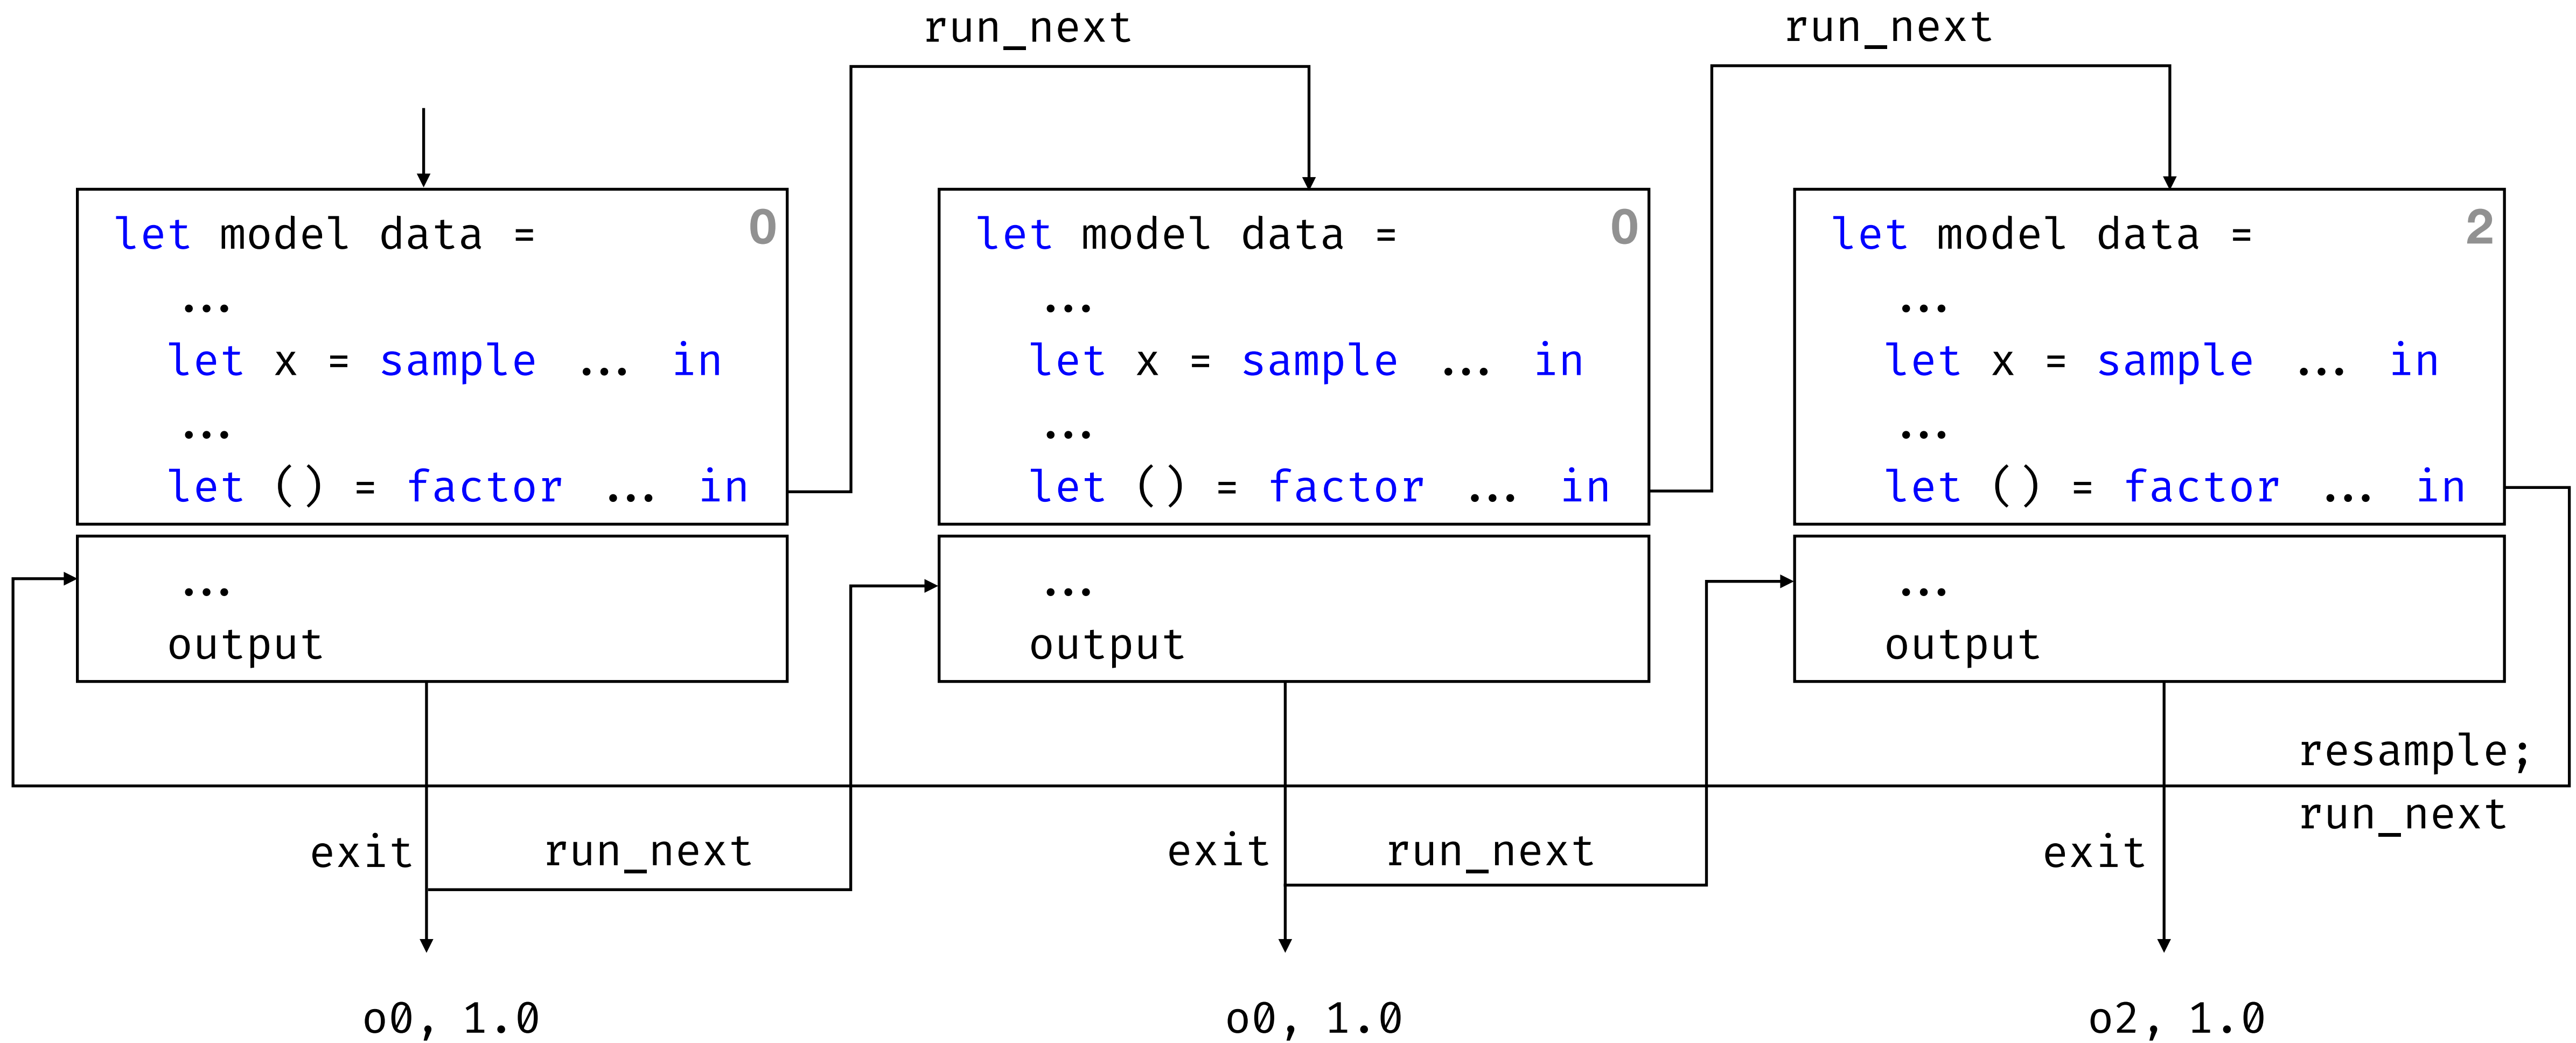
Particle Filter



Particle Filter



Particle Filter



Particle Filter

infer.ml

```
module Particle_filter = struct
  include Importance_sampling

  let resample particles =
    let logits = Array.map (fun x → x.score) particles in
    let values = Array.map (fun x → { x with score = 0. }) particles in
    let dist = Distribution.support ~values ~logits in
    Array.init (Array.length particles) (fun _ → Distribution.draw dist)

  let factor s k prob =
    let particle = prob.particles.(prob.id) in
    prob.particles.(prob.id) ← { particle with k = k (); score = s +. particle.score };
    let prob =
      if prob.id < Array.length prob.particles - 1 then prob
      else { id = -1; particles = resample prob.particles }
    in
    run_next prob
end
```

HMM: Hidden Markov Model

hmm.ml

```
open Infer.Particle_filter

let hmm prob data =
  let rec gen states data =
    match (states, data) with
    | [], y :: data → gen [ y ] data
    | states, [] → return states
    | pre_x :: _, y :: data →
      let* x = sample prob (gaussian ~mu:pre_x ~sigma:1.0) in
      let* () = observe prob (gaussian ~mu:x ~sigma:1.0) y in
      gen (x :: states) data
  in
  gen [] data

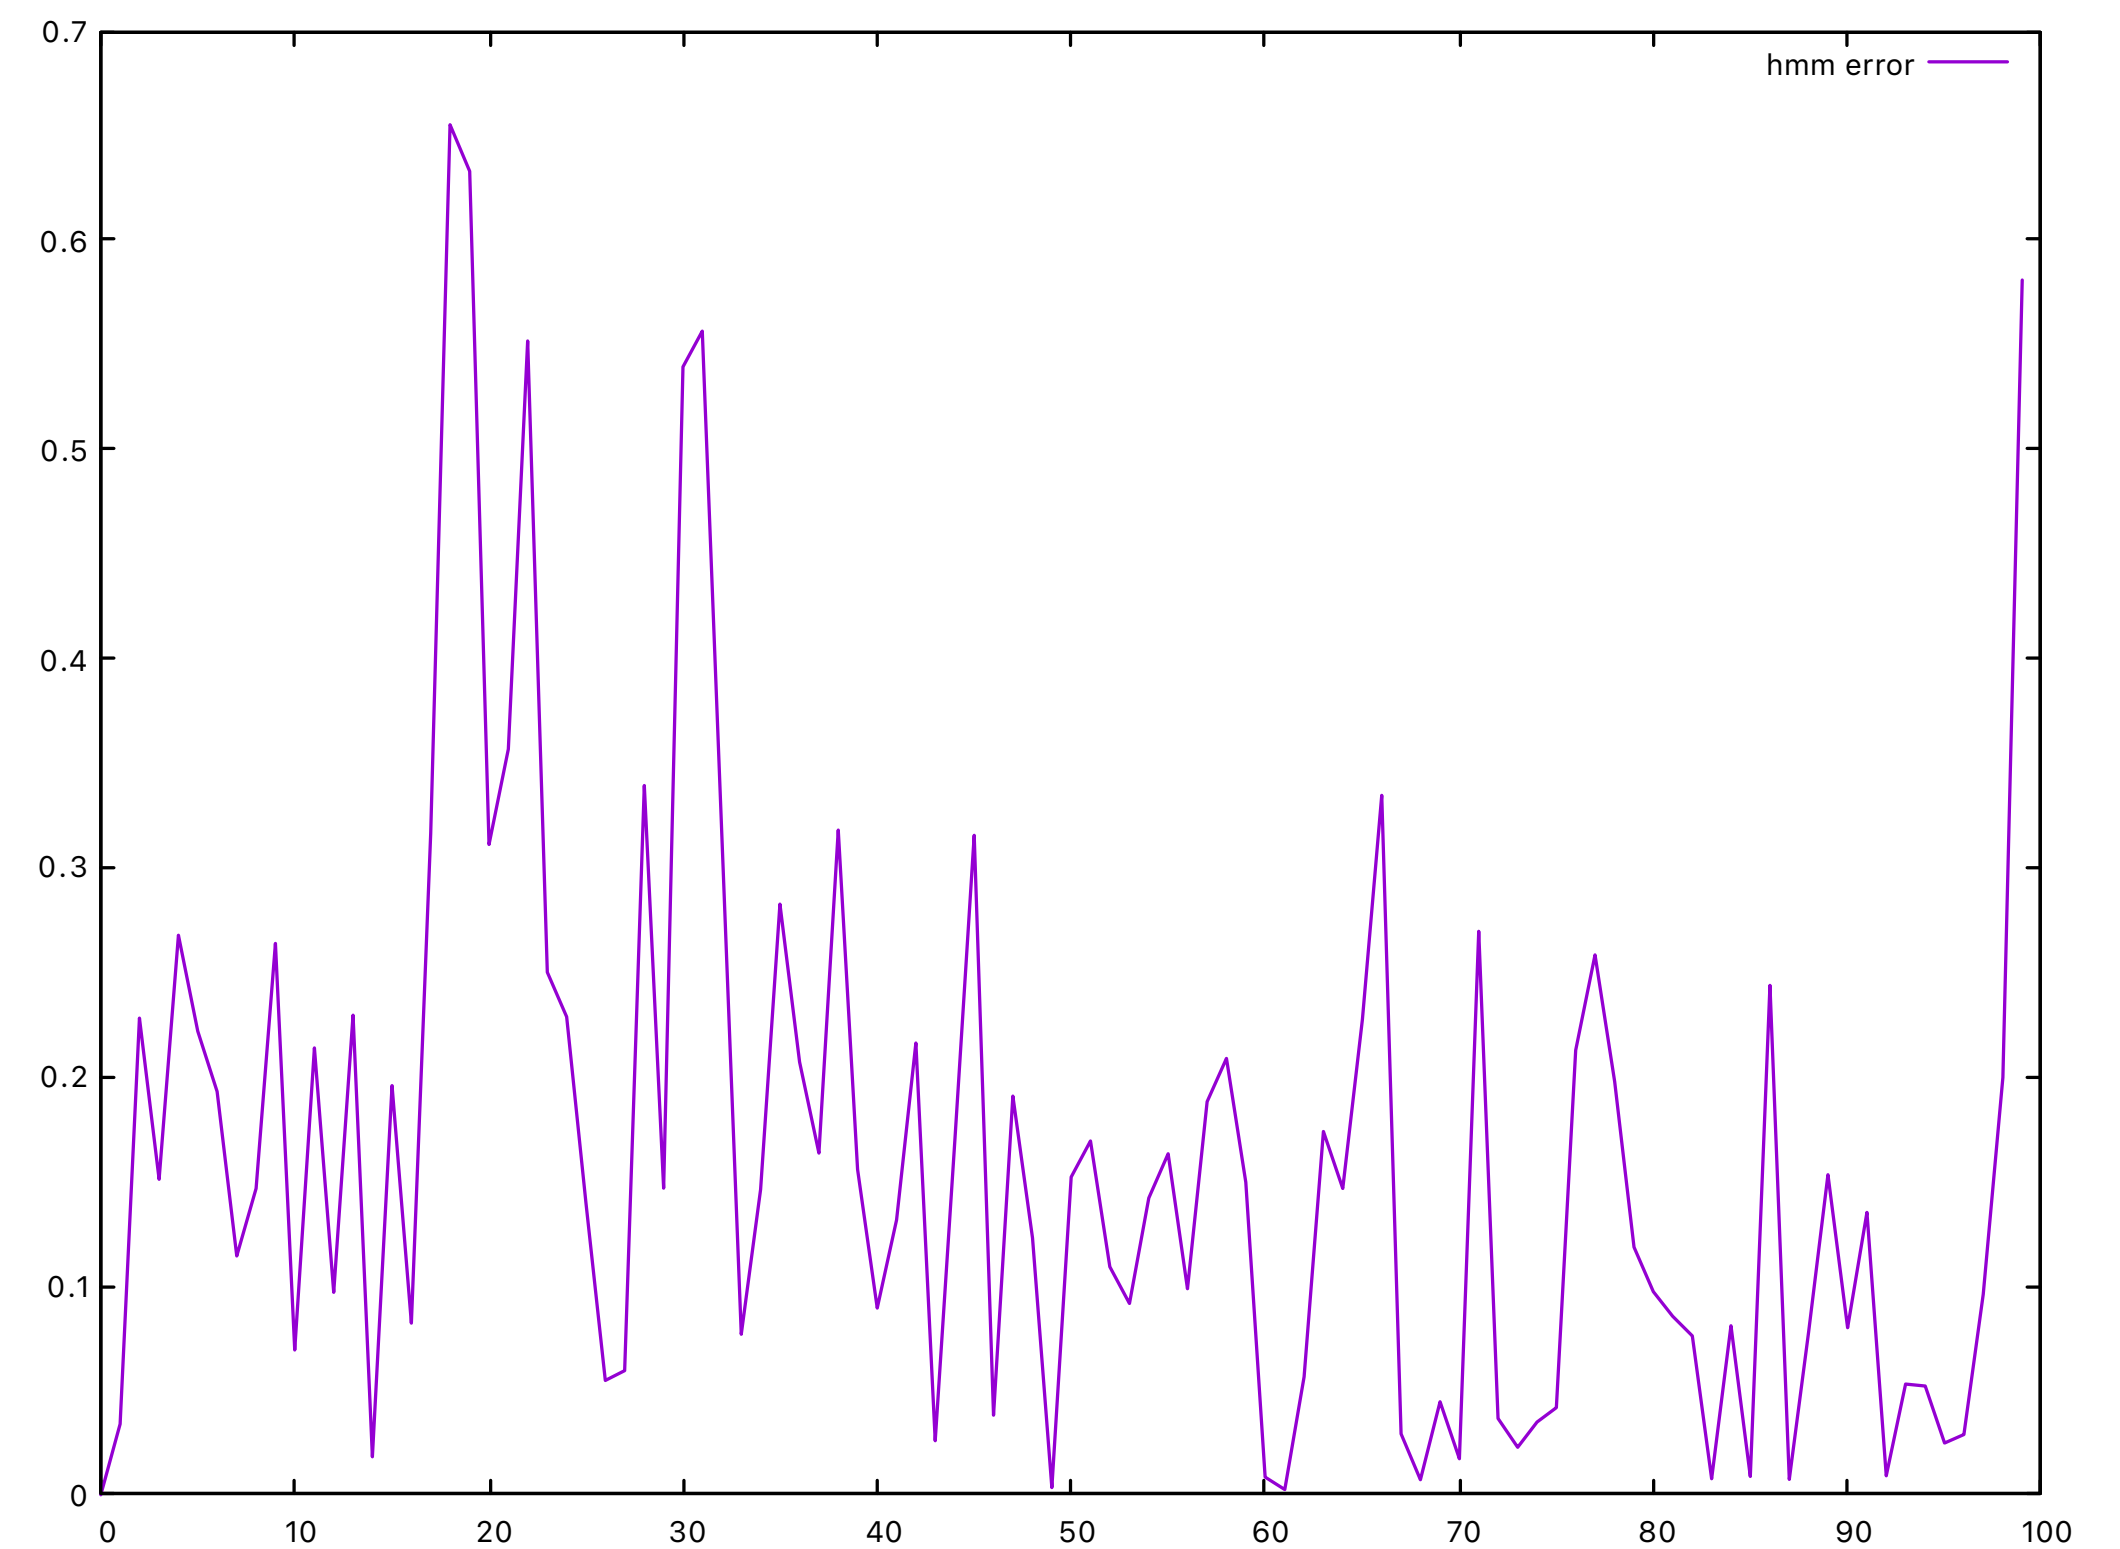
let _ =
  let data = Owl.Arr.linspace 0. 20. 20 ▷ Owl.Arr.to_array ▷ Array.to_list in
  let dist = Distribution.split_list (infer ~n:100 hmm data) in
  let m_x = List.map Distribution.mean dist in
  List.iter2 (Format.printf "%f >> %f") data m_x
```


HMM: Hidden Markov Model

```
› dune exec ./hmm.exe
```

```
0.000000 >> 0.000000
1.052632 >> 0.997546
2.105263 >> 2.300316
3.157895 >> 3.289649
4.210526 >> 4.857555
5.263158 >> 4.907179
6.315789 >> 6.254198
7.368421 >> 7.208341
8.421053 >> 8.432642
9.473684 >> 8.938143
10.526316 >> 9.555007
11.578947 >> 11.098199
12.631579 >> 12.823460
13.684211 >> 13.701444
14.736842 >> 14.934314
15.789474 >> 16.115058
```

...



Conclusion

For a given inference algorithm, how to implement `sample`, `assume`, `factor`, `observe`, and `infer`?

I - Basic inference

- Rejection sampling
- Importance sampling

II - Continuation Passing Style models

III - Inference on CPS models

- Sample generation
- Importance sampling
- Particle filter

References

WebPPL

Noah Goodman and Andreas Stuhlmüller

<http://webppl.org/>

The Design and Implementation of Probabilistic Programming Languages

Noah Goodman and Andreas Stuhlmüller

<http://dippl.org/>

An Introduction to Probabilistic Programming

Jan-Willem van de Meent, Brooks Paige, Hongseok Yang, Frank Wood

<https://arxiv.org/abs/1809.10756>

Embedded probabilistic domain-specific language HANSEI

Oleg Kiselyov, Chung-chieh Shan

<https://okmij.org/ftp/kakuritu/Hansei.html>

BYO-PPL

Build Your Own Probabilistic Language

- Clone the repo: `git clone https://github.com/mpri-probprog/byo-ppl-22-23.git`
- Install the dependencies: `opam install . --deps-only`
- Build the project: `dune build`
- Test an example: `dune exec ./examples/funny_bernoulli.exe`

Implemented as an OCaml embedded domain specific language (eDSL)

- `Distribution`: small library of probability distributions and basic statistical functions.
- `Basic`: basic inference algorithms (rejection sampling inference sampling)
- `Infer`: inference algorithms for models written in Continuation Passing Style (CPS).
- `Cps_operators`: syntactic sugar to write CPS style probabilistic models.
- `Utils`: missing utilities functions used in other modules.